


---

# Amazon EKS

## User Guide



## **Amazon EKS: User Guide**

Copyright © 2019 Amazon Web Services, Inc. and/or its affiliates. All rights reserved.

Amazon's trademarks and trade dress may not be used in connection with any product or service that is not Amazon's, in any manner that is likely to cause confusion among customers, or in any manner that disparages or discredits Amazon. All other trademarks not owned by Amazon are the property of their respective owners, who may or may not be affiliated with, connected to, or sponsored by Amazon.

# Table of Contents

What Is Amazon EKS?	1
Amazon EKS Control Plane Architecture	1
How Does Amazon EKS Work?	2
Getting Started with Amazon EKS	3
Getting Started with eksctl	3
Prerequisites	3
Create Your Amazon EKS Cluster and Worker Nodes	6
Next Steps	8
Getting Started with the Console	8
Amazon EKS Prerequisites	8
Step 1: Create Your Amazon EKS Cluster	11
Step 2: Create a kubeconfig File	13
Step 3: Launch and Configure Amazon EKS Worker Nodes	14
Next Steps	19
Clusters	20
Creating a Cluster	20
Updating Kubernetes Version	26
Cluster Endpoint Access	34
Modifying Cluster Endpoint Access	35
Accessing the API Server from within the VPC	37
Control Plane Logging	38
Enabling and Disabling Control Plane Logs	39
Viewing Cluster Control Plane Logs	40
Deleting a Cluster	41
Kubernetes Versions	43
Available Amazon EKS Kubernetes Versions	43
Kubernetes 1.13	43
Amazon EKS Version Deprecation	44
Platform Versions	44
Kubernetes version 1.13	45
Kubernetes version 1.12	46
Kubernetes version 1.11	46
Kubernetes version 1.10	47
Worker Nodes	49
Amazon EKS-Optimized AMI	49
Amazon EKS-Optimized AMI Build Scripts	52
Amazon EKS-Optimized AMI with GPU Support	52
Partner AMIs	56
Launching Amazon EKS Worker Nodes	56
Worker Node Updates	63
Migrating to a New Worker Node Group	63
Updating an Existing Worker Node Group	68
Storage Classes	73
Load Balancing and Ingress	75
Load Balancing	75
Subnet Tagging for Load Balancers	75
ALB Ingress Controller on Amazon EKS	76
Networking	79
Creating a VPC for Amazon EKS	79
Next Steps	81
Cluster VPC Considerations	81
VPC IP Addressing	82
VPC Tagging Requirement	82
Subnet Tagging Requirement	82

Cluster Security Group Considerations .....	83
Pod Networking .....	85
CNI Configuration Variables .....	87
Installing CoreDNS .....	88
External SNAT .....	91
CNI Custom Networking .....	94
CNI Upgrades .....	97
Installing Calico on Amazon EKS .....	97
Stars Policy Demo .....	98
Managing Cluster Authentication .....	104
Installing kubectl .....	104
Installing aws-iam-authenticator .....	108
Create a kubeconfig for Amazon EKS .....	111
Managing Users or IAM Roles for your Cluster .....	115
eksctl .....	119
Installing or Upgrading eksctl .....	119
Pod Security Policy .....	121
Amazon EKS Default Pod Security Policy .....	121
Guest Book .....	124
Metrics Server .....	128
Prometheus Metrics .....	130
Viewing the Raw Metrics .....	130
Deploying Prometheus .....	130
Using Helm .....	134
Tutorial: Deploy Kubernetes Dashboard .....	137
Prerequisites .....	138
Step 1: Deploy the Dashboard .....	139
Step 2: Create an eks-admin Service Account and Cluster Role Binding .....	140
Step 3: Connect to the Dashboard .....	141
Step 4: Next Steps .....	142
Getting Started with App Mesh and Kubernetes .....	143
Prerequisites .....	143
Step 1: Create Your Service Mesh .....	143
Step 2: Create Your Virtual Nodes .....	143
Step 3: Create Your Virtual Routers .....	144
Step 4: Create Your Routes .....	145
Step 5: Create Your Virtual Services .....	145
Step 6: Updating Your Microservice Pod Specifications .....	146
Tutorial: Configure App Mesh Integration with Kubernetes .....	148
Prerequisites .....	148
Step 1: Install the Controller and Custom Resources .....	148
Step 2: Install the Sidecar Injector .....	149
.....	149
Step 3: Configure App Mesh .....	150
Create Kubernetes Custom Resources .....	150
Sidecar Injection .....	151
Step 4: Remove Integration Components (Optional) .....	152
Deploy a Mesh Connected Service .....	152
Prerequisites .....	152
Deploy a Sample Application .....	152
Run Application .....	154
Change Configuration .....	154
Remove Application .....	155
Deep Learning Containers .....	156
Security .....	157
Identity and Access Management .....	157
Audience .....	158

Authenticating With Identities .....	158
Managing Access Using Policies .....	160
How Amazon EKS Works with IAM .....	161
Identity-Based Policy Examples .....	163
Service IAM Role .....	166
Worker Node IAM Role .....	168
Troubleshooting .....	170
Logging and Monitoring .....	170
Compliance Validation .....	170
Resilience .....	171
Infrastructure Security .....	171
Configuration and Vulnerability Analysis .....	172
CloudTrail .....	173
Amazon EKS Information in CloudTrail .....	173
Understanding Amazon EKS Log File Entries .....	174
Related Projects .....	175
Management Tools .....	175
eksctl .....	175
AWS Service Operator .....	175
Networking .....	175
Amazon VPC CNI plugin for Kubernetes .....	175
AWS Application Load Balancer (ALB) Ingress Controller for Kubernetes .....	175
ExternalDNS .....	176
Security .....	176
AWS IAM Authenticator .....	176
Storage .....	176
Amazon EFS CSI Driver .....	176
Machine Learning .....	176
Kubeflow .....	176
Auto Scaling .....	177
Cluster Autoscaler .....	177
Escalator .....	177
Monitoring .....	177
Prometheus .....	177
Continuous Integration / Continuous Deployment .....	177
Jenkins X .....	177
Troubleshooting .....	179
Insufficient Capacity .....	179
aws-iam-authenticator Not Found .....	179
Worker Nodes Fail to Join Cluster .....	179
Unauthorized or Access Denied (kubectl) .....	179
hostname doesn't match .....	180
getsockopt: no route to host .....	180
CNI Log Collection Tool .....	180
IAM .....	181
AccessDeniedException .....	181
I Am Not Authorized to Perform iam:PassRole .....	181
I Want to View My Access Keys .....	182
I'm an Administrator and Want to Allow Others to Access Amazon EKS .....	182
I Want to Allow People Outside of My AWS Account to Access My Amazon EKS Resources .....	182
Service Limits .....	184
Document History .....	185
AWS Glossary .....	190

# What Is Amazon EKS?

Amazon Elastic Kubernetes Service (Amazon EKS) is a managed service that makes it easy for you to run Kubernetes on AWS without needing to stand up or maintain your own Kubernetes control plane. Kubernetes is an open-source system for automating the deployment, scaling, and management of containerized applications.

Amazon EKS runs Kubernetes control plane instances across multiple Availability Zones to ensure high availability. Amazon EKS automatically detects and replaces unhealthy control plane instances, and it provides automated version upgrades and patching for them.

Amazon EKS is also integrated with many AWS services to provide scalability and security for your applications, including the following:

- Amazon ECR for container images
- Elastic Load Balancing for load distribution
- IAM for authentication
- Amazon VPC for isolation

Amazon EKS runs up-to-date versions of the open-source Kubernetes software, so you can use all the existing plugins and tooling from the Kubernetes community. Applications running on Amazon EKS are fully compatible with applications running on any standard Kubernetes environment, whether running in on-premises data centers or public clouds. This means that you can easily migrate any standard Kubernetes application to Amazon EKS without any code modification required.

## Amazon EKS Control Plane Architecture

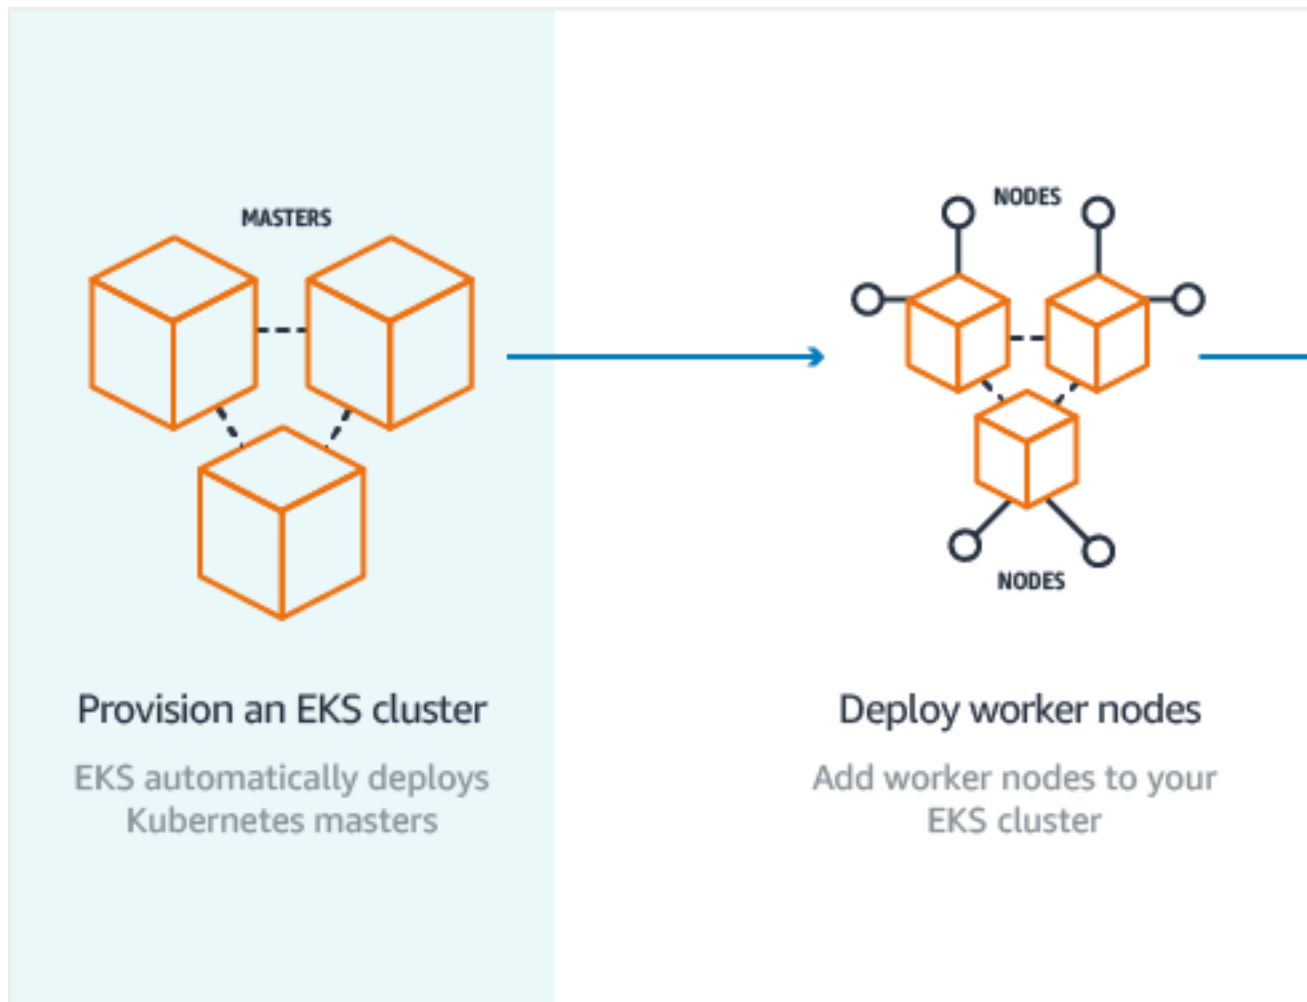
Amazon EKS runs a single tenant Kubernetes control plane for each cluster, and control plane infrastructure is not shared across clusters or AWS accounts.

This control plane consists of at least two API server nodes and three `etcd` nodes that run across three Availability Zones within a Region. Amazon EKS automatically detects and replaces unhealthy control plane instances, restarting them across the Region as needed. Amazon EKS leverages the architecture of AWS Regions in order to maintain high availability. Because of this, Amazon EKS is able to offer an [SLA for API server endpoint availability](#).

Amazon EKS uses Amazon VPC network policies to restrict traffic between control plane components to within a single cluster. Control plane components for a cluster cannot view or receive communication from other clusters or other AWS accounts, except as authorized with Kubernetes RBAC policies.

This secure and highly-available configuration makes Amazon EKS reliable and recommended for production workloads.

## How Does Amazon EKS Work?



Getting started with Amazon EKS is easy:

1. First, create an Amazon EKS cluster in the AWS Management Console or with the AWS CLI or one of the AWS SDKs.
2. Then, launch worker nodes that register with the Amazon EKS cluster. We provide you with an AWS CloudFormation template that automatically configures your nodes.
3. When your cluster is ready, you can configure your favorite Kubernetes tools (such as **kubectl**) to communicate with your cluster.
4. Deploy and manage applications on your Amazon EKS cluster the same way that you would with any other Kubernetes environment.

For more information about creating your required resources and your first Amazon EKS cluster, see [Getting Started with Amazon EKS \(p. 3\)](#).

# Getting Started with Amazon EKS

There are two getting started guides available for creating a new Kubernetes cluster with worker nodes in Amazon EKS:

- [Getting Started with eksctl \(p. 3\)](#): This getting started guide helps you to install all of the required resources to get started with Amazon EKS using `eksctl`, a simple command line utility for creating and managing Kubernetes clusters on Amazon EKS. At the end of this tutorial, you will have a running Amazon EKS cluster with worker nodes, and the `kubectl` command line utility will be configured to use your new cluster. This is the fastest and simplest way to get started with Amazon EKS.
- [Getting Started with the AWS Management Console \(p. 8\)](#): This getting started guide helps you to create all of the required resources to get started with Amazon EKS in the AWS Management Console. In this guide, you manually create each resource in the Amazon EKS or AWS CloudFormation consoles, and the workflow described here gives you complete visibility into how each resource is created and how they interact with each other; however, this is a more complicated and time consuming way to get started with Amazon EKS.

## Getting Started with eksctl

This getting started guide helps you to install all of the required resources to get started with Amazon EKS using `eksctl`, a simple command line utility for creating and managing Kubernetes clusters on Amazon EKS. At the end of this tutorial, you will have a running Amazon EKS cluster with worker nodes, and the `kubectl` command line utility will be configured to use your new cluster.

### Prerequisites

This section helps you to install and configure the binaries you need to create and manage an Amazon EKS cluster.

### Install the Latest AWS CLI

To use `kubectl` with your Amazon EKS clusters, you must install a binary that can create the required client security token for cluster API server communication. The `aws eks get-token` command, available in version 1.16.156 or greater of the AWS CLI, supports client security token creation. To install or upgrade the AWS CLI, see [Installing the AWS Command Line Interface](#) in the *AWS Command Line Interface User Guide*.

If you already have pip and a supported version of Python, you can install or upgrade the AWS CLI with the following command:

```
pip install awscli --upgrade --user
```

#### Note

Your system's Python version must be 2.7.9 or greater. Otherwise, you receive `hostname doesn't match` errors with AWS CLI calls to Amazon EKS. For more information, see [What are "hostname doesn't match" errors?](#) in the Python Requests FAQ.

For more information about other methods of installing or upgrading the AWS CLI for your platform, see the following topics in the *AWS Command Line Interface User Guide*.



- [Install the AWS Command Line Interface on macOS](#)
- [Install the AWS Command Line Interface on Linux](#)
- [Install the AWS Command Line Interface on Microsoft Windows](#)

If you are unable to install version 1.16.156 or greater of the AWS CLI on your system, you must ensure that the AWS IAM Authenticator for Kubernetes is installed on your system. For more information, see [Installing aws-iam-authenticator](#) (p. 108).

## Configure Your AWS CLI Credentials

Both `eksctl` and the AWS CLI require that you have AWS credentials configured in your environment. The `aws configure` command is the fastest way to set up your AWS CLI installation for general use.

```
$ aws configure
AWS Access Key ID [None]: AKIAIOSFODNN7EXAMPLE
AWS Secret Access Key [None]: wJalrXUtnFEMI/K7MDENG/bPxrFiCYEXAMPLEKEY
Default region name [None]: us-west-2
Default output format [None]: json
```

When you type this command, the AWS CLI prompts you for four pieces of information: access key, secret access key, AWS Region, and output format. This information is stored in a profile (a collection of settings) named `default`. This profile is used unless you specify another one.

For more information, see [Configuring the AWS CLI](#) in the *AWS Command Line Interface User Guide*.

## Install eksctl

This section helps you to install the `eksctl` command line utility. For more information, see the <https://eksctl.io/>.

Choose the tab below that best represents your client setup.

macOS

### To install or upgrade eksctl on macOS using Homebrew

The easiest way to get started with Amazon EKS and macOS is by installing `eksctl` with [Homebrew](#). The `eksctl` Homebrew recipe installs `eksctl` and any other dependencies that are required for Amazon EKS, such as `kubectl` and the `aws-iam-authenticator`.

1. If you do not already have Homebrew installed on macOS, install it with the following command.

```
/usr/bin/ruby -e "$(curl -fsSL https://raw.githubusercontent.com/Homebrew/install/master/install)"
```

2. Install the Weaveworks Homebrew tap.

```
brew tap weaveworks/tap
```

3. Install or upgrade `eksctl`.

- Install `eksctl` with the following command:

```
brew install weaveworks/tap/eksctl
```

- If `eksctl` is already installed, run the following command to upgrade:

```
brew upgrade eksctl && brew link --overwrite eksctl
```

4. Test that your installation was successful with the following command.

```
eksctl version
```

**Note**

The `GitTag` version should be at least `0.1.37`. If not, check your terminal output for any installation or upgrade errors.

## Linux

### To install or upgrade `eksctl` on Linux using `curl`

1. Download and extract the latest release of `eksctl` with the following command.

```
curl --silent --location "https://github.com/weaveworks/eksctl/releases/download/latest_release/eksctl_$(uname -s)_amd64.tar.gz" | tar xz -C /tmp
```

2. Move the extracted binary to `/usr/local/bin`.

```
sudo mv /tmp/eksctl /usr/local/bin
```

3. Test that your installation was successful with the following command.

```
eksctl version
```

**Note**

The `GitTag` version should be at least `0.1.37`. If not, check your terminal output for any installation or upgrade errors.

## Windows

### To install or upgrade `eksctl` on Windows using Chocolatey

1. If you do not already have Chocolatey installed on your Windows system, see [Installing Chocolatey](#).
2. Install or upgrade `eksctl` and the `aws-iam-authenticator`.

- Install the binaries with the following command:

```
chocolatey install -y eksctl aws-iam-authenticator
```

- If they are already installed, run the following command to upgrade:

```
chocolatey upgrade -y eksctl aws-iam-authenticator
```

3. Test that your installation was successful with the following command.

```
eksctl version
```

### Note

The `gitTag` version should be at least `0.1.37`. If not, check your terminal output for any installation or upgrade errors.

## Install and Configure `kubectl` for Amazon EKS

Kubernetes uses the `kubectl` command-line utility for communicating with the cluster API server.

### Note

If you used the preceding Homebrew instructions to install `eksctl` on macOS, then `kubectl` and the `aws-iam-authenticator` have already been installed on your system. You can skip to [Create Your Amazon EKS Cluster and Worker Nodes \(p. 6\)](#).

### To install `kubectl` for Amazon EKS

- You have multiple options to download and install **`kubectl`** for your operating system.
  - The `kubectl` binary is available in many operating system package managers, and this option is often much easier than a manual download and install process. You can follow the instructions for your specific operating system or package manager in the [Kubernetes documentation](#) to install.
  - Amazon EKS also vends **`kubectl`** binaries that you can use that are identical to the upstream **`kubectl`** binaries with the same version. To install the Amazon EKS-vended binary for your operating system, see [Installing `kubectl` \(p. 104\)](#).

## Create Your Amazon EKS Cluster and Worker Nodes

Now you can create your Amazon EKS cluster and a worker node group with the `eksctl` command line utility.

### To create your cluster and worker nodes with `eksctl`

This procedure assumes that you have installed `eksctl`, and that your `eksctl` version is at least `0.1.37`. You can check your version with the following command:

```
eksctl version
```

For more information on installing or upgrading `eksctl`, see [Installing or Upgrading `eksctl` \(p. 119\)](#).

- Create your Amazon EKS cluster and worker nodes with the following command. Substitute the red text with your own values.

### Important

Kubernetes version 1.10 is no longer supported on Amazon EKS. You can no longer create new 1.10 clusters, and all existing Amazon EKS clusters running Kubernetes version 1.10 will eventually be automatically updated to the latest available platform version of Kubernetes version 1.11. For more information, see [Amazon EKS Version Deprecation \(p. 44\)](#).

Please update any 1.10 clusters to version 1.11 or higher in order to avoid service interruption. For more information, see [Updating an Amazon EKS Cluster Kubernetes Version \(p. 26\)](#).

```
eksctl create cluster \  
--name prod \  
--version 1.13 \  
--nodegroup-name standard-workers \  
--node-type t3.medium \  

```

```
--nodes 3 \  
--nodes-min 1 \  
--nodes-max 4 \  
--node-ami auto
```

### Note

For more information on the available options for **eksctl create cluster**, see the project [README on GitHub](#) or view the help page with the following command.

```
eksctl create cluster --help
```

Output:

```
[#] using region us-west-2  
[#] setting availability zones to [us-west-2b us-west-2c us-west-2d]  
[#] subnets for us-west-2b - public:192.168.0.0/19 private:192.168.96.0/19  
[#] subnets for us-west-2c - public:192.168.32.0/19 private:192.168.128.0/19  
[#] subnets for us-west-2d - public:192.168.64.0/19 private:192.168.160.0/19  
[#] nodegroup "standard-workers" will use "ami-0923e4b35a30a5f53" [AmazonLinux2/1.12]  
[#] creating EKS cluster "prod" in "us-west-2" region  
[#] will create 2 separate CloudFormation stacks for cluster itself and the initial  
    nodegroup  
[#] if you encounter any issues, check CloudFormation console or try 'eksctl utils  
    describe-stacks --region=us-west-2 --name=prod'  
[#] building cluster stack "eksctl-prod-cluster"  
[#] creating nodegroup stack "eksctl-prod-nodegroup-standard-workers"  
[#] all EKS cluster resource for "prod" had been created  
[#] saved kubeconfig as "/Users/ericn/.kube/config"  
[#] adding role "arn:aws:iam::111122223333:role/eksctl-prod-nodegroup-standard-wo-  
NodeInstanceRole-IJP4S12W3020" to auth ConfigMap  
[#] nodegroup "standard-workers" has 0 node(s)  
[#] waiting for at least 1 node(s) to become ready in "standard-workers"  
[#] nodegroup "standard-workers" has 2 node(s)  
[#] node "ip-192-168-22-17.us-west-2.compute.internal" is not ready  
[#] node "ip-192-168-32-184.us-west-2.compute.internal" is ready  
[#] kubectl command should work with "/Users/ericn/.kube/config", try 'kubectl get  
    nodes'  
[#] EKS cluster "prod" in "us-west-2" region is ready
```

- Cluster provisioning usually takes between 10 and 15 minutes. When your cluster is ready, test that your **kubectl** configuration is correct.

```
kubectl get svc
```

### Note

If you receive the error "aws-iam-authenticator": executable file not found in \$PATH, your **kubectl** isn't configured for Amazon EKS. For more information, see [Installing aws-iam-authenticator \(p. 108\)](#).

If you receive any other authorization or resource type errors, see [Unauthorized or Access Denied \(kubectl\) \(p. 179\)](#) in the troubleshooting section.

Output:

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
svc/kubernetes	ClusterIP	10.100.0.1	<none>	443/TCP	1m

- (GPU workers only) If you chose a P2 or P3 instance type and the Amazon EKS-optimized AMI with GPU support, you must apply the [NVIDIA device plugin for Kubernetes](#) as a DaemonSet on your cluster with the following command.

```
kubectl apply -f https://raw.githubusercontent.com/NVIDIA/k8s-device-plugin/1.0.0-beta/nvidia-device-plugin.yml
```

## Next Steps

Now that you have a working Amazon EKS cluster with worker nodes, you are ready to start installing Kubernetes add-ons and deploying applications to your cluster. The following documentation topics help you to extend the functionality of your cluster.

- [Launch a Guest Book Application \(p. 124\)](#) — Create a sample guest book application to test your cluster.
- [Tutorial: Deploy the Kubernetes Web UI \(Dashboard\) \(p. 137\)](#) — This tutorial guides you through deploying the [Kubernetes dashboard](#) to your cluster.
- [Using Helm with Amazon EKS \(p. 134\)](#) — The `helm` package manager for Kubernetes helps you install and manage applications on your cluster.
- [Installing the Kubernetes Metrics Server \(p. 128\)](#) — The Kubernetes metrics server is an aggregator of resource usage data in your cluster.
- [Control Plane Metrics with Prometheus \(p. 130\)](#) — This topic helps you deploy Prometheus into your cluster with `helm`.

## Getting Started with the AWS Management Console

This getting started guide helps you to create all of the required resources to get started with Amazon EKS in the AWS Management Console. In this guide, you manually create each resource in the Amazon EKS or AWS CloudFormation consoles, and the workflow described here gives you complete visibility into how each resource is created and how they interact with each other.

For a simpler and more automated getting started experience, see [Getting Started with `eksctl` \(p. 3\)](#).

## Amazon EKS Prerequisites

Before you can create an Amazon EKS cluster, you must create an IAM role that Kubernetes can assume to create AWS resources. For example, when a load balancer is created, Kubernetes assumes the role to create an Elastic Load Balancing load balancer in your account. This only needs to be done one time and can be used for multiple EKS clusters.

You must also create a VPC and a security group for your cluster to use. Although the VPC and security groups can be used for multiple EKS clusters, we recommend that you use a separate VPC for each EKS cluster to provide better network isolation.

This section also helps you to install the `kubectl` binary and configure it to work with Amazon EKS.

## Create your Amazon EKS Service Role

**To create your Amazon EKS service role in the IAM console**

1. Open the IAM console at <https://console.aws.amazon.com/iam/>.
2. Choose **Roles**, then **Create role**.

3. Choose **EKS** from the list of services, then **Allows Amazon EKS to manage your clusters on your behalf** for your use case, then **Next: Permissions**.
4. Choose **Next: Tags**.
5. (Optional) Add metadata to the role by attaching tags as key-value pairs. For more information about using tags in IAM, see [Tagging IAM Entities](#) in the *IAM User Guide*.
6. Choose **Next: Review**.
7. For **Role name**, enter a unique name for your role, such as `eksServiceRole`, then choose **Create role**.

## Create your Amazon EKS Cluster VPC

This section guides you through creating a VPC for your cluster with either 3 public subnets, or two public subnets and two private subnets, which are provided with internet access through a NAT gateway. We recommend a network architecture that uses private subnets for your worker nodes, and public subnets for Kubernetes to create public load balancers within.

Choose the tab below that represents your desired VPC configuration.

Only public subnets

### To create your cluster VPC with only public subnets

1. Open the AWS CloudFormation console at <https://console.aws.amazon.com/cloudformation>.
2. From the navigation bar, select a Region that supports Amazon EKS.
3. Choose **Create stack**.
4. For **Choose a template**, select **Specify an Amazon S3 template URL**.
5. Paste the following URL into the text area and choose **Next**:

```
https://amazon-eks.s3-us-west-2.amazonaws.com/cloudformation/2019-02-11/amazon-eks-vpc-sample.yaml
```

6. On the **Specify Details** page, fill out the parameters accordingly, and then choose **Next**.
  - **Stack name**: Choose a stack name for your AWS CloudFormation stack. For example, you can call it **eks-vpc**.
  - **VpcBlock**: Choose a CIDR range for your VPC. You can keep the default value.
  - **Subnet01Block**: Specify a CIDR range for subnet 1. We recommend that you keep the default value so that you have plenty of IP addresses for pods to use.
  - **Subnet02Block**: Specify a CIDR range for subnet 2. We recommend that you keep the default value so that you have plenty of IP addresses for pods to use.
  - **Subnet03Block**: Specify a CIDR range for subnet 3. We recommend that you keep the default value so that you have plenty of IP addresses for pods to use.
7. (Optional) On the **Options** page, tag your stack resources. Choose **Next**.
8. On the **Review** page, choose **Create**.
9. When your stack is created, select it in the console and choose **Outputs**.
10. Record the **SecurityGroups** value for the security group that was created. You need this when you create your EKS cluster; this security group is applied to the cross-account elastic network interfaces that are created in your subnets that allow the Amazon EKS control plane to communicate with your worker nodes.
11. Record the **VpcId** for the VPC that was created. You need this when you launch your worker node group template.
12. Record the **SubnetIds** for the subnets that were created. You need this when you create your EKS cluster; these are the subnets that your worker nodes are launched into.

## Public and private subnets

### To create your cluster VPC with public and private subnets

1. Open the AWS CloudFormation console at <https://console.aws.amazon.com/cloudformation>.
2. From the navigation bar, select a Region that supports Amazon EKS.
3. Choose **Create stack**.
4. For **Choose a template**, select **Specify an Amazon S3 template URL**.
5. Paste the following URL into the text area and choose **Next**:

```
https://amazon-eks.s3-us-west-2.amazonaws.com/cloudformation/2019-02-11/amazon-eks-vpc-private-subnets.yaml
```

6. On the **Specify Details** page, fill out the parameters accordingly, and then choose **Next**.
  - **Stack name**: Choose a stack name for your AWS CloudFormation stack. For example, you can call it **eks-vpc**.
  - **VpcBlock**: Choose a CIDR range for your VPC. You can keep the default value.
  - **PublicSubnet01Block**: Specify a CIDR range for public subnet 1. We recommend that you keep the default value so that you have plenty of IP addresses for pods to use.
  - **PublicSubnet02Block**: Specify a CIDR range for public subnet 2. We recommend that you keep the default value so that you have plenty of IP addresses for pods to use.
  - **PrivateSubnet01Block**: Specify a CIDR range for private subnet 1. We recommend that you keep the default value so that you have plenty of IP addresses for pods to use.
  - **PrivateSubnet02Block**: Specify a CIDR range for private subnet 2. We recommend that you keep the default value so that you have plenty of IP addresses for pods to use.
7. (Optional) On the **Options** page, tag your stack resources. Choose **Next**.
8. On the **Review** page, choose **Create**.
9. When your stack is created, select it in the console and choose **Outputs**.
10. Record the **SecurityGroups** value for the security group that was created. You need this when you create your EKS cluster; this security group is applied to the cross-account elastic network interfaces that are created in your subnets that allow the Amazon EKS control plane to communicate with your worker nodes.
11. Record the **VpcId** for the VPC that was created. You need this when you launch your worker node group template.
12. Record the **SubnetIds** for the subnets that were created. You need this when you create your EKS cluster; these are the subnets that your worker nodes are launched into.
13. Tag your private subnets so that Kubernetes knows that it can use them for internal load balancers.
  - a. Open the Amazon VPC console at <https://console.aws.amazon.com/vpc/>.
  - b. Choose **Subnets** in the left navigation.
  - c. Select one of the private subnets for your Amazon EKS cluster's VPC (you can filter them with the string `PrivateSubnet`), and choose the **Tags** tab, and then **Add/Edit Tags**.
  - d. Choose **Create Tag** and add the following key and value, and then choose **Save**.

Key	Value
kubernetes.io/role/internal-elb	1

- e. Repeat these substeps for each private subnet in your VPC.

## Install and Configure **kubectl** for Amazon EKS

Kubernetes uses a command-line utility called `kubectl` for communicating with the cluster API server.

### To install **kubectl** for Amazon EKS

- You have multiple options to download and install **kubectl** for your operating system.
  - The `kubectl` binary is available in many operating system package managers, and this option is often much easier than a manual download and install process. You can follow the instructions for your specific operating system or package manager in the [Kubernetes documentation](#) to install.
  - Amazon EKS also vends **kubectl** binaries that you can use that are identical to the upstream **kubectl** binaries with the same version. To install the Amazon EKS-vended binary for your operating system, see [Installing `kubectl` \(p. 104\)](#).

## Install the Latest AWS CLI

To use `kubectl` with your Amazon EKS clusters, you must install a binary that can create the required client security token for cluster API server communication. The **aws eks get-token** command, available in version 1.16.156 or greater of the AWS CLI, supports client security token creation. To install or upgrade the AWS CLI, see [Installing the AWS Command Line Interface](#) in the *AWS Command Line Interface User Guide*.

### Important

Package managers such **yum**, **apt-get**, or Homebrew for macOS are often behind several versions of the AWS CLI. To ensure that you have the latest version, see [Installing the AWS Command Line Interface](#) in the *AWS Command Line Interface User Guide*.

You can check your AWS CLI version with the following command:

```
aws --version
```

### Note

Your system's Python version must be 2.7.9 or greater. Otherwise, you receive `hostname doesn't match` errors with AWS CLI calls to Amazon EKS. For more information, see [What are "hostname doesn't match" errors?](#) in the Python Requests FAQ.

If you are unable to install version 1.16.156 or greater of the AWS CLI on your system, you must ensure that the AWS IAM Authenticator for Kubernetes is installed on your system. For more information, see [Installing `aws-iam-authenticator` \(p. 108\)](#).

## Step 1: Create Your Amazon EKS Cluster

Now you can create your Amazon EKS cluster.

### Important

When an Amazon EKS cluster is created, the IAM entity (user or role) that creates the cluster is added to the Kubernetes RBAC authorization table as the administrator (with `system:master` permissions. Initially, only that IAM user can make calls to the Kubernetes API server using **kubectl**. For more information, see [Managing Users or IAM Roles for your Cluster \(p. 115\)](#). If you use the console to create the cluster, you must ensure that the same IAM user credentials are in the AWS SDK credential chain when you are running **kubectl** commands on your cluster. If you install and configure the AWS CLI, you can configure the IAM credentials for your user. If the AWS CLI is configured properly for your user, then `eksctl` and the [AWS IAM Authenticator](#)



for [Kubernetes](#) can find those credentials as well. For more information, see [Configuring the AWS CLI](#) in the *AWS Command Line Interface User Guide*.

### To create your cluster with the console

1. Open the Amazon EKS console at <https://console.aws.amazon.com/eks/home#/clusters>.
2. Choose **Create cluster**.

#### Note

If your IAM user does not have administrative privileges, you must explicitly add permissions for that user to call the Amazon EKS API operations. For more information, see [Amazon EKS Identity-Based Policy Examples](#) (p. 163).

3. On the **Create cluster** page, fill in the following fields and then choose **Create**:

- **Cluster name:** A unique name for your cluster.
- **Kubernetes version:** The version of Kubernetes to use for your cluster. By default, the latest available version is selected.
- **Role ARN:** Select the IAM role that you created with [Create your Amazon EKS Service Role](#) (p. 8).
- **VPC:** The VPC you created with [Create your Amazon EKS Cluster VPC](#) (p. 9). You can find the name of your VPC in the drop-down list.
- **Subnets:** The **SubnetIds** values (comma-separated) from the AWS CloudFormation output that you generated with [Create your Amazon EKS Cluster VPC](#) (p. 9). Specify all subnets that will host resources for your cluster (such as private subnets for worker nodes and public subnets for load balancers). By default, the available subnets in the VPC specified in the previous field are preselected.
- **Security Groups:** The **SecurityGroups** value from the AWS CloudFormation output that you generated with [Create your Amazon EKS Cluster VPC](#) (p. 9). This security group has **ControlPlaneSecurityGroup** in the drop-down name.

#### Important

The worker node AWS CloudFormation template modifies the security group that you specify here, so **Amazon EKS strongly recommends that you use a dedicated security group for each cluster control plane (one per cluster)**. If this security group is shared with other resources, you might block or disrupt connections to those resources.

- **Endpoint private access:** Choose whether to enable or disable private access for your cluster's Kubernetes API server endpoint. If you enable private access, Kubernetes API requests that originate from within your cluster's VPC will use the private VPC endpoint. For more information, see [Amazon EKS Cluster Endpoint Access Control](#) (p. 34).
- **Endpoint public access:** Choose whether to enable or disable public access for your cluster's Kubernetes API server endpoint. If you disable public access, your cluster's Kubernetes API server can only receive requests from within the cluster VPC. For more information, see [Amazon EKS Cluster Endpoint Access Control](#) (p. 34).
- **Logging** – For each individual log type, choose whether the log type should be **Enabled** or **Disabled**. By default, each log type is **Disabled**. For more information, see [Amazon EKS Control Plane Logging](#) (p. 38)

#### Note

You might receive an error that one of the Availability Zones in your request doesn't have sufficient capacity to create an Amazon EKS cluster. If this happens, the error output contains the Availability Zones that can support a new cluster. Retry creating your cluster with at least two subnets that are located in the supported Availability Zones for your account. For more information, see [Insufficient Capacity](#) (p. 179).

4. On the **Clusters** page, choose the name of your newly created cluster to view the cluster information.

5. The **Status** field shows **CREATING** until the cluster provisioning process completes. Cluster provisioning usually takes between 10 and 15 minutes.

## Step 2: Create a kubeconfig File

In this section, you create a kubeconfig file for your cluster with the AWS CLI **update-kubeconfig** command. If you do not want to install the AWS CLI, or if you would prefer to create or update your kubeconfig manually, see [Create a kubeconfig for Amazon EKS \(p. 111\)](#).

### To create your kubeconfig file with the AWS CLI

1. Ensure that you have at least version 1.16.156 of the AWS CLI installed. To install or upgrade the AWS CLI, see [Installing the AWS Command Line Interface](#) in the *AWS Command Line Interface User Guide*.

#### Note

Your system's Python version must be 2.7.9 or greater. Otherwise, you receive `hostname doesn't match` errors with AWS CLI calls to Amazon EKS. For more information, see [What are "hostname doesn't match" errors?](#) in the Python Requests FAQ.

You can check your AWS CLI version with the following command:

```
aws --version
```

#### Important

Package managers such **yum**, **apt-get**, or Homebrew for macOS are often behind several versions of the AWS CLI. To ensure that you have the latest version, see [Installing the AWS Command Line Interface](#) in the *AWS Command Line Interface User Guide*.

2. Use the AWS CLI **update-kubeconfig** command to create or update your kubeconfig for your cluster.
  - By default, the resulting configuration file is created at the default kubeconfig path (`.kube/config`) in your home directory or merged with an existing kubeconfig at that location. You can specify another path with the `--kubeconfig` option.
  - You can specify an IAM role ARN with the `--role-arn` option to use for authentication when you issue **kubectl** commands. Otherwise, the IAM entity in your default AWS CLI or SDK credential chain is used. You can view your default AWS CLI or SDK identity by running the **aws sts get-caller-identity** command.
  - For more information, see the help page with the **aws eks update-kubeconfig help** command or see [update-kubeconfig](#) in the *AWS CLI Command Reference*.

```
aws eks --region region update-kubeconfig --name cluster_name
```

3. Test your configuration.

```
kubectl get svc
```

#### Note

If you receive the error `"aws-iam-authenticator": executable file not found in $PATH`, your **kubectl** isn't configured for Amazon EKS. For more information, see [Installing aws-iam-authenticator \(p. 108\)](#).

If you receive any other authorization or resource type errors, see [Unauthorized or Access Denied \(kubectl\) \(p. 179\)](#) in the troubleshooting section.

Output:

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
svc/kubernetes	ClusterIP	10.100.0.1	<none>	443/TCP	1m

## Step 3: Launch and Configure Amazon EKS Worker Nodes

Now that your VPC and Kubernetes control plane are created, you can launch and configure your worker nodes.

### Important

Amazon EKS worker nodes are standard Amazon EC2 instances, and you are billed for them based on normal Amazon EC2 On-Demand Instance prices. For more information, see [Amazon EC2 Pricing](#).

### To launch your worker nodes

1. Wait for your cluster status to show as **ACTIVE**. If you launch your worker nodes before the cluster is active, the worker nodes will fail to register with the cluster and you will have to relaunch them.
2. Open the AWS CloudFormation console at <https://console.aws.amazon.com/cloudformation>.
3. From the navigation bar, select a Region that supports Amazon EKS.
4. Choose **Create stack**.
5. For **Choose a template**, select **Specify an Amazon S3 template URL**.
6. Paste the following URL into the text area and choose **Next**:

```
https://amazon-eks.s3-us-west-2.amazonaws.com/cloudformation/2019-02-11/amazon-eks-nodegroup.yaml
```

### Note

If you intend to only deploy worker nodes to private subnets, you should edit this template in the AWS CloudFormation designer and modify the `AssociatePublicIpAddress` parameter in the `NodeLaunchConfig` to be `false`.

```
AssociatePublicIpAddress: 'false'
```

7. On the **Specify Details** page, fill out the following parameters accordingly, and choose **Next**.
  - **Stack name:** Choose a stack name for your AWS CloudFormation stack. For example, you can call it **<cluster-name>-worker-nodes**.
  - **ClusterName:** Enter the name that you used when you created your Amazon EKS cluster.

### Important

This name must exactly match the name you used in [Step 1: Create Your Amazon EKS Cluster \(p. 11\)](#); otherwise, your worker nodes cannot join the cluster.

- **ClusterControlPlaneSecurityGroup:** Choose the **SecurityGroups** value from the AWS CloudFormation output that you generated with [Create your Amazon EKS Cluster VPC \(p. 9\)](#).
- **NodeGroupName:** Enter a name for your node group. This name can be used later to identify the Auto Scaling node group that is created for your worker nodes.
- **NodeAutoScalingGroupMinSize:** Enter the minimum number of nodes that your worker node Auto Scaling group can scale in to.
- **NodeAutoScalingGroupDesiredCapacity:** Enter the desired number of nodes to scale to when your stack is created.

- **NodeAutoScalingGroupMaxSize:** Enter the maximum number of nodes that your worker node Auto Scaling group can scale out to.
- **NodeInstanceType:** Choose an instance type for your worker nodes.

**Important**

Some instance types might not be available in all regions.

- **NodeImageId:** Enter the current Amazon EKS worker node AMI ID for your Region. The AMI IDs for the latest Amazon EKS-optimized AMI (with and without [GPU support \(p. 52\)](#)) are shown in the following table.

**Note**

The Amazon EKS-optimized AMI with GPU support only supports P2 and P3 instance types. Be sure to specify these instance types in your worker node AWS CloudFormation template. By using the Amazon EKS-optimized AMI with GPU support, you agree to [NVIDIA's end user license agreement \(EULA\)](#).

Kubernetes version 1.13.7

Region	Amazon EKS-optimized AMI	with GPU support
US East (Ohio) (us-east-2)	ami-0485258c2d1c3608f	ami-0ccac9d9b57864000
US East (N. Virginia) (us-east-1)	ami-0f2e8e5663e16b436	ami-0017d945a10387606
US West (Oregon) (us-west-2)	ami-03a55127c613349a7	ami-08335952e837d087b
Asia Pacific (Hong Kong) (ap-east-1)	ami-032850771ac6f8ae2	N/A*
Asia Pacific (Mumbai) (ap-south-1)	ami-0a9b1c1807b1a40ab	ami-005b754faac73f0cc
Asia Pacific (Tokyo) (ap-northeast-1)	ami-0fde798d17145fae1	ami-04cf69bbd6c0fae0b
Asia Pacific (Seoul) (ap-northeast-2)	ami-07fd7609df6c8e39b	ami-0730e699ed0118737
Asia Pacific (Singapore) (ap-southeast-1)	ami-0361e14efd56a71c7	ami-07be5e97a529cd146
Asia Pacific (Sydney) (ap-southeast-2)	ami-0237d87bc27daba65	ami-0a2f4c3aeb596aa7e
EU (Frankfurt) (eu-central-1)	ami-0b7127e7a2a38802a	ami-0fbbd205f797eccccd
EU (Ireland) (eu-west-1)	ami-00ac2e6b3cb38a9b9	ami-0f9571a3e65dc4e20
EU (London) (eu-west-2)	ami-0147919d2ff9a6ad5	ami-032348bd69c5dd665
EU (Paris) (eu-west-3)	ami-0537ee9329c1628a2	ami-053962359d6859fec
EU (Stockholm) (eu-north-1)	ami-0fd05922165907b85	ami-0641def7f02a4cac5

Kubernetes version 1.12.7

Region	Amazon EKS-optimized AMI	with GPU support
US East (Ohio) (us-east-2)	ami-0fe61ae4c397e710d	ami-067d88fb64d3d7990
US East (N. Virginia) (us-east-1)	ami-0e380e0a62d368837	ami-06e46a15650294dfa
US West (Oregon) (us-west-2)	ami-0355c210cb3f58aa2	ami-084e8e620163aa50e
Asia Pacific (Hong Kong) (ap-east-1)	ami-03885e0a7493b073f	N/A*
Asia Pacific (Mumbai) (ap-south-1)	ami-01b6a163133c31994	ami-09ad3a49fb13389a0
Asia Pacific (Tokyo) (ap-northeast-1)	ami-0a9b3f8b4b65b402b	ami-0cd09d7293f31df8a
Asia Pacific (Seoul) (ap-northeast-2)	ami-069f6a654a8795f72	ami-006549812c03748cb
Asia Pacific (Singapore) (ap-southeast-1)	ami-03737a1ac334a5767	ami-01be8fddd9b16320c
Asia Pacific (Sydney) (ap-southeast-2)	ami-07580768e8538626f	ami-0a1bf783357dd8492
EU (Frankfurt) (eu-central-1)	ami-0ee5ca4231511cafc	ami-0ae5976723472b6d4
EU (Ireland) (eu-west-1)	ami-0404d23c7e8188740	ami-042f9abf2f96a0097
EU (London) (eu-west-2)	ami-07346d8553f83f9d6	ami-0b87e9246afd42760
EU (Paris) (eu-west-3)	ami-038cb36289174bac4	ami-0d9405868a6e9ee11
EU (Stockholm) (eu-north-1)	ami-03e60b5a990893129	ami-0122b7e2a6736e3c5

Kubernetes version 1.11.9

Region	Amazon EKS-optimized AMI	with GPU support
US East (Ohio) (us-east-2)	ami-03c6648b74285020f	ami-0b87186dda80931ee
US East (N. Virginia) (us-east-1)	ami-0a5f5d5b0f6f58199	ami-07207754196c1a8fc
US West (Oregon) (us-west-2)	ami-057d1c0dcb254a878	ami-052da6a4e0ae156ad
Asia Pacific (Hong Kong) (ap-east-1)	ami-006e7a52666ab7494	N/A*

Region	Amazon EKS-optimized AMI	with GPU support
Asia Pacific (Mumbai) (ap-south-1)	ami-00f1adebe5ab9a431	ami-04645af6384529c5d
Asia Pacific (Tokyo) (ap-northeast-1)	ami-0a0b6606652f9b3b9	ami-0a8f4e1f9bf09a81f
Asia Pacific (Seoul) (ap-northeast-2)	ami-0c84b3f055cda1afb	ami-01db6bb089f6adfcf
Asia Pacific (Singapore) (ap-southeast-1)	ami-05e92412054db3f87	ami-0e001196bd450aa0c
Asia Pacific (Sydney) (ap-southeast-2)	ami-07eb76498b1ba6cd6	ami-0c7132a332aa55aa6
EU (Frankfurt) (eu-central-1)	ami-0234bc9c2b341aa02	ami-05cb4f6e8be8b83f1
EU (Ireland) (eu-west-1)	ami-06902949103360023	ami-02f337476a5c33f1b
EU (London) (eu-west-2)	ami-0db100ad46c7966d2	ami-0aa2208dbb9bb7cc5
EU (Paris) (eu-west-3)	ami-052046d313576d0ba	ami-0f6ea479cb4e7a4d2
EU (Stockholm) (eu-north-1)	ami-02ebf24da505128f9	ami-078c260b9a737fc35

\* GPU instance types are not available in the Asia Pacific (Hong Kong) (ap-east-1) region, so Amazon EKS does not publish the Amazon EKS-optimized AMI with GPU support in that region.

#### Note

The Amazon EKS worker node AMI is based on Amazon Linux 2. You can track security or privacy events for Amazon Linux 2 at the [Amazon Linux Security Center](#) or subscribe to the associated [RSS feed](#). Security and privacy events include an overview of the issue, what packages are affected, and how to update your instances to correct the issue.

- **KeyName:** Enter the name of an Amazon EC2 SSH key pair that you can use to connect using SSH into your worker nodes with after they launch. If you don't already have an Amazon EC2 keypair, you can create one in the AWS Management Console. For more information, see [Amazon EC2 Key Pairs](#) in the *Amazon EC2 User Guide for Linux Instances*.

#### Note

If you do not provide a keypair here, the AWS CloudFormation stack creation fails.

- **BootstrapArguments:** Specify any optional arguments to pass to the worker node bootstrap script, such as extra **kubelet** arguments. For more information, view the bootstrap script usage information at <https://github.com/aws-labs/amazon-eks-ami/blob/master/files/bootstrap.sh>
- **VpcId:** Enter the ID for the VPC that you created in [Create your Amazon EKS Cluster VPC \(p. 9\)](#).
- **Subnets:** Choose the subnets that you created in [Create your Amazon EKS Cluster VPC \(p. 9\)](#). If you created your VPC using the steps described at [Creating a VPC for Your Amazon EKS Cluster \(p. 79\)](#), then specify only the private subnets within the VPC for your worker nodes to launch into.

8. On the **Options** page, you can choose to tag your stack resources. Choose **Next**.
9. On the **Review** page, review your information, acknowledge that the stack might create IAM resources, and then choose **Create**.
10. When your stack has finished creating, select it in the console and choose the **Outputs** tab.

11. Record the **NodeInstanceRole** for the node group that was created. You need this when you configure your Amazon EKS worker nodes.

### To enable worker nodes to join your cluster

1. Download, edit, and apply the AWS authenticator configuration map:
  - a. Download the configuration map with the following command:

```
curl -o aws-auth-cm.yaml https://amazon-eks.s3-us-west-2.amazonaws.com/cloudformation/2019-02-11/aws-auth-cm.yaml
```

- b. Open the file with your favorite text editor. Replace the *<ARN of instance role (not instance profile)>* snippet with the **NodeInstanceRole** value that you recorded in the previous procedure, and save the file.

#### Important

Do not modify any other lines in this file.

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: aws-auth
  namespace: kube-system
data:
  mapRoles: |
    - rolearn: <ARN of instance role (not instance profile)>
      username: system:node:{{EC2PrivateDNSName}}
      groups:
        - system:bootstrappers
        - system:nodes
```

- c. Apply the configuration. This command might take a few minutes to finish.

```
kubectl apply -f aws-auth-cm.yaml
```

#### Note

If you receive the error "aws-iam-authenticator": executable file not found in \$PATH, your **kubectl** isn't configured for Amazon EKS. For more information, see [Installing aws-iam-authenticator \(p. 108\)](#).

If you receive any other authorization or resource type errors, see [Unauthorized or Access Denied \(kubectl\) \(p. 179\)](#) in the troubleshooting section.

2. Watch the status of your nodes and wait for them to reach the Ready status.

```
kubectl get nodes --watch
```

3. (GPU workers only) If you chose a P2 or P3 instance type and the Amazon EKS-optimized AMI with GPU support, you must apply the [NVIDIA device plugin for Kubernetes](#) as a DaemonSet on your cluster with the following command.

```
kubectl apply -f https://raw.githubusercontent.com/NVIDIA/k8s-device-plugin/1.0.0-beta/nvidia-device-plugin.yml
```

## Next Steps

Now that you have a working Amazon EKS cluster with worker nodes, you are ready to start installing Kubernetes add-ons and deploying applications to your cluster. The following documentation topics help you to extend the functionality of your cluster.

- [Launch a Guest Book Application \(p. 124\)](#) — Create a sample guest book application to test your cluster.
- [Tutorial: Deploy the Kubernetes Web UI \(Dashboard\) \(p. 137\)](#) — This tutorial guides you through deploying the [Kubernetes dashboard](#) to your cluster.
- [Using Helm with Amazon EKS \(p. 134\)](#) — The `helm` package manager for Kubernetes helps you install and manage applications on your cluster.
- [Installing the Kubernetes Metrics Server \(p. 128\)](#) — The Kubernetes metrics server is an aggregator of resource usage data in your cluster.
- [Control Plane Metrics with Prometheus \(p. 130\)](#) — This topic helps you deploy Prometheus into your cluster with `helm`.



# Amazon EKS Clusters

An Amazon EKS cluster consists of two primary components:

- The Amazon EKS control plane
- Amazon EKS worker nodes that are registered with the control plane

The Amazon EKS control plane consists of control plane nodes that run the Kubernetes software, such as `etcd` and the Kubernetes API server. The control plane runs in an account managed by AWS, and the Kubernetes API is exposed via the Amazon EKS endpoint associated with your cluster. Each Amazon EKS cluster control plane is single-tenant and unique, and runs on its own set of Amazon EC2 instances.

The cluster control plane is provisioned across multiple Availability Zones and fronted by an Elastic Load Balancing Network Load Balancer. Amazon EKS also provisions elastic network interfaces in your VPC subnets to provide connectivity from the control plane instances to the worker nodes (for example, to support **kubectl exec**, **logs**, and **proxy** data flows).

Amazon EKS worker nodes run in your AWS account and connect to your cluster's control plane via the API server endpoint and a certificate file that is created for your cluster.

## Topics

- [Creating an Amazon EKS Cluster \(p. 20\)](#)
- [Updating an Amazon EKS Cluster Kubernetes Version \(p. 26\)](#)
- [Amazon EKS Cluster Endpoint Access Control \(p. 34\)](#)
- [Amazon EKS Control Plane Logging \(p. 38\)](#)
- [Deleting a Cluster \(p. 41\)](#)
- [Amazon EKS Kubernetes Versions \(p. 43\)](#)
- [Platform Versions \(p. 44\)](#)

## Creating an Amazon EKS Cluster

This topic walks you through creating an Amazon EKS cluster.

If this is your first time creating an Amazon EKS cluster, we recommend that you follow one of our [Getting Started with Amazon EKS \(p. 3\)](#) guides instead. They provide complete end-to-end walkthroughs for creating an Amazon EKS cluster with worker nodes.

### Important

When an Amazon EKS cluster is created, the IAM entity (user or role) that creates the cluster is added to the Kubernetes RBAC authorization table as the administrator (with `system:master` permissions. Initially, only that IAM user can make calls to the Kubernetes API server using **kubectl**. For more information, see [Managing Users or IAM Roles for your Cluster \(p. 115\)](#). If you use the console to create the cluster, you must ensure that the same IAM user credentials are in the AWS SDK credential chain when you are running **kubectl** commands on your cluster. If you install and configure the AWS CLI, you can configure the IAM credentials for your user. If the AWS CLI is configured properly for your user, then `eksctl` and the [AWS IAM Authenticator for Kubernetes](#) can find those credentials as well. For more information, see [Configuring the AWS CLI in the AWS Command Line Interface User Guide](#).

Choose the tab below that corresponds to your desired cluster creation method:

eksctl

## To create your cluster and worker nodes with eksctl

This procedure assumes that you have installed eksctl, and that your eksctl version is at least 0.1.37. You can check your version with the following command:

```
eksctl version
```

For more information on installing or upgrading eksctl, see [Installing or Upgrading eksctl](#) (p. 119).

1. Create your Amazon EKS cluster and worker nodes with the following command. Substitute the red text with your own values.

### Important

Kubernetes version 1.10 is no longer supported on Amazon EKS. You can no longer create new 1.10 clusters, and all existing Amazon EKS clusters running Kubernetes version 1.10 will eventually be automatically updated to the latest available platform version of Kubernetes version 1.11. For more information, see [Amazon EKS Version Deprecation](#) (p. 44).

Please update any 1.10 clusters to version 1.11 or higher in order to avoid service interruption. For more information, see [Updating an Amazon EKS Cluster Kubernetes Version](#) (p. 26).

```
eksctl create cluster \
--name prod \
--version 1.13 \
--nodegroup-name standard-workers \
--node-type t3.medium \
--nodes 3 \
--nodes-min 1 \
--nodes-max 4 \
--node-ami auto
```

### Note

For more information on the available options for **eksctl create cluster**, see the project [README on GitHub](#) or view the help page with the following command.

```
eksctl create cluster --help
```

Output:

```
[#] using region us-west-2
[#] setting availability zones to [us-west-2b us-west-2c us-west-2d]
[#] subnets for us-west-2b - public:192.168.0.0/19 private:192.168.96.0/19
[#] subnets for us-west-2c - public:192.168.32.0/19 private:192.168.128.0/19
[#] subnets for us-west-2d - public:192.168.64.0/19 private:192.168.160.0/19
[#] nodegroup "standard-workers" will use
    "ami-0923e4b35a30a5f53" [AmazonLinux2/1.12]
[#] creating EKS cluster "prod" in "us-west-2" region
[#] will create 2 separate CloudFormation stacks for cluster itself and the
    initial nodegroup
[#] if you encounter any issues, check CloudFormation console or try 'eksctl utils
    describe-stacks --region=us-west-2 --name=prod'
[#] building cluster stack "eksctl-prod-cluster"
[#] creating nodegroup stack "eksctl-prod-nodegroup-standard-workers"
[#] all EKS cluster resource for "prod" had been created
[#] saved kubeconfig as "/Users/ericn/.kube/config"
```

```
[#] adding role "arn:aws:iam::111122223333:role/eksctl-prod-nodegroup-standard-wo-NodeInstanceRole-IJP4S12W3020" to auth ConfigMap
[#] nodegroup "standard-workers" has 0 node(s)
[#] waiting for at least 1 node(s) to become ready in "standard-workers"
[#] nodegroup "standard-workers" has 2 node(s)
[#] node "ip-192-168-22-17.us-west-2.compute.internal" is not ready
[#] node "ip-192-168-32-184.us-west-2.compute.internal" is ready
[#] kubectl command should work with "/Users/ericn/.kube/config", try 'kubectl get nodes'
[#] EKS cluster "prod" in "us-west-2" region is ready
```

- Cluster provisioning usually takes between 10 and 15 minutes. When your cluster is ready, test that your `kubectl` configuration is correct.

```
kubectl get svc
```

### Note

If you receive the error `"aws-iam-authenticator": executable file not found in $PATH`, your `kubectl` isn't configured for Amazon EKS. For more information, see [Installing aws-iam-authenticator \(p. 108\)](#).

If you receive any other authorization or resource type errors, see [Unauthorized or Access Denied \(kubectl\) \(p. 179\)](#) in the troubleshooting section.

Output:

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
svc/kubernetes	ClusterIP	10.100.0.1	<none>	443/TCP	1m

- (GPU workers only) If you chose a P2 or P3 instance type and the Amazon EKS-optimized AMI with GPU support, you must apply the [NVIDIA device plugin for Kubernetes](#) as a DaemonSet on your cluster with the following command.

```
kubectl apply -f https://raw.githubusercontent.com/NVIDIA/k8s-device-plugin/1.0.0-beta/nvidia-device-plugin.yml
```

## AWS Management Console

### To create your cluster with the console

This procedure has the following prerequisites:

- You have created a VPC and a dedicated security group that meet the requirements for an Amazon EKS cluster. For more information, see [Cluster VPC Considerations \(p. 81\)](#) and [Cluster Security Group Considerations \(p. 83\)](#). The [Getting Started with the AWS Management Console \(p. 8\)](#) guide creates a VPC that meets the requirements, or you can also follow [Creating a VPC for Your Amazon EKS Cluster \(p. 79\)](#) to create one.
- You have created an Amazon EKS service role to apply to your cluster. The [Getting Started with Amazon EKS \(p. 3\)](#) guide creates a service role for you, or you can also follow [Amazon EKS IAM Roles \(p. 162\)](#) to create one manually.

- Open the Amazon EKS console at <https://console.aws.amazon.com/eks/home#/clusters>.
- Choose **Create cluster**.

### Note

If your IAM user doesn't have administrative privileges, you must explicitly add permissions for that user to call the Amazon EKS API operations. For more information, see [Amazon EKS Identity-Based Policy Examples \(p. 163\)](#).

3. On the **Create cluster** page, fill in the following fields and then choose **Create**:

- **Cluster name** – A unique name for your cluster.
- **Kubernetes version** – The version of Kubernetes to use for your cluster. Unless you require a specific Kubernetes version for your application, we recommend that you use the latest version available in Amazon EKS.

**Important**

Kubernetes version 1.10 is no longer supported on Amazon EKS. You can no longer create new 1.10 clusters, and all existing Amazon EKS clusters running Kubernetes version 1.10 will eventually be automatically updated to the latest available platform version of Kubernetes version 1.11. For more information, see [Amazon EKS Version Deprecation](#) (p. 44).

Please update any 1.10 clusters to version 1.11 or higher in order to avoid service interruption. For more information, see [Updating an Amazon EKS Cluster Kubernetes Version](#) (p. 26).

- **Role name** – Choose the Amazon EKS service role to allow Amazon EKS and the Kubernetes control plane to manage AWS resources on your behalf. For more information, see [Amazon EKS IAM Roles](#) (p. 162).
- **VPC** – The VPC to use for your cluster.
- **Subnets** – The subnets within the preceding VPC to use for your cluster. By default, the available subnets in the VPC are preselected. Specify all subnets that will host resources for your cluster (such as private subnets for worker nodes and public subnets for load balancers). Your subnets must meet the requirements for an Amazon EKS cluster. For more information, see [Cluster VPC Considerations](#) (p. 81).
- **Security Groups** – Specify one or more (up to a limit of five) security groups within the preceding VPC to apply to the cross-account elastic network interfaces for your cluster. Your cluster and worker node security groups must meet the requirements for an Amazon EKS cluster. For more information, see [Cluster Security Group Considerations](#) (p. 83).

**Important**

The worker node AWS CloudFormation template modifies the security group that you specify here, so **Amazon EKS strongly recommends that you use a dedicated security group for each cluster control plane (one per cluster)**. If this security group is shared with other resources, you might block or disrupt connections to those resources.

- **Endpoint private access** – Choose whether to enable or disable private access for your cluster's Kubernetes API server endpoint. If you enable private access, Kubernetes API requests that originate from within your cluster's VPC use the private VPC endpoint. For more information, see [Amazon EKS Cluster Endpoint Access Control](#) (p. 34).
- **Endpoint public access** – Choose whether to enable or disable public access for your cluster's Kubernetes API server endpoint. If you disable public access, your cluster's Kubernetes API server can receive only requests from within the cluster VPC. For more information, see [Amazon EKS Cluster Endpoint Access Control](#) (p. 34).
- **Logging** – For each individual log type, choose whether the log type should be **Enabled** or **Disabled**. By default, each log type is **Disabled**. For more information, see [Amazon EKS Control Plane Logging](#) (p. 38).

**Note**

You might receive an error that one of the Availability Zones in your request doesn't have sufficient capacity to create an Amazon EKS cluster. If this happens, the error output contains the Availability Zones that can support a new cluster. Retry creating your cluster with at least two subnets that are located in the supported Availability Zones for your account. For more information, see [Insufficient Capacity](#) (p. 179).

4. On the **Clusters** page, choose the name of your new cluster to view the cluster information.

5. The **Status** field shows **CREATING** until the cluster provisioning process completes. When your cluster provisioning is complete (usually between 10 and 15 minutes), note the **API server endpoint** and **Certificate authority** values. These are used in your **kubectl** configuration.
6. Now that you have created your cluster, follow the procedures in [Installing aws-iam-authenticator \(p. 108\)](#) and [Create a kubeconfig for Amazon EKS \(p. 111\)](#) to enable communication with your new cluster.
7. After you enable communication, follow the procedures in [Launching Amazon EKS Worker Nodes \(p. 56\)](#) to add worker nodes to your cluster to support your workloads.

## AWS CLI

### To create your cluster with the AWS CLI

This procedure has the following prerequisites:

- You have created a VPC and a dedicated security group that meets the requirements for an Amazon EKS cluster. For more information, see [Cluster VPC Considerations \(p. 81\)](#) and [Cluster Security Group Considerations \(p. 83\)](#). The [Getting Started with the AWS Management Console \(p. 8\)](#) guide creates a VPC that meets the requirements, or you can also follow [Creating a VPC for Your Amazon EKS Cluster \(p. 79\)](#) to create one.
  - You have created an Amazon EKS service role to apply to your cluster. The [Getting Started with Amazon EKS \(p. 3\)](#) guide creates a service role for you, or you can also follow [Amazon EKS IAM Roles \(p. 162\)](#) to create one manually.
1. Create your cluster with the following command. Substitute your cluster name, the Amazon Resource Name (ARN) of your Amazon EKS service role that you created in [Create your Amazon EKS Service Role \(p. 8\)](#), and the subnet and security group IDs for the VPC that you created in [Create your Amazon EKS Cluster VPC \(p. 9\)](#).

#### Important

Kubernetes version 1.10 is no longer supported on Amazon EKS. You can no longer create new 1.10 clusters, and all existing Amazon EKS clusters running Kubernetes version 1.10 will eventually be automatically updated to the latest available platform version of Kubernetes version 1.11. For more information, see [Amazon EKS Version Deprecation \(p. 44\)](#).

Please update any 1.10 clusters to version 1.11 or higher in order to avoid service interruption. For more information, see [Updating an Amazon EKS Cluster Kubernetes Version \(p. 26\)](#).

```
aws eks --region region create-cluster --name devel --kubernetes-version 1.13 \
--role-arn arn:aws:iam::111122223333:role/eks-service-role-
AWSServiceRoleForAmazonEKS-EXAMPLEBKZRQR \
--resources-vpc-config subnetIds=subnet-
a9189fe2,subnet-50432629,securityGroupIds=sg-f5c54184
```

#### Important

If you receive a syntax error similar to the following, you might be using a preview version of the AWS CLI for Amazon EKS. The syntax for many Amazon EKS commands has changed since the public service launch. Update your AWS CLI version to the latest available and delete the custom service model directory at `~/.aws/models/eks`.

```
aws: error: argument --cluster-name is required
```

### Note

If your IAM user doesn't have administrative privileges, you must explicitly add permissions for that user to call the Amazon EKS API operations. For more information, see [Amazon EKS Identity-Based Policy Examples \(p. 163\)](#).

Output:

```
{
  "cluster": {
    "name": "devel",
    "arn": "arn:aws:eks:us-west-2:111122223333:cluster/devel",
    "createdAt": 1527785885.159,
    "version": "1.13",
    "roleArn": "arn:aws:iam::111122223333:role/eks-service-role-AWSServiceRoleForAmazonEKS-AFNL4H8HB71F",
    "resourcesVpcConfig": {
      "subnetIds": [
        "subnet-a9189fe2",
        "subnet-50432629"
      ],
      "securityGroupIds": [
        "sg-f5c54184"
      ],
      "vpcId": "vpc-a54041dc",
      "endpointPublicAccess": true,
      "endpointPrivateAccess": false
    },
    "status": "CREATING",
    "certificateAuthority": {}
  }
}
```

### Note

You might receive an error that one of the Availability Zones in your request doesn't have sufficient capacity to create an Amazon EKS cluster. If this happens, the error output contains the Availability Zones that can support a new cluster. Retry creating your cluster with at least two subnets that are located in the supported Availability Zones for your account. For more information, see [Insufficient Capacity \(p. 179\)](#).

2. Cluster provisioning usually takes between 10 and 15 minutes. You can query the status of your cluster with the following command. When your cluster status is `ACTIVE`, you can proceed.

```
aws eks --region region describe-cluster --name devel --query cluster.status
```

3. When your cluster provisioning is complete, retrieve the endpoint and `certificateAuthority.data` values with the following commands. You must add these values to your **kubectl** configuration so that you can communicate with your cluster.

- a. Retrieve the endpoint.

```
aws eks --region region describe-cluster --name devel --query cluster.endpoint
--output text
```

- b. Retrieve the `certificateAuthority.data`.

```
aws eks --region region describe-cluster --name devel --query
cluster.certificateAuthority.data --output text
```

4. Now that you have created your cluster, follow the procedures in [Installing aws-iam-authenticator \(p. 108\)](#) and [Create a kubeconfig for Amazon EKS \(p. 111\)](#) to enable communication with your new cluster.
5. After you enable communication, follow the procedures in [Launching Amazon EKS Worker Nodes \(p. 56\)](#) to add worker nodes to your cluster to support your workloads.

## Updating an Amazon EKS Cluster Kubernetes Version

When a new Kubernetes version is available in Amazon EKS, you can update your cluster to the latest version. New Kubernetes versions introduce significant changes, so we recommend that you test the behavior of your applications against a new Kubernetes version before performing the update on your production clusters. You can achieve this by building a continuous integration workflow to test your application behavior end-to-end before moving to a new Kubernetes version.

The update process consists of Amazon EKS launching new API server nodes with the updated Kubernetes version to replace the existing ones. Amazon EKS performs standard infrastructure and readiness health checks for network traffic on these new nodes to verify that they are working as expected. If any of these checks fail, Amazon EKS reverts the infrastructure deployment, and your cluster remains on the prior Kubernetes version. Running applications are not affected, and your cluster is never left in a non-deterministic or unrecoverable state. Amazon EKS regularly backs up all managed clusters, and mechanisms exist to recover clusters if necessary. We are constantly evaluating and improving our Kubernetes infrastructure management processes.

In order to upgrade the cluster, Amazon EKS requires 2-3 free IP addresses from the subnets which were provided when you created the cluster. If these subnets do not have available IP addresses, then the upgrade can fail. Additionally, if any of the subnets or security groups that were provided during cluster creation have been deleted, the cluster upgrade process can fail.

### Note

Although Amazon EKS runs a highly available control plane, you might experience minor service interruptions during an update. For example, if you attempt to connect to an API server just before or just after it's terminated and replaced by a new API server running the new version of Kubernetes, you might experience API call errors or connectivity issues. If this happens, retry your API operations until they succeed.

Amazon EKS does not modify any of your Kubernetes add-ons when you update a cluster. After updating your cluster, we recommend that you update your add-ons to the versions listed in the following table for the new Kubernetes version that you're updating to (steps to accomplish this are included in the update procedures).

Kubernetes Version	1.11	1.12	1.13
Amazon VPC CNI plug-in	We recommend the latest available CNI version (1.5.2)		
DNS	CoreDNS 1.1.3	CoreDNS 1.2.2	CoreDNS 1.2.6
KubeProxy	1.11.8	1.12.6	1.13.7

### Important

Kubernetes version 1.10 is no longer supported on Amazon EKS. You can no longer create new 1.10 clusters, and all existing Amazon EKS clusters running Kubernetes version 1.10 will eventually be automatically updated to the latest available platform version of Kubernetes version 1.11. For more information, see [Amazon EKS Version Deprecation \(p. 44\)](#).

Please update any 1.10 clusters to version 1.11 or higher in order to avoid service interruption. For more information, see [Updating an Amazon EKS Cluster Kubernetes Version \(p. 26\)](#).

If you're using additional add-ons for your cluster that aren't listed in the previous table, update them to the latest compatible versions after updating your cluster.

Choose the tab below that corresponds to your desired cluster update method:

eksctl

### To update an existing cluster with eksctl

This procedure assumes that you have installed eksctl, and that your eksctl version is at least 0.1.37. You can check your version with the following command:

```
eksctl version
```

For more information on installing or upgrading eksctl, see [Installing or Upgrading eksctl \(p. 119\)](#).

#### Note

This procedure only works for clusters that were created with eksctl.

1. Compare the Kubernetes version of your cluster control plane to the Kubernetes version of your worker nodes.

- Get the Kubernetes version of your cluster control plane with the following command.

```
kubectl version --short
```

- Get the Kubernetes version of your worker nodes with the following command.

```
kubectl get nodes
```

If your worker nodes are more than one Kubernetes minor version older than your control plane, then you must upgrade your worker nodes to a newer Kubernetes minor version before you update your cluster's Kubernetes version. For more information, see [Kubernetes version and version skew support policy](#) in the Kubernetes documentation.

We recommend that you update your worker nodes to your cluster's current pre-update Kubernetes minor version prior to your cluster update. Your worker nodes must not run a newer Kubernetes version than your control plane. For example, if your control plane is running version 1.12 and your workers are running version 1.10, update your worker nodes to version 1.11 or 1.12 (recommended) before you update your cluster's Kubernetes version to 1.13. For more information, see [Worker Node Updates \(p. 63\)](#).

2. Update your Amazon EKS cluster Kubernetes version with the following command, replacing the red text with your cluster name:

```
eksctl update cluster --name dev --approve
```

This process takes several minutes to complete.

3. Patch the kube-proxy daemonset for your updated Kubernetes version with the following command, replacing the red text with your cluster name:

```
eksctl utils update-kube-proxy --name dev --approve
```



4. Check your cluster's DNS provider. Clusters that were created with Kubernetes version 1.10 shipped with `kube-dns` as the default DNS and service discovery provider. If you have updated a 1.10 cluster to a newer version and you want to use CoreDNS for DNS and service discovery, you must install CoreDNS and remove `kube-dns`.

To check if your cluster is already running CoreDNS, use the following command.

```
kubectl get pod -n kube-system -l k8s-app=kube-dns
```

If the output shows `coredns` in the pod names, you're already running CoreDNS in your cluster. If not, run the following command to install `coredns`, replacing the red text with your cluster name:

```
eksctl utils install-coredns --name dev --approve
```

5. Check the current version of your cluster's `coredns` deployment.

```
kubectl describe deployment coredns --namespace kube-system | grep Image | cut -d "/" -f 3
```

The recommended `coredns` versions for their corresponding Kubernetes versions are as follows:

- **Kubernetes 1.13:** 1.2.6
- **Kubernetes 1.12:** 1.2.2
- **Kubernetes 1.11:** 1.1.3

If your current `coredns` version doesn't match the recommendation for your cluster version, update the `coredns` deployment to use the recommended image.

```
kubectl set image --namespace kube-system deployment.apps/coredns \  
coredns=602401143452.dkr.ecr.us-west-2.amazonaws.com/eks/coredns:v1.2.2
```

6. Check the version of your cluster's Amazon VPC CNI Plugin for Kubernetes. Use the following command to print your cluster's CNI version.

```
kubectl describe daemonset aws-node --namespace kube-system | grep Image | cut -d "/" -f 2
```

Output:

```
amazon-k8s-cni:1.4.1
```

If your CNI version is earlier than 1.5.2, use the following command to upgrade your CNI version to the latest version:

- For Kubernetes 1.10 clusters:

```
kubectl apply -f https://raw.githubusercontent.com/aws/amazon-vpc-cni-k8s/  
release-1.5/config/v1.5/aws-k8s-cni-1.10.yaml
```

- For all other Kubernetes versions:

```
kubectl apply -f https://raw.githubusercontent.com/aws/amazon-vpc-cni-k8s/release-1.5/config/v1.5/aws-k8s-cni.yaml
```

7. (Clusters with GPU workers only) If your cluster has worker node groups with GPU support (for example, `p3.2xlarge`), you must update the [NVIDIA device plugin for Kubernetes](#) DaemonSet on your cluster with the following command.

```
kubectl apply -f https://raw.githubusercontent.com/NVIDIA/k8s-device-plugin/1.0.0-beta/nvidia-device-plugin.yml
```

8. After your cluster update is complete, update your worker nodes to the same Kubernetes version of your updated cluster. For more information, see [Worker Node Updates](#) (p. 63).

## AWS Management Console

### To update an existing cluster with the console

1. Compare the Kubernetes version of your cluster control plane to the Kubernetes version of your worker nodes.
  - Get the Kubernetes version of your cluster control plane with the following command.

```
kubectl version --short
```

- Get the Kubernetes version of your worker nodes with the following command.

```
kubectl get nodes
```

If your worker nodes are more than one Kubernetes minor version older than your control plane, then you must upgrade your worker nodes to a newer Kubernetes minor version before you update your cluster's Kubernetes version. For more information, see [Kubernetes version and version skew support policy](#) in the Kubernetes documentation.

We recommend that you update your worker nodes to your cluster's current pre-update Kubernetes minor version prior to your cluster update. Your worker nodes must not run a newer Kubernetes version than your control plane. For example, if your control plane is running version 1.12 and your workers are running version 1.10, update your worker nodes to version 1.11 or 1.12 (recommended) before you update your cluster's Kubernetes version to 1.13. For more information, see [Worker Node Updates](#) (p. 63).

2. Open the Amazon EKS console at <https://console.aws.amazon.com/eks/home#/clusters>.
3. Choose the name of the cluster to update and choose **Update cluster version**.
4. For **Kubernetes version**, select the version to update your cluster to and choose **Update**.

#### Important

Kubernetes version 1.10 is no longer supported on Amazon EKS. You can no longer create new 1.10 clusters, and all existing Amazon EKS clusters running Kubernetes version 1.10 will eventually be automatically updated to the latest available platform version of Kubernetes version 1.11. For more information, see [Amazon EKS Version Deprecation](#) (p. 44).

Please update any 1.10 clusters to version 1.11 or higher in order to avoid service interruption. For more information, see [Updating an Amazon EKS Cluster Kubernetes Version](#) (p. 26).

### Important

Because Amazon EKS runs a highly available control plane, you must update only one minor version at a time. See [Kubernetes Version and Version Skew Support Policy](#) for the rationale behind this requirement. Therefore, if your current version is 1.11 and you want to upgrade to 1.13, you must first upgrade your cluster to 1.12 and then upgrade it from 1.12 to 1.13. If you try to update directly from 1.11 to 1.13, the update version command throws an error.

- For **Cluster name**, type the name of your cluster and choose **Confirm**.

### Note

The cluster update should finish in a few minutes.

- After the update is complete, find the current Kubernetes version for your cluster.

```
kubectl version --short
Client Version: v1.13.4
Server Version: v1.12.6-eks-d69f1b
```

- Patch the kube-proxy daemonset to use the image that corresponds to your current cluster Kubernetes version (in this example, 1.12.6).

```
kubectl patch daemonset kube-proxy \
-n kube-system \
-p '{"spec": {"template": {"spec": {"containers": [{"image":
"602401143452.dkr.ecr.us-west-2.amazonaws.com/eks/kube-
proxy:v1.12.6", "name": "kube-proxy"}]}}}}'
```

- Check your cluster's DNS provider. Clusters that were created with Kubernetes version 1.10 shipped with kube-dns as the default DNS and service discovery provider. If you have updated a 1.10 cluster to a newer version and you want to use CoreDNS for DNS and service discovery, you must install CoreDNS and remove kube-dns.

To check if your cluster is already running CoreDNS, use the following command.

```
kubectl get pod -n kube-system -l k8s-app=kube-dns
```

If the output shows `coredns` in the pod names, you're already running CoreDNS in your cluster. If not, see [Installing CoreDNS \(p. 88\)](#) to install CoreDNS on your cluster and then return here.

- Check the current version of your cluster's `coredns` deployment.

```
kubectl describe deployment coredns --namespace kube-system | grep Image | cut -d
"/" -f 3
```

The recommended `coredns` versions for their corresponding Kubernetes versions are as follows:

- **Kubernetes 1.13:** 1.2.6
- **Kubernetes 1.12:** 1.2.2
- **Kubernetes 1.11:** 1.1.3

If your current `coredns` version doesn't match the recommendation for your cluster version, update the `coredns` deployment to use the recommended image.

```
kubectl set image --namespace kube-system deployment.apps/coredns \
coredns=602401143452.dkr.ecr.us-west-2.amazonaws.com/eks/coredns:v1.2.2
```

10. Check the version of your cluster's Amazon VPC CNI Plugin for Kubernetes. Use the following command to print your cluster's CNI version.

```
kubectl describe daemonset aws-node --namespace kube-system | grep Image | cut -d "/" -f 2
```

Output:

```
amazon-k8s-cni:1.4.1
```

If your CNI version is earlier than 1.5.2, use the following command to upgrade your CNI version to the latest version.

- For Kubernetes 1.10 clusters:

```
kubectl apply -f https://raw.githubusercontent.com/aws/amazon-vpc-cni-k8s/release-1.5/config/v1.5/aws-k8s-cni-1.10.yaml
```

- For all other Kubernetes versions:

```
kubectl apply -f https://raw.githubusercontent.com/aws/amazon-vpc-cni-k8s/release-1.5/config/v1.5/aws-k8s-cni.yaml
```

11. (Clusters with GPU workers only) If your cluster has worker node groups with GPU support (for example, p3.2xlarge), you must update the [NVIDIA device plugin for Kubernetes](#) DaemonSet on your cluster with the following command.

```
kubectl apply -f https://raw.githubusercontent.com/NVIDIA/k8s-device-plugin/1.0.0-beta/nvidia-device-plugin.yaml
```

12. After your cluster update is complete, update your worker nodes to the same Kubernetes version of your updated cluster. For more information, see [Worker Node Updates](#) (p. 63).

## AWS CLI

### To update an existing cluster with the AWS CLI

1. Compare the Kubernetes version of your cluster control plane to the Kubernetes version of your worker nodes.

- Get the Kubernetes version of your cluster control plane with the following command.

```
kubectl version --short
```

- Get the Kubernetes version of your worker nodes with the following command.

```
kubectl get nodes
```

If your worker nodes are more than one Kubernetes minor version older than your control plane, then you must upgrade your worker nodes to a newer Kubernetes minor version before you update your cluster's Kubernetes version. For more information, see [Kubernetes version and version skew support policy](#) in the Kubernetes documentation.

We recommend that you update your worker nodes to your cluster's current pre-update Kubernetes minor version prior to your cluster update. Your worker nodes must not run a newer

Kubernetes version than your control plane. For example, if your control plane is running version 1.12 and your workers are running version 1.10, update your worker nodes to version 1.11 or 1.12 (recommended) before you update your cluster's Kubernetes version to 1.13. For more information, see [Worker Node Updates \(p. 63\)](#).

2. Update your cluster with the following AWS CLI command. Substitute your cluster name and desired Kubernetes minor version.

### Important

Kubernetes version 1.10 is no longer supported on Amazon EKS. You can no longer create new 1.10 clusters, and all existing Amazon EKS clusters running Kubernetes version 1.10 will eventually be automatically updated to the latest available platform version of Kubernetes version 1.11. For more information, see [Amazon EKS Version Deprecation \(p. 44\)](#).

Please update any 1.10 clusters to version 1.11 or higher in order to avoid service interruption. For more information, see [Updating an Amazon EKS Cluster Kubernetes Version \(p. 26\)](#).

### Important

Because Amazon EKS runs a highly available control plane, you must update only one minor version at a time. See [Kubernetes Version and Version Skew Support Policy](#) for the rationale behind this requirement. Therefore, if your current version is 1.11 and you want to upgrade to 1.13, you must first upgrade your cluster to 1.12 and then upgrade it from 1.12 to 1.13. If you try to update directly from 1.11 to 1.13, the update version command throws an error.

```
aws eks --region region update-cluster-version --name prod --kubernetes-  
version 1.13
```

Output:

```
{  
  "update": {  
    "id": "b5f0ba18-9a87-4450-b5a0-825e6e84496f",  
    "status": "InProgress",  
    "type": "VersionUpdate",  
    "params": [  
      {  
        "type": "Version",  
        "value": "1.13"  
      },  
      {  
        "type": "PlatformVersion",  
        "value": "eks.1"  
      }  
    ],  
    "createdAt": 1544051347.305,  
    "errors": []  
  }  
}
```

3. Monitor the status of your cluster update with the following command, using the cluster name and update ID that the previous command returned. Your update is complete when the status appears as Successful.

### Note

The cluster update should finish in a few minutes.

```
aws eks --region region describe-update --name prod --update-id b5f0ba18-9a87-4450-  
b5a0-825e6e84496f
```

Output:

```
{
  "update": {
    "id": "b5f0ba18-9a87-4450-b5a0-825e6e84496f",
    "status": "Successful",
    "type": "VersionUpdate",
    "params": [
      {
        "type": "Version",
        "value": "1.13"
      },
      {
        "type": "PlatformVersion",
        "value": "eks.1"
      }
    ],
    "createdAt": 1544051347.305,
    "errors": []
  }
}
```

- After the update is complete, find the current Kubernetes version for your cluster.

```
kubectl version --short
Client Version: v1.13.4
Server Version: v1.13.7-eks-nnnnnn
```

- Patch the kube-proxy daemonset to use the image that corresponds to your current cluster Kubernetes version (in this example, 1.13.7).

```
kubectl patch daemonset kube-proxy \
-n kube-system \
-p '{"spec": {"template": {"spec": {"containers": [{"image":
  "602401143452.dkr.ecr.us-west-2.amazonaws.com/eks/kube-
proxy:v1.13.7", "name": "kube-proxy"}]}}}}'
```

- Check your cluster's DNS provider. Clusters that were created with Kubernetes version 1.10 shipped with kube-dns as the default DNS and service discovery provider. If you have updated a 1.10 cluster to a newer version and you want to use CoreDNS for DNS and service discovery, you must install CoreDNS and remove kube-dns.

To check if your cluster is already running CoreDNS, use the following command.

```
kubectl get pod -n kube-system -l k8s-app=kube-dns
```

If the output shows `coredns` in the pod names, you're already running CoreDNS in your cluster. If not, see [Installing CoreDNS \(p. 88\)](#) to install CoreDNS on your cluster and then return here.

- Check the current version of your cluster's `coredns` deployment.

```
kubectl describe deployment coredns --namespace kube-system | grep Image | cut -d
"/" -f 3
```

The recommended `coredns` versions for their corresponding Kubernetes versions are as follows:

- **Kubernetes 1.13:** 1.2.6
- **Kubernetes 1.12:** 1.2.2

- **Kubernetes 1.11:** 1.1.3

If your current `coredns` version doesn't match the recommendation for your cluster version, update the `coredns` deployment to use the recommended image.

```
kubectl set image --namespace kube-system deployment.apps/coredns \
coredns=602401143452.dkr.ecr.us-west-2.amazonaws.com/eks/coredns:v1.2.6
```

8. Check the version of your cluster's Amazon VPC CNI Plugin for Kubernetes. Use the following command to print your cluster's CNI version.

```
kubectl describe daemonset aws-node --namespace kube-system | grep Image | cut -d
"/" -f 2
```

Output:

```
amazon-k8s-cni:1.4.1
```

If your CNI version is earlier than 1.5.2, use the following command to upgrade your CNI version to the latest version.

- For Kubernetes 1.10 clusters:

```
kubectl apply -f https://raw.githubusercontent.com/aws/amazon-vpc-cni-k8s/
release-1.5/config/v1.5/aws-k8s-cni-1.10.yaml
```

- For all other Kubernetes versions:

```
kubectl apply -f https://raw.githubusercontent.com/aws/amazon-vpc-cni-k8s/
release-1.5/config/v1.5/aws-k8s-cni.yaml
```

9. (Clusters with GPU workers only) If your cluster has worker node groups with GPU support (for example, `p3.2xlarge`), you must update the [NVIDIA device plugin for Kubernetes](#) DaemonSet on your cluster with the following command.

```
kubectl apply -f https://raw.githubusercontent.com/NVIDIA/k8s-device-plugin/1.0.0-
beta/nvidia-device-plugin.yml
```

10. After your cluster update is complete, update your worker nodes to the same Kubernetes version of your updated cluster. For more information, see [Worker Node Updates](#) (p. 63).

## Amazon EKS Cluster Endpoint Access Control

This topic helps you to enable private access for your Amazon EKS cluster's Kubernetes API server endpoint and completely disable public access so that it's not accessible from the internet.

When you create a new cluster, Amazon EKS creates an endpoint for the managed Kubernetes API server that you use to communicate with your cluster (using Kubernetes management tools such as `kubectl`). By default, this API server endpoint is public to the internet, and access to the API server is secured using a combination of AWS Identity and Access Management (IAM) and native Kubernetes [Role Based Access Control](#) (RBAC).

You can enable private access to the Kubernetes API server so that all communication between your worker nodes and the API server stays within your VPC. You can also completely disable public access to your API server so that it's not accessible from the internet.

**Note**

Because this endpoint is for the Kubernetes API server and not a traditional AWS PrivateLink endpoint for communicating with an AWS API, it doesn't appear as an endpoint in the Amazon VPC console.

When you enable endpoint private access for your cluster, Amazon EKS creates a Route 53 private hosted zone on your behalf and associates it with your cluster's VPC. This private hosted zone is managed by Amazon EKS, and it doesn't appear in your account's Route 53 resources. In order for the private hosted zone to properly route traffic to your API server, your VPC must have `enableDnsHostnames` and `enableDnsSupport` set to `true`, and the DHCP options set for your VPC must include `AmazonProvidedDNS` in its domain name servers list. For more information, see [Updating DNS Support for Your VPC](#) in the *Amazon VPC User Guide*.

**Note**

In addition to standard Amazon EKS permissions, your IAM user or role must have `route53:AssociateVPCWithHostedZone` permissions to enable the cluster's endpoint private access.

You can define your API server endpoint access requirements when you create a new cluster, and you can update the API server endpoint access for a cluster at any time.

## Modifying Cluster Endpoint Access

Use the procedures in this section to modify the endpoint access for an existing cluster. The following table shows the supported API server endpoint access combinations and their associated behavior.

**API server endpoint access options**

Endpoint Public Access	Endpoint Private Access	Behavior
Enabled	Disabled	<ul style="list-style-type: none"><li>This is the default behavior for new Amazon EKS clusters.</li><li>Kubernetes API requests that originate from within your cluster's VPC (such as worker node to control plane communication) leave the VPC but not Amazon's network.</li><li>Your cluster API server is accessible from the internet.</li></ul>
Enabled	Enabled	<ul style="list-style-type: none"><li>Kubernetes API requests within your cluster's VPC (such as worker node to control plane communication) use the private VPC endpoint.</li><li>Your cluster API server is accessible from the internet.</li></ul>
Disabled	Enabled	<ul style="list-style-type: none"><li>All traffic to your cluster API server must come from within your cluster's VPC.</li><li>There is no public access to your API server from the internet. Any <code>kubectl</code> commands must come from within the VPC as well. For</li></ul>



Endpoint Public Access	Endpoint Private Access	Behavior
		connectivity options, see <a href="#">Accessing the API Server from within the VPC (p. 37)</a> .

### To modify your cluster API server endpoint access with the console

1. Open the Amazon EKS console at <https://console.aws.amazon.com/eks/home#/clusters>.
2. Choose the name of the cluster to display your cluster information.
3. Under **Networking**, choose **Update**.
4. For **Endpoint private access**, choose whether to enable or disable private access for your cluster's Kubernetes API server endpoint. If you enable private access, Kubernetes API requests that originate from within your cluster's VPC use the private VPC endpoint. You must enable private access to disable public access.
5. For **Endpoint public access**, choose whether to enable or disable public access for your cluster's Kubernetes API server endpoint. If you disable public access, your cluster's Kubernetes API server can only receive requests from within the cluster VPC.
6. Choose **Update** to finish.

### To modify your cluster API server endpoint access with the AWS CLI

1. Update your cluster API server endpoint access with the following AWS CLI command. Substitute your cluster name and desired endpoint access values.

#### Note

The following command enables private access for the API server endpoint and completely disables public access.

```
aws eks --region region update-cluster-config --name dev --resources-vpc-config  
endpointPublicAccess=false,endpointPrivateAccess=true
```

Output:

```
{  
  "update": {  
    "id": "70e7ad6d-8de4-4ed3-9040-1ced27f8c332",  
    "status": "InProgress",  
    "type": "EndpointAccessUpdate",  
    "params": [  
      {  
        "type": "EndpointPublicAccess",  
        "value": "false"  
      },  
      {  
        "type": "EndpointPrivateAccess",  
        "value": "true"  
      }  
    ],  
    "createdAt": 1551817408.563,  
    "errors": []  
  }  
}
```

2. Monitor the status of your endpoint access update with the following command, using the cluster name and update ID that was returned by the previous command. Your update is complete when the status is shown as **Successful**.

```
aws eks --region region describe-update --name dev --update-id 70e7ad6d-8de4-4ed3-9040-1ced27f8c332
```

Output:

```
{
  "update": {
    "id": "70e7ad6d-8de4-4ed3-9040-1ced27f8c332",
    "status": "Successful",
    "type": "EndpointAccessUpdate",
    "params": [
      {
        "type": "EndpointPublicAccess",
        "value": "false"
      },
      {
        "type": "EndpointPrivateAccess",
        "value": "true"
      }
    ],
    "createdAt": 1551817408.563,
    "errors": []
  }
}
```

## Accessing the API Server from within the VPC

If you have disabled public access for your cluster's Kubernetes API server endpoint, you can only access the API server from within your VPC. Here are a few possible ways to access the Kubernetes API server endpoint from within the VPC:

### Note

You must ensure that your Amazon EKS control plane security group contains rules to allow ingress traffic for the following solutions. For example, if you are using an Amazon EC2 bastion host or AWS Cloud9 IDE to communicate with your cluster, then your control plane security group must allow ingress traffic on port 443 from your bastion host or IDE security group. For more information, see [Cluster Security Group Considerations \(p. 83\)](#).

The DNS name of the Kubernetes cluster endpoint is only resolvable from the worker node VPC, for the following reasons:

- The Route 53 private hosted zone that is created for the endpoint is only associated with the worker node VPC.
- The private hosted zone is created in a separate AWS managed account and cannot be altered.

If you want to reach the cluster endpoint from a peered VPC or your on premises network through AWS Direct Connect or a transit gateway, you must enable DNS resolution for the cluster endpoint to work outside of the worker node VPC. For more information, see [Enabling DNS resolution for Amazon EKS cluster endpoints](#).

- **Amazon EC2 bastion host:** You can launch an Amazon EC2 instance into a public subnet in your cluster's VPC and then log in via SSH into that instance to run `kubectl` commands. For more information, see [Linux Bastion Hosts on AWS](#).

When you configure `kubectl` for your bastion host, be sure to use AWS credentials that are already mapped to your cluster's RBAC configuration, or add the IAM user or role that your bastion will use to the RBAC configuration before you remove endpoint public access. For more information,

see [Managing Users or IAM Roles for your Cluster \(p. 115\)](#) and [Unauthorized or Access Denied \(kubectl\) \(p. 179\)](#).

- **Transit Gateway:** A transit gateway is a network transit hub that you can use to interconnect your VPCs and on-premises networks. For more information, see [What is a Transit Gateway?](#) in the *Amazon VPC Transit Gateways* documentation.
- **Amazon VPC connectivity options:** Amazon VPC provides multiple network connectivity options for you to leverage depending on your current network designs and requirements. These connectivity options include leveraging either the internet or an AWS Direct Connect connection as the network backbone and terminating the connection into either AWS or user-managed network endpoints. For more information, see [Amazon Virtual Private Cloud Connectivity Options](#).
- **AWS Cloud9 IDE:** AWS Cloud9 is a cloud-based integrated development environment (IDE) that lets you write, run, and debug your code with just a browser. You can create an AWS Cloud9 IDE in your cluster's VPC and use the IDE to communicate with your cluster. For more information, see [Creating an Environment in AWS Cloud9](#).

When you configure `kubectl` for your AWS Cloud9 IDE, be sure to use AWS credentials that are already mapped to your cluster's RBAC configuration, or add the IAM user or role that your IDE will use to the RBAC configuration before you remove endpoint public access. For more information, see [Managing Users or IAM Roles for your Cluster \(p. 115\)](#) and [Unauthorized or Access Denied \(kubectl\) \(p. 179\)](#).

## Amazon EKS Control Plane Logging

Amazon EKS control plane logging provides audit and diagnostic logs directly from the Amazon EKS control plane to CloudWatch Logs in your account. These logs make it easy for you to secure and run your clusters. You can select the exact log types you need, and logs are sent as log streams to a group for each Amazon EKS cluster in CloudWatch.

You can start using Amazon EKS control plane logging by choosing which log types you want to enable for each new or existing Amazon EKS cluster. You can enable or disable each log type on a per-cluster basis using the AWS Management Console, AWS CLI (version 1.16.139 or higher), or through the Amazon EKS API. When enabled, logs are automatically sent from the Amazon EKS cluster to CloudWatch Logs in the same account.

When you use Amazon EKS control plane logging, you're charged standard Amazon EKS pricing for each cluster that you run. You are charged the standard CloudWatch Logs data ingestion and storage costs for any logs sent to CloudWatch Logs from your clusters. You are also charged for any AWS resources, such as Amazon EC2 instances or Amazon EBS volumes, that you provision as part of your cluster.

The following cluster control plane log types are available. Each log type corresponds to a component of the Kubernetes control plane. To learn more about these components, see [Kubernetes Components](#) in the Kubernetes documentation.

- **Kubernetes API server component logs (api)** – Your cluster's API server is the control plane component that exposes the Kubernetes API. For more information, see [kube-apiserver](#) in the Kubernetes documentation.
- **Audit (audit)** – Kubernetes audit logs provide a record of the individual users, administrators, or system components that have affected your cluster. For more information, see [Auditing](#) in the Kubernetes documentation.
- **Authenticator (authenticator)** – Authenticator logs are unique to Amazon EKS. These logs represent the control plane component that Amazon EKS uses for Kubernetes [Role Based Access Control](#) (RBAC) authentication using IAM credentials. For more information, see [Managing Cluster Authentication \(p. 104\)](#).

- **Controller manager (`controllerManager`)** – The controller manager manages the core control loops that are shipped with Kubernetes. For more information, see [kube-controller-manager](#) in the Kubernetes documentation.
- **Scheduler (`scheduler`)** – The scheduler component manages when and where to run pods in your cluster. For more information, see [kube-scheduler](#) in the Kubernetes documentation.

## Enabling and Disabling Control Plane Logs

By default, cluster control plane logs aren't sent to CloudWatch Logs. You must enable each log type individually to send logs for your cluster. CloudWatch Logs ingestion, archive storage, and data scanning rates apply to enabled control plane logs. For more information, see [CloudWatch Pricing](#).

When you enable a log type, the logs are sent with a log verbosity level of 2.

### To enable or disable control plane logs with the console

1. Open the Amazon EKS console at <https://console.aws.amazon.com/eks/home#/clusters>.
2. Choose the name of the cluster to display your cluster information.
3. Under **Logging**, choose **Update**.
4. For each individual log type, choose whether the log type should be **Enabled** or **Disabled**. By default, each log type is **Disabled**.
5. Choose **Update** to finish.

### To enable or disable control plane logs with the AWS CLI

1. Check your AWS CLI version with the following command.

```
aws --version
```

If your AWS CLI version is below 1.16.139, you must first update to the latest version. To install or upgrade the AWS CLI, see [Installing the AWS Command Line Interface](#) in the *AWS Command Line Interface User Guide*.

2. Update your cluster's control plane log export configuration with the following AWS CLI command. Substitute your cluster name and desired endpoint access values.

#### Note

The following command sends all available log types to CloudWatch Logs.

```
aws eks --region us-west-2 update-cluster-config --name prod \
--logging '{"clusterLogging":[{"types":
["api","audit","authenticator","controllerManager","scheduler"],"enabled":true}]}'
```

Output:

```
{
  "update": {
    "id": "883405c8-65c6-4758-8cee-2a7c1340a6d9",
    "status": "InProgress",
    "type": "LoggingUpdate",
    "params": [
      {
        "type": "ClusterLogging",
        "value": "{\"clusterLogging\":{\"types\":[\"api\",\"audit\",
\"authenticator\",\"controllerManager\",\"scheduler\"],\"enabled\":true}}"
```

```
    },
    "createdAt": 1553271814.684,
    "errors": []
  }
}
```

3. Monitor the status of your log configuration update with the following command, using the cluster name and the update ID that were returned by the previous command. Your update is complete when the status appears as Successful.

```
aws eks --region us-west-2 describe-update --name prod --update-id 883405c8-65c6-4758-8cee-2a7c1340a6d9
```

Output:

```
{
  "update": {
    "id": "883405c8-65c6-4758-8cee-2a7c1340a6d9",
    "status": "Successful",
    "type": "LoggingUpdate",
    "params": [
      {
        "type": "ClusterLogging",
        "value": "{\"clusterLogging\": [{\"types\": [\"api\", \"audit\", \"authenticator\", \"controllerManager\", \"scheduler\"], \"enabled\": true}]}"
      }
    ],
    "createdAt": 1553271814.684,
    "errors": []
  }
}
```

## Viewing Cluster Control Plane Logs

After you have enabled any of the control plane log types for your Amazon EKS cluster, you can view them on the CloudWatch console.

To learn more about viewing, analyzing, and managing logs in CloudWatch, see the [Amazon CloudWatch Logs User Guide](#).

### To view your cluster control plane logs on the CloudWatch console

1. Open the CloudWatch console at <https://console.aws.amazon.com/cloudwatch/home#logs:prefix=/aws/eks>. This URL displays your current available log groups and filters them with the `/aws/eks` prefix.
2. Choose the cluster that you want to view logs for. The log group name format is `/aws/eks/cluster-name/cluster`.
3. Choose the log stream to view. The following list describes the log stream name format for each log type.

#### Note

As log stream data grows, the log stream names are rotated. When multiple log streams exist for a particular log type, you can view the latest log stream by looking for the log stream name with the latest **Last Event Time**.

- **Kubernetes API server component logs (api)** – kube-apiserver-**nnn...**
- **Audit (audit)** – kube-apiserver-audit-**nnn...**

- **Authenticator (authenticator)** – authenticator-*nnn...*
- **Controller manager (controllerManager)** – kube-apiserver-*nnn...*
- **Scheduler (scheduler)** – kube-apiserver-*nnn...*

## Deleting a Cluster

When you're done using an Amazon EKS cluster, you should delete the resources associated with it so that you don't incur any unnecessary costs.

### Important

If you have active services in your cluster that are associated with a load balancer, you must delete those services before deleting the cluster so that the load balancers are deleted properly. Otherwise, you can have orphaned resources in your VPC that prevent you from being able to delete the VPC.

Choose the tab below that corresponds to your preferred cluster deletion method.

eksctl

### To delete an Amazon EKS cluster and worker nodes with eksctl

This procedure assumes that you have installed eksctl, and that your eksctl version is at least 0.1.37. You can check your version with the following command:

```
eksctl version
```

For more information on installing or upgrading eksctl, see [Installing or Upgrading eksctl \(p. 119\)](#).

### Note

This procedure only works for clusters that were created with eksctl.

1. List all services running in your cluster.

```
kubectl get svc --all-namespaces
```

2. Delete any services that have an associated `EXTERNAL-IP` value. These services are fronted by an Elastic Load Balancing load balancer, and you must delete them in Kubernetes to allow the load balancer and associated resources to be properly released.

```
kubectl delete svc service-name
```

3. Delete the cluster and its associated worker nodes with the following command, substituting the red text with your cluster name.

```
eksctl delete cluster --name prod
```

Output:

```
[#] using region us-west-2
[#] deleting EKS cluster "prod"
[#] will delete stack "eksctl-prod-nodegroup-standard-workers"
[#] waiting for stack "eksctl-prod-nodegroup-standard-workers" to get deleted
[#] will delete stack "eksctl-prod-cluster"
[#] the following EKS cluster resource(s) for "prod" will be deleted: cluster. If
    in doubt, check CloudFormation console
```

## AWS Management Console

### To delete an Amazon EKS cluster with the AWS Management Console

1. List all services running in your cluster.

```
kubectl get svc --all-namespaces
```

2. Delete any services that have an associated `EXTERNAL-IP` value. These services are fronted by an Elastic Load Balancing load balancer, and you must delete them in Kubernetes to allow the load balancer and associated resources to be properly released.

```
kubectl delete svc service-name
```

3. Delete the worker node AWS CloudFormation stack.
  - a. Open the AWS CloudFormation console at <https://console.aws.amazon.com/cloudformation>.
  - b. Select the worker node stack to delete and then choose **Actions, Delete Stack**.
  - c. On the **Delete Stack** confirmation screen, choose **Yes, Delete**.
4. Delete the cluster.
  - a. Open the Amazon EKS console at <https://console.aws.amazon.com/eks/home#/clusters>.
  - b. Select the cluster to delete and choose **Delete**.
  - c. On the delete cluster confirmation screen, choose **Delete**.
5. (Optional) Delete the VPC AWS CloudFormation stack.
  - a. Select the VPC stack to delete and choose **Actions** and then **Delete Stack**.
  - b. On the **Delete Stack** confirmation screen, choose **Yes, Delete**.

## AWS CLI

### To delete an Amazon EKS cluster with the AWS CLI

1. List all services running in your cluster.

```
kubectl get svc --all-namespaces
```

2. Delete any services that have an associated `EXTERNAL-IP` value. These services are fronted by an Elastic Load Balancing load balancer, and you must delete them in Kubernetes to allow the load balancer and associated resources to be properly released.

```
kubectl delete svc service-name
```

3. Delete the worker node AWS CloudFormation stack.
  - a. List your available AWS CloudFormation stacks with the following command. Find the worker node template name in the resulting output.

```
aws cloudformation list-stacks --query StackSummaries[].StackName
```

- b. Delete the worker node stack with the following command, substituting the red text with your worker node stack name.

```
aws cloudformation delete-stack --stack-name worker-node-stack
```

4. Delete the cluster with the following command, substituting the red text with your cluster name.

```
aws eks delete-cluster --name my-cluster
```

5. (Optional) Delete the VPC AWS CloudFormation stack.
  - a. List your available AWS CloudFormation stacks with the following command. Find the VPC template name in the resulting output.

```
aws cloudformation list-stacks --query StackSummaries[].StackName
```

- b. Delete the VPC stack with the following command, substituting the red text with your VPC stack name.

```
aws cloudformation delete-stack --stack-name my-vpc-stack
```

## Amazon EKS Kubernetes Versions

The Kubernetes project is rapidly evolving with new features, design updates, and bug fixes. The community releases new Kubernetes minor versions, such as 1.13, as generally available approximately every three months, and each minor version is supported for approximately one year after it is first released.

### Available Amazon EKS Kubernetes Versions

The following Kubernetes versions are currently available for new clusters in Amazon EKS:

- 1.13.7
- 1.12.6
- 1.11.8

#### Important

Kubernetes version 1.10 is no longer supported on Amazon EKS. You can no longer create new 1.10 clusters, and all existing Amazon EKS clusters running Kubernetes version 1.10 will eventually be automatically updated to the latest available platform version of Kubernetes version 1.11. For more information, see [Amazon EKS Version Deprecation \(p. 44\)](#).

Please update any 1.10 clusters to version 1.11 or higher in order to avoid service interruption. For more information, see [Updating an Amazon EKS Cluster Kubernetes Version \(p. 26\)](#).

Unless your application requires a specific version of Kubernetes, we recommend that you choose the latest available Kubernetes version supported by Amazon EKS for your clusters. As new Kubernetes versions become available in Amazon EKS, we recommend that you proactively update your clusters to use the latest available version. For more information, see [Updating an Amazon EKS Cluster Kubernetes Version \(p. 26\)](#).

### Kubernetes 1.13

The following features are now supported in Kubernetes 1.13 Amazon EKS clusters:

- The `PodSecurityPolicy` admission controller is now enabled. This admission controller allows fine-grained control over pod creation and updates. For more information, see [Pod Security Policy \(p. 121\)](#).



- Amazon ECR interface VPC endpoints (AWS PrivateLink) are supported. When you enable these endpoints in your VPC, all network traffic between your VPC and Amazon ECR is restricted to the Amazon network. For more information, see [Amazon ECR Interface VPC Endpoints \(AWS PrivateLink\)](#) in the *Amazon Elastic Container Registry User Guide*.
- The `DryRun` feature is in beta in Kubernetes 1.13 and is enabled by default for Amazon EKS clusters. For more information, see [Dry run](#) in the Kubernetes documentation.
- The `TaintBasedEvictions` feature is in beta in Kubernetes 1.13 and is enabled by default for Amazon EKS clusters. For more information, see [Taint based Evictions](#) in the Kubernetes documentation.
- Raw block volume support is in beta in Kubernetes 1.13 and is enabled by default for Amazon EKS clusters. This is accessible via the `volumeDevices` container field in pod specs, and the `volumeMode` field in persistent volume and persistent volume claim definitions. For more information, see [Raw Block Volume Support](#) in the Kubernetes documentation.
- Node lease renewal is treated as the heartbeat signal from the node, in addition to its `NodeStatus` update. This reduces load on the control plane for large clusters. For more information, see <https://github.com/kubernetes/kubernetes/pull/69241>.

For the complete Kubernetes 1.13 changelog, see <https://github.com/kubernetes/kubernetes/blob/master/CHANGELOG-1.13.md>

## Amazon EKS Version Deprecation

In line with the Kubernetes community support for Kubernetes versions, Amazon EKS is committed to running at least three production-ready versions of Kubernetes at any given time, with a fourth version in deprecation.

We will announce the deprecation of a given Kubernetes minor version at least 60 days before the deprecation date. Because of the Amazon EKS qualification and release process for new Kubernetes versions, the deprecation of a Kubernetes version on Amazon EKS will be on or after the date the Kubernetes project stops supporting the version upstream.

On the deprecation date, Amazon EKS clusters running the version targeted for deprecation will begin to be updated to the next Amazon EKS-supported version of Kubernetes. This means that if the deprecated version is 1.10, clusters will eventually be automatically updated to version 1.11. If a cluster is automatically updated by Amazon EKS, you must also update the version of your worker nodes after the update is complete. For more information, see [Worker Node Updates](#) (p. 63).

Kubernetes supports compatibility between masters and workers for at least 2 minor versions, so 1.10 workers will continue to operate when orchestrated by a 1.11 control plane. For more information, see [Kubernetes Version and Version Skew Support Policy](#) in the Kubernetes documentation.

## Platform Versions

Amazon EKS platform versions represent the capabilities of the cluster control plane, such as which Kubernetes API server flags are enabled, as well as the current Kubernetes patch version. Each Kubernetes minor version has one or more associated Amazon EKS platform versions. The platform versions for different Kubernetes minor versions are independent.

When a new Kubernetes minor version is available in Amazon EKS, such as 1.13, the initial Amazon EKS platform version for that Kubernetes minor version starts at `eks . 1`. However, Amazon EKS releases new platform versions periodically to enable new Kubernetes control plane settings and to provide security fixes.

When new Amazon EKS platform versions become available for a minor version:

- The Amazon EKS platform version number is incremented (`eks . n+1`).
- Amazon EKS automatically upgrades all existing clusters to the latest Amazon EKS platform version for their corresponding Kubernetes minor version.
- Amazon EKS might publish a new worker AMI with a corresponding patch version. However, all patch versions are compatible between the EKS control plane and worker AMIs for a given Kubernetes minor version.

New Amazon EKS platform versions don't introduce breaking changes or cause service interruptions.

**Note**

Automatic upgrades of existing Amazon EKS platform versions are rolled out incrementally. The roll-out process might take some time. If you need the latest Amazon EKS platform version features immediately, you should create a new Amazon EKS cluster.

Clusters are always created with the latest available Amazon EKS platform version (`eks . n`) for the specified Kubernetes version. If you update your cluster to a new Kubernetes minor version, your cluster receives the current Amazon EKS platform version for the Kubernetes minor version that you updated to.

The current and recent Amazon EKS platform versions are described in the following tables.

## Kubernetes version 1.13

Kubernetes Version	Amazon EKS Platform Version	Enabled Admission Controllers	Release Notes	
1.13.8	eks . 2	NamespaceLifecycle, LimitRanger, ServiceAccount, DefaultStorageClass, ResourceQuota, DefaultToleration, NodeRestriction, MutatingAdmissionWebhook, ValidatingAdmissionWebhook, PodSecurityPolicy	New platform version updating Amazon EKS Kubernetes 1.13 clusters to a patched version of Kubernetes 1.13.8 to address CVE-2019-11247.	
1.13.7	eks . 1	NamespaceLifecycle, LimitRanger, ServiceAccount, DefaultStorageClass, ResourceQuota, DefaultToleration, NodeRestriction, MutatingAdmissionWebhook, ValidatingAdmissionWebhook, PodSecurityPolicy	Initial release of Kubernetes 1.13 for Amazon EKS. For more information, see <a href="#">Kubernetes 1.13 (p. 43)</a> .	

## Kubernetes version 1.12

Kubernetes Version	Amazon EKS Platform Version	Enabled Admission Controllers	Release Notes	
1.12.10	eks.3	NamespaceLifecycle, LimitRanger, ServiceAccount, DefaultStorageClass, ResourceQuota, DefaultTolerationSeconds, NodeRestriction, MutatingAdmissionWebhook, ValidatingAdmissionWebhook	New platform version updating Amazon EKS Kubernetes 1.12 clusters to a patched version of Kubernetes 1.12.10 to address CVE-2019-11247.	
1.12.6	eks.2	NamespaceLifecycle, LimitRanger, ServiceAccount, DefaultStorageClass, ResourceQuota, DefaultTolerationSeconds, NodeRestriction, MutatingAdmissionWebhook, ValidatingAdmissionWebhook	New platform version to support custom DNS names in the Kubelet certificate and improve etcd performance. This fixes a bug that caused worker node Kubelet daemons to request a new certificate every few seconds.	
1.12.6	eks.1	NamespaceLifecycle, LimitRanger, ServiceAccount, DefaultStorageClass, ResourceQuota, DefaultTolerationSeconds, NodeRestriction, MutatingAdmissionWebhook, ValidatingAdmissionWebhook	Initial release of Kubernetes 1.12 for Amazon EKS.	

## Kubernetes version 1.11

Kubernetes Version	Amazon EKS Platform Version	Enabled Admission Controllers	Release Notes	
1.11.10	eks.4	NamespaceLifecycle, LimitRanger, ServiceAccount, DefaultStorageClass, ResourceQuota, DefaultTolerationSeconds, NodeRestriction,	New platform version updating Amazon EKS Kubernetes 1.11 clusters to a patched version of	

Kubernetes Version	Amazon EKS Platform Version	Enabled Admission Controllers	Release Notes	
		MutatingAdmissionWebhook, ValidatingAdmissionWebhook	1.10.0 to address CVE-2019-11247.	
1.11.8	eks.3	NamespaceLifecycle, LimitRanger, ServiceAccount, DefaultStorageClass, ResourceQuota, DefaultTolerationSeconds, NodeRestriction, MutatingAdmissionWebhook, ValidatingAdmissionWebhook	New platform version to support custom DNS names in the Kubelet certificate and improve etcd performance.	
1.11.8	eks.2	NamespaceLifecycle, LimitRanger, ServiceAccount, DefaultStorageClass, ResourceQuota, DefaultTolerationSeconds, NodeRestriction, MutatingAdmissionWebhook, ValidatingAdmissionWebhook	New platform version updating Amazon EKS Kubernetes 1.11 clusters to patch level 1.11.8 to address CVE-2019-1002100.	
1.11.5	eks.1	NamespaceLifecycle, LimitRanger, ServiceAccount, DefaultStorageClass, ResourceQuota, DefaultTolerationSeconds, NodeRestriction, MutatingAdmissionWebhook, ValidatingAdmissionWebhook	Initial release of Kubernetes 1.11 for Amazon EKS.	

## Kubernetes version 1.10

### Important

Kubernetes version 1.10 is no longer supported on Amazon EKS. You can no longer create new 1.10 clusters, and all existing Amazon EKS clusters running Kubernetes version 1.10 will eventually be automatically updated to the latest available platform version of Kubernetes version 1.11. For more information, see [Amazon EKS Version Deprecation \(p. 44\)](#). Please update any 1.10 clusters to version 1.11 or higher in order to avoid service interruption. For more information, see [Updating an Amazon EKS Cluster Kubernetes Version \(p. 26\)](#).

Kubernetes Version	Amazon EKS Platform Version	Enabled Admission Controllers	Release Notes	
1.10.13	eks.5	NamespaceLifecycle, LimitRanger, ServiceAccount, DefaultStorageClass, ResourceQuota,	New platform version to support custom DNS names in the Kubelet certificate	

Kubernetes Version	Amazon EKS Platform Version	Enabled Admission Controllers	Release Notes	
		DefaultTolerationSeconds, etcd NodeRestriction, performance. MutatingAdmissionWebhook, ValidatingAdmissionWebhook	Added support for etcd NodeRestriction, performance. MutatingAdmissionWebhook, ValidatingAdmissionWebhook: "v1.10.13-eks-4a9600" to address <a href="#">CVE-2019-11247</a> .	
1.10.13	eks.4	NamespaceLifecycle, LimitRanger, ServiceAccount, DefaultStorageClass, ResourceQuota, DefaultTolerationSeconds, NodeRestriction, MutatingAdmissionWebhook, ValidatingAdmissionWebhook	New platform version updating Kubernetes to patch level 1.10.13 and a patch to address <a href="#">CVE-2019-1002100</a> .	
1.10.11	eks.3	NamespaceLifecycle, LimitRanger, ServiceAccount, DefaultStorageClass, ResourceQuota, DefaultTolerationSeconds, NodeRestriction, MutatingAdmissionWebhook, ValidatingAdmissionWebhook	New platform version updating Kubernetes to patch level 1.10.11 to address <a href="#">CVE-2019-1002105</a> .	
1.10.3	eks.2	NamespaceLifecycle, LimitRanger, ServiceAccount, DefaultStorageClass, ResourceQuota, DefaultTolerationSeconds, NodeRestriction, MutatingAdmissionWebhook, ValidatingAdmissionWebhook	<ul style="list-style-type: none"> <li>Added support for Kubernetes <a href="#">aggregation layer</a>.</li> <li>Added support for Kubernetes <a href="#">Horizontal Pod Autoscaler (HPA)</a>.</li> <li>Kubernetes Metrics Server 0.3.0 or greater is compatible with EKS platform version eks.2.</li> </ul>	
1.10.3	eks.1	NamespaceLifecycle, LimitRanger, ServiceAccount, DefaultStorageClass, ResourceQuota, DefaultTolerationSeconds, NodeRestriction	Initial launch of Amazon EKS.	

# Worker Nodes

Worker machines in Kubernetes are called nodes. Amazon EKS worker nodes run in your AWS account and connect to your cluster's control plane via the cluster API server endpoint.

Amazon EKS worker nodes are standard Amazon EC2 instances, and you are billed for them based on normal EC2 On-Demand prices. For more information, see [Amazon EC2 Pricing](#).

By default, Amazon EKS provides AWS CloudFormation templates to spin up worker nodes in your Amazon EKS cluster. This AMI is built on top of Amazon Linux 2, and is configured to serve as the base image for Amazon EKS worker nodes. The AMI is configured to work with Amazon EKS out of the box, and it includes Docker, **kubelet**, and the AWS IAM Authenticator. The AMI also contains a specialized bootstrap script that allows it to discover and connect to your cluster's control plane automatically.

## Note

You can track security or privacy events for Amazon Linux 2 at the [Amazon Linux Security Center](#) or subscribe to the associated [RSS feed](#). Security and privacy events include an overview of the issue, what packages are affected, and how to update your instances to correct the issue.

The AWS CloudFormation worker node template launches your worker nodes with specialized Amazon EC2 user data. This user data triggers a specialized [bootstrap script](#) that allows your worker nodes to discover and connect to your cluster's control plane automatically. For more information, see [Launching Amazon EKS Worker Nodes \(p. 56\)](#).

For more information about worker nodes from a general Kubernetes perspective, see [Nodes](#) in the Kubernetes documentation.

## Topics

- [Amazon EKS-Optimized AMI \(p. 49\)](#)
- [Amazon EKS Partner AMIs \(p. 56\)](#)
- [Launching Amazon EKS Worker Nodes \(p. 56\)](#)
- [Worker Node Updates \(p. 63\)](#)

## Amazon EKS-Optimized AMI

The Amazon EKS-optimized AMI is built on top of Amazon Linux 2, and is configured to serve as the base image for Amazon EKS worker nodes. The AMI is configured to work with Amazon EKS out of the box, and it includes Docker, **kubelet**, and the AWS IAM Authenticator.

## Note

You can track security or privacy events for Amazon Linux 2 at the [Amazon Linux Security Center](#) or subscribe to the associated [RSS feed](#). Security and privacy events include an overview of the issue, what packages are affected, and how to update your instances to correct the issue.

The AMI IDs for the latest Amazon EKS-optimized AMI (with and without [GPU support \(p. 52\)](#)) are shown in the following table.

## Note

The Amazon EKS-optimized AMI with GPU support only supports P2 and P3 instance types. Be sure to specify these instance types in your worker node AWS CloudFormation template. By using the Amazon EKS-optimized AMI with GPU support, you agree to [NVIDIA's end user license agreement \(EULA\)](#).

Kubernetes version 1.13.7

Region	Amazon EKS-optimized AMI	with GPU support
US East (Ohio) (us-east-2)	ami-0485258c2d1c3608f	ami-0ccac9d9b57864000
US East (N. Virginia) (us-east-1)	ami-0f2e8e5663e16b436	ami-0017d945a10387606
US West (Oregon) (us-west-2)	ami-03a55127c613349a7	ami-08335952e837d087b
Asia Pacific (Hong Kong) (ap-east-1)	ami-032850771ac6f8ae2	N/A*
Asia Pacific (Mumbai) (ap-south-1)	ami-0a9b1c1807b1a40ab	ami-005b754faac73f0cc
Asia Pacific (Tokyo) (ap-northeast-1)	ami-0fde798d17145fae1	ami-04cf69bbd6c0fae0b
Asia Pacific (Seoul) (ap-northeast-2)	ami-07fd7609df6c8e39b	ami-0730e699ed0118737
Asia Pacific (Singapore) (ap-southeast-1)	ami-0361e14efd56a71c7	ami-07be5e97a529cd146
Asia Pacific (Sydney) (ap-southeast-2)	ami-0237d87bc27daba65	ami-0a2f4c3aeb596aa7e
EU (Frankfurt) (eu-central-1)	ami-0b7127e7a2a38802a	ami-0fbbd205f797ecccd
EU (Ireland) (eu-west-1)	ami-00ac2e6b3cb38a9b9	ami-0f9571a3e65dc4e20
EU (London) (eu-west-2)	ami-0147919d2ff9a6ad5	ami-032348bd69c5dd665
EU (Paris) (eu-west-3)	ami-0537ee9329c1628a2	ami-053962359d6859fec
EU (Stockholm) (eu-north-1)	ami-0fd05922165907b85	ami-0641def7f02a4cac5

Kubernetes version 1.12.7

Region	Amazon EKS-optimized AMI	with GPU support
US East (Ohio) (us-east-2)	ami-0fe61ae4c397e710d	ami-067d88fb64d3d7990
US East (N. Virginia) (us-east-1)	ami-0e380e0a62d368837	ami-06e46a15650294dfa
US West (Oregon) (us-west-2)	ami-0355c210cb3f58aa2	ami-084e8e620163aa50e
Asia Pacific (Hong Kong) (ap-east-1)	ami-03885e0a7493b073f	N/A*
Asia Pacific (Mumbai) (ap-south-1)	ami-01b6a163133c31994	ami-09ad3a49fb13389a0

Region	Amazon EKS-optimized AMI	with GPU support
Asia Pacific (Tokyo) (ap-northeast-1)	ami-0a9b3f8b4b65b402b	ami-0cd09d7293f31df8a
Asia Pacific (Seoul) (ap-northeast-2)	ami-069f6a654a8795f72	ami-006549812c03748cb
Asia Pacific (Singapore) (ap-southeast-1)	ami-03737a1ac334a5767	ami-01be8fddd9b16320c
Asia Pacific (Sydney) (ap-southeast-2)	ami-07580768e8538626f	ami-0a1bf783357dd8492
EU (Frankfurt) (eu-central-1)	ami-0ee5ca4231511cafc	ami-0ae5976723472b6d4
EU (Ireland) (eu-west-1)	ami-0404d23c7e8188740	ami-042f9abf2f96a0097
EU (London) (eu-west-2)	ami-07346d8553f83f9d6	ami-0b87e9246afd42760
EU (Paris) (eu-west-3)	ami-038cb36289174bac4	ami-0d9405868a6e9ee11
EU (Stockholm) (eu-north-1)	ami-03e60b5a990893129	ami-0122b7e2a6736e3c5

Kubernetes version 1.11.9

Region	Amazon EKS-optimized AMI	with GPU support
US East (Ohio) (us-east-2)	ami-03c6648b74285020f	ami-0b87186dda80931ee
US East (N. Virginia) (us-east-1)	ami-0a5f5d5b0f6f58199	ami-07207754196c1a8fc
US West (Oregon) (us-west-2)	ami-057d1c0dcb254a878	ami-052da6a4e0ae156ad
Asia Pacific (Hong Kong) (ap-east-1)	ami-006e7a52666ab7494	N/A*
Asia Pacific (Mumbai) (ap-south-1)	ami-00f1adebe5ab9a431	ami-04645af6384529c5d
Asia Pacific (Tokyo) (ap-northeast-1)	ami-0a0b6606652f9b3b9	ami-0a8f4e1f9bf09a81f
Asia Pacific (Seoul) (ap-northeast-2)	ami-0c84b3f055cda1afb	ami-01db6bb089f6adfcf
Asia Pacific (Singapore) (ap-southeast-1)	ami-05e92412054db3f87	ami-0e001196bd450aa0c
Asia Pacific (Sydney) (ap-southeast-2)	ami-07eb76498b1ba6cd6	ami-0c7132a332aa55aa6
EU (Frankfurt) (eu-central-1)	ami-0234bc9c2b341aa02	ami-05cb4f6e8be8b83f1
EU (Ireland) (eu-west-1)	ami-06902949103360023	ami-02f337476a5c33f1b



Region	Amazon EKS-optimized AMI	with GPU support
EU (London) (eu-west-2)	ami-0db100ad46c7966d2	ami-0aa2208dbb9bb7cc5
EU (Paris) (eu-west-3)	ami-052046d313576d0ba	ami-0f6ea479cb4e7a4d2
EU (Stockholm) (eu-north-1)	ami-02ebf24da505128f9	ami-078c260b9a737fc35

\* GPU instance types are not available in the Asia Pacific (Hong Kong) (ap-east-1) region, so Amazon EKS does not publish the Amazon EKS-optimized AMI with GPU support in that region.

### Important

These AMIs require the latest AWS CloudFormation worker node template. You can't use these AMIs with a previous version of the worker node template; they will fail to join your cluster. Be sure to upgrade any existing AWS CloudFormation worker stacks with the latest template (URL shown below) before you attempt to use these AMIs.

```
https://amazon-eks.s3-us-west-2.amazonaws.com/cloudformation/2019-02-11/amazon-eks-nodegroup.yaml
```

The AWS CloudFormation worker node template launches your worker nodes with Amazon EC2 user data that triggers a specialized [bootstrap script](#). This script allows your worker nodes to discover and connect to your cluster's control plane automatically. For more information, see [Launching Amazon EKS Worker Nodes](#) (p. 56).

## Amazon EKS-Optimized AMI Build Scripts

Amazon Elastic Kubernetes Service (Amazon EKS) has open-sourced the build scripts that are used to build the Amazon EKS-optimized AMI. These build scripts are now available [on GitHub](#).

The Amazon EKS-optimized AMI is built on top of Amazon Linux 2, specifically for use as a worker node in Amazon EKS clusters. You can use this repository to view the specifics of how the Amazon EKS team configures **kubelet**, Docker, the AWS IAM Authenticator for Kubernetes, and more.

The build scripts repository includes a [HashiCorp Packer](#) template and build scripts to generate an AMI. These scripts are the source of truth for Amazon EKS-optimized AMI builds, so you can follow the GitHub repository to monitor changes to our AMIs. For example, perhaps you want your own AMI to use the same version of Docker that the EKS team uses for the official AMI.

The GitHub repository also contains the specialized [bootstrap script](#) that runs at boot time to configure your instance's certificate data, control plane endpoint, cluster name, and more.

Additionally, the GitHub repository contains our Amazon EKS worker node AWS CloudFormation templates. These templates make it easier to spin up an instance running the Amazon EKS-optimized AMI and register it with a cluster.

For more information, see the repositories on GitHub at <https://github.com/aws-labs/amazon-eks-ami>.

## Amazon EKS-Optimized AMI with GPU Support

The Amazon EKS-optimized AMI with GPU support is built on top of the standard Amazon EKS-optimized AMI, and is configured to serve as an optional image for Amazon EKS worker nodes to support GPU workloads.

In addition to the standard Amazon EKS-optimized AMI configuration, the GPU AMI includes the following:

- NVIDIA drivers
- The `nvidia-docker2` package
- The `nvidia-container-runtime` (as the default runtime)

The AMI IDs for the latest Amazon EKS-optimized AMI with GPU support are shown in the following table.

**Note**

The Amazon EKS-optimized AMI with GPU support only supports P2 and P3 instance types. Be sure to specify these instance types in your worker node AWS CloudFormation template. By using the Amazon EKS-optimized AMI with GPU support, you agree to [NVIDIA's end user license agreement \(EULA\)](#).

Kubernetes version 1.13.7

Region	Amazon EKS-optimized AMI with GPU support
US East (Ohio) (us-east-2)	ami-0ccac9d9b57864000
US East (N. Virginia) (us-east-1)	ami-0017d945a10387606
US West (Oregon) (us-west-2)	ami-08335952e837d087b
Asia Pacific (Hong Kong) (ap-east-1)	N/A*
Asia Pacific (Mumbai) (ap-south-1)	ami-005b754faac73f0cc
Asia Pacific (Tokyo) (ap-northeast-1)	ami-04cf69bbd6c0fae0b
Asia Pacific (Seoul) (ap-northeast-2)	ami-0730e699ed0118737
Asia Pacific (Singapore) (ap-southeast-1)	ami-07be5e97a529cd146
Asia Pacific (Sydney) (ap-southeast-2)	ami-0a2f4c3aeb596aa7e
EU (Frankfurt) (eu-central-1)	ami-0fbbd205f797ecccd
EU (Ireland) (eu-west-1)	ami-0f9571a3e65dc4e20
EU (London) (eu-west-2)	ami-032348bd69c5dd665
EU (Paris) (eu-west-3)	ami-053962359d6859fec
EU (Stockholm) (eu-north-1)	ami-0641def7f02a4cac5

Kubernetes version 1.12.7

Region	Amazon EKS-optimized AMI with GPU support
US East (Ohio) (us-east-2)	ami-067d88fb64d3d7990
US East (N. Virginia) (us-east-1)	ami-06e46a15650294dfa
US West (Oregon) (us-west-2)	ami-084e8e620163aa50e
Asia Pacific (Hong Kong) (ap-east-1)	N/A*
Asia Pacific (Mumbai) (ap-south-1)	ami-09ad3a49fb13389a0

Region	Amazon EKS-optimized AMI with GPU support
Asia Pacific (Tokyo) (ap-northeast-1)	ami-0cd09d7293f31df8a
Asia Pacific (Seoul) (ap-northeast-2)	ami-006549812c03748cb
Asia Pacific (Singapore) (ap-southeast-1)	ami-01be8fddd9b16320c
Asia Pacific (Sydney) (ap-southeast-2)	ami-0a1bf783357dd8492
EU (Frankfurt) (eu-central-1)	ami-0ae5976723472b6d4
EU (Ireland) (eu-west-1)	ami-042f9abf2f96a0097
EU (London) (eu-west-2)	ami-0b87e9246afd42760
EU (Paris) (eu-west-3)	ami-0d9405868a6e9ee11
EU (Stockholm) (eu-north-1)	ami-0122b7e2a6736e3c5

#### Kubernetes version 1.11.9

Region	Amazon EKS-optimized AMI with GPU support
US East (Ohio) (us-east-2)	ami-0b87186dda80931ee
US East (N. Virginia) (us-east-1)	ami-07207754196c1a8fc
US West (Oregon) (us-west-2)	ami-052da6a4e0ae156ad
Asia Pacific (Hong Kong) (ap-east-1)	N/A*
Asia Pacific (Mumbai) (ap-south-1)	ami-04645af6384529c5d
Asia Pacific (Tokyo) (ap-northeast-1)	ami-0a8f4e1f9bf09a81f
Asia Pacific (Seoul) (ap-northeast-2)	ami-01db6bb089f6adfcf
Asia Pacific (Singapore) (ap-southeast-1)	ami-0e001196bd450aa0c
Asia Pacific (Sydney) (ap-southeast-2)	ami-0c7132a332aa55aa6
EU (Frankfurt) (eu-central-1)	ami-05cb4f6e8be8b83f1
EU (Ireland) (eu-west-1)	ami-02f337476a5c33f1b
EU (London) (eu-west-2)	ami-0aa2208dbb9bb7cc5
EU (Paris) (eu-west-3)	ami-0f6ea479cb4e7a4d2
EU (Stockholm) (eu-north-1)	ami-078c260b9a737fc35

\* GPU instance types are not available in the Asia Pacific (Hong Kong) (ap-east-1) region, so Amazon EKS does not publish the Amazon EKS-optimized AMI with GPU support in that region.

#### Important

These AMIs require the latest AWS CloudFormation worker node template. You can't use these AMIs with a previous version of the worker node template; they will fail to join your cluster. Be sure to upgrade any existing AWS CloudFormation worker stacks with the latest template (URL shown below) before you attempt to use these AMIs.

```
https://amazon-eks.s3-us-west-2.amazonaws.com/cloudformation/2019-02-11/amazon-eks-  
nodegroup.yaml
```

The AWS CloudFormation worker node template launches your worker nodes with Amazon EC2 user data that triggers a specialized [bootstrap script](#). This script allows your worker nodes to discover and connect to your cluster's control plane automatically. For more information, see [Launching Amazon EKS Worker Nodes](#) (p. 56).

After your GPU worker nodes join your cluster, you must apply the [NVIDIA device plugin for Kubernetes](#) as a DaemonSet on your cluster with the following command.

```
kubectl apply -f https://raw.githubusercontent.com/NVIDIA/k8s-device-plugin/1.0.0-beta/  
nvidia-device-plugin.yaml
```

You can verify that your nodes have allocatable GPUs with the following command:

```
kubectl get nodes "-o=custom-columns=NAME:.metadata.name,GPU:.status.allocatable.nvidia  
\.com/gpu"
```

## Example GPU Manifest

This section provides an example pod manifest for you to test that your GPU workers are configured properly.

### Example Get `nvidia-smi` output

This example pod manifest launches a Cuda container that runs `nvidia-smi` on a worker node. Create a file called `nvidia-smi.yaml`, copy and paste the following manifest into it, and save the file.

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: nvidia-smi  
spec:  
  restartPolicy: OnFailure  
  containers:  
    - name: nvidia-smi  
      image: nvidia/cuda:9.2-devel  
      args:  
        - "nvidia-smi"  
      resources:  
        limits:  
          nvidia.com/gpu: 1
```

Apply the manifest with the following command:

```
kubectl apply -f nvidia-smi.yaml
```

After the pod has finished running, view its logs with the following command:

```
kubectl logs nvidia-smi
```

Output:

```
Mon Aug  6 20:23:31 2018  
+-----+
```

NVIDIA-SMI 396.26				Driver Version: 396.26			
-----							
GPU Name		Persistence-M		Bus-Id	Disp.A	Volatile	Uncorr. ECC
Fan	Temp	Perf	Pwr:Usage/Cap	Memory-Usage		GPU-Util	Compute M.
=====							
0	Tesla	V100-SXM2...	On	00000000:00:1C.0	Off		0
N/A	46C	P0	47W / 300W	0MiB / 16160MiB		0%	Default
-----							
-----							
Processes:						GPU Memory	
GPU	PID	Type	Process name			Usage	
=====							
No running processes found							
-----							

## Amazon EKS Partner AMIs

In addition to the official Amazon EKS-optimized, Canonical has partnered with Amazon EKS to create worker node AMIs that you can use in your clusters.

[Canonical](#) delivers a built-for-purpose Kubernetes Node OS image. This minimized Ubuntu image is optimized for Amazon EKS and includes the custom AWS kernel that is jointly developed with AWS. For more information, see [Ubuntu and Amazon Elastic Kubernetes Service](#) and [Optimized Support for Amazon EKS on Ubuntu 18.04](#).

## Launching Amazon EKS Worker Nodes

This topic helps you to launch an Auto Scaling group of worker nodes that register with your Amazon EKS cluster. After the nodes join the cluster, you can deploy Kubernetes applications to them.

If this is your first time launching Amazon EKS worker nodes, we recommend that you follow one of our [Getting Started with Amazon EKS \(p. 3\)](#) guides instead. They provide complete end-to-end walkthroughs for creating an Amazon EKS cluster with worker nodes.

### Important

Amazon EKS worker nodes are standard Amazon EC2 instances, and you are billed for them based on normal Amazon EC2 On-Demand Instance prices. For more information, see [Amazon EC2 Pricing](#).

Choose the tab below that corresponds to your desired worker node creation method:

`eksctl`

### To launch worker nodes with `eksctl`

This procedure assumes that you have installed `eksctl`, and that your `eksctl` version is at least 0.1.37. You can check your version with the following command:

```
eksctl version
```

For more information on installing or upgrading `eksctl`, see [Installing or Upgrading eksctl \(p. 119\)](#).

### Note

This procedure only works for clusters that were created with `eksctl`.

- Create your worker node group with the following command. Substitute the red text with your own values.

```
eksctl create nodegroup \  
--cluster default \  
--version auto \  
--name standard-workers \  
--node-type t3.medium \  
--node-ami auto \  
--nodes 3 \  
--nodes-min 1 \  
--nodes-max 4
```

### Note

For more information on the available options for **eksctl create nodegroup**, see the project [README on GitHub](#) or view the help page with the following command.

```
eksctl create nodegroup --help
```

Output:

```
[#] using region us-west-2  
[#] will use version 1.12 for new nodegroup(s) based on control plane version  
[#] nodegroup "standard-workers" will use  
    "ami-0923e4b35a30a5f53" [AmazonLinux2/1.12]  
[#] 1 nodegroup (standard-workers) was included  
[#] will create a CloudFormation stack for each of 1 nodegroups in cluster  
    "default"  
[#] 1 task: { create nodegroup "standard-workers" }  
[#] building nodegroup stack "eksctl-default-nodegroup-standard-workers"  
[#] deploying stack "eksctl-default-nodegroup-standard-workers"  
[#] adding role "arn:aws:iam::111122223333:role/eksctl-default-nodegroup-standard-  
NodeInstanceRole-12C2J0814XSEE" to auth ConfigMap  
[#] nodegroup "standard-workers" has 0 node(s)  
[#] waiting for at least 1 node(s) to become ready in "standard-workers"  
[#] nodegroup "standard-workers" has 3 node(s)  
[#] node "ip-192-168-52-42.us-west-2.compute.internal" is ready  
[#] node "ip-192-168-7-27.us-west-2.compute.internal" is not ready  
[#] node "ip-192-168-76-138.us-west-2.compute.internal" is not ready  
[#] created 1 nodegroup(s) in cluster "default"  
[#] checking security group configuration for all nodegroups  
[#] all nodegroups have up-to-date configuration
```

## AWS Management Console

### To launch your worker nodes with the AWS Management Console

These procedures have the following prerequisites:

- You have created a VPC and security group that meet the requirements for an Amazon EKS cluster. For more information, see [Cluster VPC Considerations \(p. 81\)](#) and [Cluster Security Group Considerations \(p. 83\)](#). The [Getting Started with Amazon EKS \(p. 3\)](#) guide creates a VPC that meets the requirements, or you can also follow [Creating a VPC for Your Amazon EKS Cluster \(p. 79\)](#) to create one manually.
- You have created an Amazon EKS cluster and specified that it use the VPC and security group that meet the requirements of an Amazon EKS cluster. For more information, see [Creating an Amazon EKS Cluster \(p. 20\)](#).

1. Wait for your cluster status to show as **ACTIVE**. If you launch your worker nodes before the cluster is active, the worker nodes will fail to register with the cluster and you will have to relaunch them.
2. Open the AWS CloudFormation console at <https://console.aws.amazon.com/cloudformation>.
3. From the navigation bar, select a Region that supports Amazon EKS.
4. Choose **Create stack**.
5. For **Choose a template**, select **Specify an Amazon S3 template URL**.
6. Paste the following URL into the text area and choose **Next**.

```
https://amazon-eks.s3-us-west-2.amazonaws.com/cloudformation/2019-02-11/amazon-eks-nodegroup.yaml
```

#### Note

If you intend to only deploy worker nodes to private subnets, you should edit this template in the AWS CloudFormation designer and modify the `AssociatePublicIpAddress` parameter in the `NodeLaunchConfig` to be `false`.

```
AssociatePublicIpAddress: 'false'
```

7. On the **Specify Details** page, fill out the following parameters accordingly, and choose **Next**:
  - **Stack name** – Choose a stack name for your AWS CloudFormation stack. For example, you can call it **<cluster-name>-worker-nodes**.
  - **ClusterName** – Enter the name that you used when you created your Amazon EKS cluster.

#### Important

This name must exactly match your Amazon EKS cluster name. Otherwise, your worker nodes will be unable to join it.

- **ClusterControlPlaneSecurityGroup** – Enter the security group or groups that you used when you created your Amazon EKS cluster. This AWS CloudFormation template creates a worker node security group that allows traffic to and from the cluster control plane security group specified.

#### Important

The worker node AWS CloudFormation template modifies the security group that you specify here, so **Amazon EKS strongly recommends that you use a dedicated security group for each cluster control plane (one per cluster)**. If this security group is shared with other resources, you might block or disrupt connections to those resources.

- **NodeGroupName** – Enter a name for your node group. This name can be used later to identify the Auto Scaling node group that is created for your worker nodes.
- **NodeAutoScalingGroupMinSize** – Enter the minimum number of nodes to which your worker node Auto Scaling group can scale in.
- **NodeAutoScalingGroupDesiredCapacity** – Enter the desired number of nodes to scale to when your stack is created.
- **NodeAutoScalingGroupMaxSize** – Enter the maximum number of nodes to which your worker node Auto Scaling group can scale out. This value must be at least one node greater than your desired capacity so that you can perform a rolling update of your worker nodes without reducing your node count during the update.
- **NodeInstanceType** – Choose an instance type for your worker nodes. The instance type and size that you choose determines how many IP addresses are available per worker node for the containers in your pods. For more information, see [IP Addresses Per Network Interface Per Instance Type](#) in the *Amazon EC2 User Guide for Linux Instances*.

### Note

The supported instance types for the latest version of the [Amazon VPC CNI plugin for Kubernetes](#) are shown [here](#). You may need to update your CNI version to take advantage of the latest supported instance types. For more information, see [Amazon VPC CNI Plugin for Kubernetes Upgrades \(p. 97\)](#).

### Important

Some instance types might not be available in all regions.

- **NodeImageId** – Enter the current Amazon EKS worker node AMI ID for your Region. The AMI IDs for the latest Amazon EKS-optimized AMI (with and without [GPU support \(p. 52\)](#)) are shown in the following table. Be sure to choose the correct AMI ID for your desired Kubernetes version and AWS region.

### Note

The Amazon EKS-optimized AMI with GPU support only supports P2 and P3 instance types. Be sure to specify these instance types in your worker node AWS CloudFormation template. By using the Amazon EKS-optimized AMI with GPU support, you agree to [NVIDIA's end user license agreement \(EULA\)](#).

Kubernetes version 1.13.7

Region	Amazon EKS-optimized AMI	with GPU support
US East (Ohio) (us-east-2)	ami-0485258c2d1c3608f	ami-0ccac9d9b57864000
US East (N. Virginia) (us-east-1)	ami-0f2e8e5663e16b436	ami-0017d945a10387606
US West (Oregon) (us-west-2)	ami-03a55127c613349a7	ami-08335952e837d087b
Asia Pacific (Hong Kong) (ap-east-1)	ami-032850771ac6f8ae2	N/A*
Asia Pacific (Mumbai) (ap-south-1)	ami-0a9b1c1807b1a40ab	ami-005b754faac73f0cc
Asia Pacific (Tokyo) (ap-northeast-1)	ami-0fde798d17145fae1	ami-04cf69bbd6c0fae0b
Asia Pacific (Seoul) (ap-northeast-2)	ami-07fd7609df6c8e39b	ami-0730e699ed0118737
Asia Pacific (Singapore) (ap-southeast-1)	ami-0361e14efd56a71c7	ami-07be5e97a529cd146
Asia Pacific (Sydney) (ap-southeast-2)	ami-0237d87bc27daba65	ami-0a2f4c3aeb596aa7e
EU (Frankfurt) (eu-central-1)	ami-0b7127e7a2a38802a	ami-0fbbd205f797eccccd
EU (Ireland) (eu-west-1)	ami-00ac2e6b3cb38a9b9	ami-0f9571a3e65dc4e20
EU (London) (eu-west-2)	ami-0147919d2ff9a6ad5	ami-032348bd69c5dd665
EU (Paris) (eu-west-3)	ami-0537ee9329c1628a2	ami-053962359d6859fec



Region	Amazon EKS-optimized AMI	with GPU support
EU (Stockholm) (eu-north-1)	ami-0fd05922165907b85	ami-0641def7f02a4cac5

Kubernetes version 1.12.7

Region	Amazon EKS-optimized AMI	with GPU support
US East (Ohio) (us-east-2)	ami-0fe61ae4c397e710d	ami-067d88fb64d3d7990
US East (N. Virginia) (us-east-1)	ami-0e380e0a62d368837	ami-06e46a15650294dfa
US West (Oregon) (us-west-2)	ami-0355c210cb3f58aa2	ami-084e8e620163aa50e
Asia Pacific (Hong Kong) (ap-east-1)	ami-03885e0a7493b073f	N/A*
Asia Pacific (Mumbai) (ap-south-1)	ami-01b6a163133c31994	ami-09ad3a49fb13389a0
Asia Pacific (Tokyo) (ap-northeast-1)	ami-0a9b3f8b4b65b402b	ami-0cd09d7293f31df8a
Asia Pacific (Seoul) (ap-northeast-2)	ami-069f6a654a8795f72	ami-006549812c03748cb
Asia Pacific (Singapore) (ap-southeast-1)	ami-03737a1ac334a5767	ami-01be8fddd9b16320c
Asia Pacific (Sydney) (ap-southeast-2)	ami-07580768e8538626f	ami-0a1bf783357dd8492
EU (Frankfurt) (eu-central-1)	ami-0ee5ca4231511cafc	ami-0ae5976723472b6d4
EU (Ireland) (eu-west-1)	ami-0404d23c7e8188740	ami-042f9abf2f96a0097
EU (London) (eu-west-2)	ami-07346d8553f83f9d6	ami-0b87e9246afd42760
EU (Paris) (eu-west-3)	ami-038cb36289174bac4	ami-0d9405868a6e9ee11
EU (Stockholm) (eu-north-1)	ami-03e60b5a990893129	ami-0122b7e2a6736e3c5

Kubernetes version 1.11.9

Region	Amazon EKS-optimized AMI	with GPU support
US East (Ohio) (us-east-2)	ami-03c6648b74285020f	ami-0b87186dda80931ee
US East (N. Virginia) (us-east-1)	ami-0a5f5d5b0f6f58199	ami-07207754196c1a8fc
US West (Oregon) (us-west-2)	ami-057d1c0dcb254a878	ami-052da6a4e0ae156ad
Asia Pacific (Hong Kong) (ap-east-1)	ami-006e7a52666ab7494	N/A*
Asia Pacific (Mumbai) (ap-south-1)	ami-00f1adebe5ab9a431	ami-04645af6384529c5d
Asia Pacific (Tokyo) (ap-northeast-1)	ami-0a0b6606652f9b3b9	ami-0a8f4e1f9bf09a81f
Asia Pacific (Seoul) (ap-northeast-2)	ami-0c84b3f055cda1afb	ami-01db6bb089f6adfcf
Asia Pacific (Singapore) (ap-southeast-1)	ami-05e92412054db3f87	ami-0e001196bd450aa0c
Asia Pacific (Sydney) (ap-southeast-2)	ami-07eb76498b1ba6cd6	ami-0c7132a332aa55aa6
EU (Frankfurt) (eu-central-1)	ami-0234bc9c2b341aa02	ami-05cb4f6e8be8b83f1
EU (Ireland) (eu-west-1)	ami-06902949103360023	ami-02f337476a5c33f1b
EU (London) (eu-west-2)	ami-0db100ad46c7966d2	ami-0aa2208dbb9bb7cc5
EU (Paris) (eu-west-3)	ami-052046d313576d0ba	ami-0f6ea479cb4e7a4d2
EU (Stockholm) (eu-north-1)	ami-02ebf24da505128f9	ami-078c260b9a737fc35

\* GPU instance types are not available in the Asia Pacific (Hong Kong) (ap-east-1) region, so Amazon EKS does not publish the Amazon EKS-optimized AMI with GPU support in that region.

#### Note

The Amazon EKS worker node AMI is based on Amazon Linux 2. You can track security or privacy events for Amazon Linux 2 at the [Amazon Linux Security Center](#) or subscribe to the associated [RSS feed](#). Security and privacy events include an overview of the issue, what packages are affected, and how to update your instances to correct the issue.

- **KeyName** – Enter the name of an Amazon EC2 SSH key pair that you can use to connect using SSH into your worker nodes with after they launch. If you don't already have an Amazon

EC2 keypair, you can create one in the AWS Management Console. For more information, see [Amazon EC2 Key Pairs](#) in the *Amazon EC2 User Guide for Linux Instances*.

**Note**

If you do not provide a keypair here, the AWS CloudFormation stack creation fails.

- **BootstrapArguments** – Specify any optional arguments to pass to the worker node bootstrap script, such as extra **kubelet** arguments. For more information, view the bootstrap script usage information at <https://github.com/awslabs/amazon-eks-ami/blob/master/files/bootstrap.sh>
  - **VpcId** – Enter the ID for the VPC that your worker nodes should launch into.
  - **Subnets** – Choose the subnets within the preceding VPC that your worker nodes should launch into. If you are launching worker nodes into only private subnets, do not include public subnets here.
8. On the **Options** page, you can choose to tag your stack resources. Choose **Next**.
  9. On the **Review** page, review your information, acknowledge that the stack might create IAM resources, and then choose **Create**.
  10. When your stack has finished creating, select it in the console and choose **Outputs**.
  11. Record the **NodeInstanceRole** for the node group that was created. You need this when you configure your Amazon EKS worker nodes.

**To enable worker nodes to join your cluster**

1. Download, edit, and apply the AWS IAM Authenticator configuration map.
  - a. Use the following command to download the configuration map:

```
curl -o aws-auth-cm.yaml https://amazon-eks.s3-us-west-2.amazonaws.com/cloudformation/2019-02-11/aws-auth-cm.yaml
```

- b. Open the file with your favorite text editor. Replace the *<ARN of instance role (not instance profile)>* snippet with the **NodeInstanceRole** value that you recorded in the previous procedure, and save the file.

**Important**

Do not modify any other lines in this file.

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: aws-auth
  namespace: kube-system
data:
  mapRoles: |
    - rolearn: <ARN of instance role (not instance profile)>
      username: system:node:{{EC2PrivateDNSName}}
      groups:
        - system:bootstrappers
        - system:nodes
```

- c. Apply the configuration. This command may take a few minutes to finish.

```
kubectl apply -f aws-auth-cm.yaml
```

**Note**

If you receive the error "aws-iam-authenticator": executable file not found in \$PATH, your **kubectl** isn't configured for Amazon EKS. For more information, see [Installing aws-iam-authenticator](#) (p. 108).

If you receive any other authorization or resource type errors, see [Unauthorized or Access Denied \(kubectl\)](#) (p. 179) in the troubleshooting section.

2. Watch the status of your nodes and wait for them to reach the Ready status.

```
kubectl get nodes --watch
```

3. (GPU workers only) If you chose a P2 or P3 instance type and the Amazon EKS-optimized AMI with GPU support, you must apply the [NVIDIA device plugin for Kubernetes](#) as a DaemonSet on your cluster with the following command.

```
kubectl apply -f https://raw.githubusercontent.com/NVIDIA/k8s-device-plugin/1.0.0-beta/nvidia-device-plugin.yml
```

## Worker Node Updates

When a new Amazon EKS-optimized AMI is released, you should consider replacing the nodes in your worker node group with the new AMI. Likewise, if you have updated the Kubernetes version for your Amazon EKS cluster, you should also update the worker nodes to use worker nodes with the same Kubernetes version.

There are two basic ways to update the worker nodes in your clusters to use a new AMI: create a new worker node group and migrate your pods to that group, or update the AWS CloudFormation stack for an existing worker node group to use the new AMI. This latter method is not supported for worker node groups that were created with `eksctl`.

Migrating to a new worker node group is more graceful than simply updating the AMI ID in an existing AWS CloudFormation stack, because the migration process taints the old node group as `NoSchedule` and drains the nodes after a new stack is ready to accept the existing pod workload.

### Topics

- [Migrating to a New Worker Node Group](#) (p. 63)
- [Updating an Existing Worker Node Group](#) (p. 68)

## Migrating to a New Worker Node Group

This topic helps you to create a new worker node group, gracefully migrate your existing applications to the new group, and then remove the old worker node group from your cluster.

`eksctl`

### To migrate your applications to a new worker node group with `eksctl`

This procedure assumes that you have installed `eksctl`, and that your `eksctl` version is at least 0.1.37. You can check your version with the following command:

```
eksctl version
```

For more information on installing or upgrading `eksctl`, see [Installing or Upgrading eksctl](#) (p. 119).

#### Note

This procedure only works for clusters and worker node groups that were created with `eksctl`.

1. Retrieve the name of your existing worker node groups, substituting the red text with your cluster name.

```
eksctl get nodegroups --cluster=default
```

Output:

CLUSTER	NODEGROUP	CREATED	MIN SIZE	MAX SIZE
	DESIRED CAPACITY	INSTANCE TYPE	IMAGE ID	
default	standard-workers	2019-05-01T22:26:58Z	1	4
	t3.medium	ami-05a71d034119ffc12		3

2. Launch a new worker node group with eksctl with the following command, substituting the red text with your own values.

**Note**

For more available flags and their descriptions, see <https://eksctl.io/>.

```
eksctl create nodegroup \  
--cluster default \  
--version 1.13 \  
--name standard-1-13 \  
--node-type t3.medium \  
--nodes 3 \  
--nodes-min 1 \  
--nodes-max 4 \  
--node-ami auto
```

3. When the previous command completes, verify that all of your worker nodes have reached the Ready state with the following command:

```
kubectl get nodes
```

4. Delete the original node group with the following command, substituting the red text with your cluster and nodegroup names:

```
eksctl delete nodegroup --cluster default --name standard-workers
```

## AWS Management Console

### To migrate your applications to a new worker node group with the AWS Management Console

1. Launch a new worker node group by following the steps outlined in [Launching Amazon EKS Worker Nodes \(p. 56\)](#).
2. When your stack has finished creating, select it in the console and choose **Outputs**.
3. Record the **NodeInstanceRole** for the node group that was created. You need this to add the new Amazon EKS worker nodes to your cluster.

**Note**

If you have attached any additional IAM policies to your old node group IAM role, such as adding permissions for the Kubernetes [Cluster Autoscaler](#), you should attach those same policies to your new node group IAM role to maintain that functionality on the new group.

4. Update the security groups for both worker node groups so that they can communicate with each other. For more information, see [Cluster Security Group Considerations \(p. 83\)](#).

- a. Record the security group IDs for both worker node groups. This is shown as the **NodeSecurityGroup** value in the AWS CloudFormation stack outputs.

You can use the following AWS CLI commands to get the security group IDs from the stack names. In these commands, `oldNodes` is the AWS CloudFormation stack name for your older worker node stack, and `newNodes` is the name of the stack that you are migrating to.

```
oldNodes="<old_node_CFN_stack_name>"
newNodes="<new_node_CFN_stack_name>"

oldSecGroup=$(aws cloudformation describe-stack-resources --stack-name
  $oldNodes \
  --query 'StackResources[?
ResourceType==`AWS::EC2::SecurityGroup`.PhysicalResourceId' \
  --output text)
newSecGroup=$(aws cloudformation describe-stack-resources --stack-name
  $newNodes \
  --query 'StackResources[?
ResourceType==`AWS::EC2::SecurityGroup`.PhysicalResourceId' \
  --output text)
```

- b. Add ingress rules to each worker node security group so that they accept traffic from each other.

The following AWS CLI commands add ingress rules to each security group that allow all traffic on all protocols from the other security group. This configuration allows pods in each worker node group to communicate with each other while you are migrating your workload to the new group.

```
aws ec2 authorize-security-group-ingress --group-id $oldSecGroup \
--source-group $newSecGroup --protocol -1
aws ec2 authorize-security-group-ingress --group-id $newSecGroup \
--source-group $oldSecGroup --protocol -1
```

5. Edit the `aws-auth` configmap to map the new worker node instance role in RBAC.

```
kubectl edit configmap -n kube-system aws-auth
```

Add a new `mapRoles` entry for the new worker node group.

```
apiVersion: v1
data:
  mapRoles: |
    - rolearn: <ARN of instance role (not instance profile)>
      username: system:node:{{EC2PrivateDNSName}}
      groups:
        - system:bootstrappers
        - system:nodes
    - rolearn: arn:aws:iam::111122223333:role/workers-1-10-NodeInstanceRole-
      U11V27W93CX5
      username: system:node:{{EC2PrivateDNSName}}
      groups:
        - system:bootstrappers
        - system:nodes
```

Replace the *<ARN of instance role (not instance profile)>* snippet with the **NodeInstanceRole** value that you recorded in [Step 3](#) (p. 64), then save and close the file to apply the updated configmap.

6. Watch the status of your nodes and wait for your new worker nodes to join your cluster and reach the Ready status.

```
kubectl get nodes --watch
```

7. (Optional) If you are using the Kubernetes [Cluster Autoscaler](#), scale the deployment down to 0 replicas to avoid conflicting scaling actions.

```
kubectl scale deployments/cluster-autoscaler --replicas=0 -n kube-system
```

8. Use the following command to taint each of the nodes that you want to remove with NoSchedule so that new pods are not scheduled or rescheduled on the nodes you are replacing:

```
kubectl taint nodes node_name key=value:NoSchedule
```

If you are upgrading your worker nodes to a new Kubernetes version, you can identify and taint all of the nodes of a particular Kubernetes version (in this case, 1.10.3) with the following code snippet.

```
K8S_VERSION=1.10.3
nodes=$(kubectl get nodes -o jsonpath="{.items[?(@.status.nodeInfo.kubeletVersion==\v$K8S_VERSION\")]}.metadata.name}")
for node in ${nodes[@]}
do
    echo "Tainting $node"
    kubectl taint nodes $node key=value:NoSchedule
done
```

9. Determine your cluster's DNS provider.

```
kubectl get deployments -l k8s-app=kube-dns -n kube-system
```

Output (this cluster is using kube-dns for DNS resolution, but your cluster may return coredns instead):

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
<b>kube-dns</b>	1	1	1	1	31m

10. If your current deployment is running fewer than two replicas, scale out the deployment to two replicas. Substitute coredns for kube-dns if your previous command output returned that instead.

```
kubectl scale deployments/kube-dns --replicas=2 -n kube-system
```

11. Drain each of the nodes that you want to remove from your cluster with the following command:

```
kubectl drain node_name --ignore-daemonsets --delete-local-data
```

If you are upgrading your worker nodes to a new Kubernetes version, you can identify and drain all of the nodes of a particular Kubernetes version (in this case, 1.10.3) with the following code snippet.

```
K8S_VERSION=1.10.3
```

```
nodes=$(kubectl get nodes -o jsonpath="{.items[?(@.status.nodeInfo.kubeletVersion==\n\"v$K8S_VERSION\")].metadata.name}")
for node in ${nodes[@]}
do
    echo "Draining $node"
    kubectl drain $node --ignore-daemonsets --delete-local-data
done
```

12. After your old worker nodes have finished draining, revoke the security group ingress rules you authorized earlier, and then delete the AWS CloudFormation stack to terminate the instances.

**Note**

If you have attached any additional IAM policies to your old node group IAM role, such as adding permissions for the Kubernetes [Cluster Autoscaler](#), you must detach those additional policies from the role before you can delete your AWS CloudFormation stack.

- a. Revoke the ingress rules that you created for your worker node security groups earlier. In these commands, `oldNodes` is the AWS CloudFormation stack name for your older worker node stack, and `newNodes` is the name of the stack that you are migrating to.

```
oldNodes="<old_node_CFN_stack_name>"
newNodes="<new_node_CFN_stack_name>"

oldSecGroup=$(aws cloudformation describe-stack-resources --stack-name
  $oldNodes \
  --query 'StackResources[?
ResourceType==`AWS::EC2::SecurityGroup`.PhysicalResourceId' \
  --output text)
newSecGroup=$(aws cloudformation describe-stack-resources --stack-name
  $newNodes \
  --query 'StackResources[?
ResourceType==`AWS::EC2::SecurityGroup`.PhysicalResourceId' \
  --output text)
aws ec2 revoke-security-group-ingress --group-id $oldSecGroup \
  --source-group $newSecGroup --protocol -1
aws ec2 revoke-security-group-ingress --group-id $newSecGroup \
  --source-group $oldSecGroup --protocol -1
```

- b. Open the AWS CloudFormation console at <https://console.aws.amazon.com/cloudformation>.
  - c. Select your old worker node stack.
  - d. Choose **Actions**, then **Delete stack**.
13. Edit the `aws-auth` configmap to remove the old worker node instance role from RBAC.

```
kubectl edit configmap -n kube-system aws-auth
```

Delete the `mapRoles` entry for the old worker node group.

```
apiVersion: v1
data:
  mapRoles: |
    - rolearn: arn:aws:iam::111122223333:role/workers-1-11-NodeInstanceRole-
W70725MZQFF8
      username: system:node:{{EC2PrivateDNSName}}
      groups:
        - system:bootstrappers
        - system:nodes
    - rolearn: arn:aws:iam::111122223333:role/workers-1-10-NodeInstanceRole-
U11V27W93CX5
      username: system:node:{{EC2PrivateDNSName}}
      groups:
```



```
- system:bootstrappers
- system:nodes
```

Save and close the file to apply the updated configmap.

14. (Optional) If you are using the Kubernetes [Cluster Autoscaler](#), scale the deployment back to one replica.

**Note**

You must also tag your new Auto Scaling group appropriately (for example, `k8s.io/cluster-autoscaler/enabled,k8s.io/cluster-autoscaler/<YOUR CLUSTER NAME>`) and update your Cluster Autoscaler deployment's command to point to the newly tagged Auto Scaling group. For more information, see [Cluster Autoscaler on AWS](#).

```
kubectl scale deployments/cluster-autoscaler --replicas=1 -n kube-system
```

15. (Optional) Verify that you are using the latest version of the [Amazon VPC CNI plugin for Kubernetes](#). You may need to update your CNI version to take advantage of the latest supported instance types. For more information, see [Amazon VPC CNI Plugin for Kubernetes Upgrades \(p. 97\)](#).
16. If your cluster is using `kube-dns` for DNS resolution (see step [Step 9 \(p. 66\)](#)), scale in the `kube-dns` deployment to one replica.

```
kubectl scale deployments/kube-dns --replicas=1 -n kube-system
```

## Updating an Existing Worker Node Group

This topic helps you to update an existing AWS CloudFormation worker node stack with a new AMI. You can use this procedure to update your worker nodes to a new version of Kubernetes following a cluster update, or you can update to the latest Amazon EKS-optimized AMI for an existing Kubernetes version.

The latest default Amazon EKS worker node AWS CloudFormation template is configured to launch an instance with the new AMI into your cluster before removing an old one, one at a time. This configuration ensures that you always have your Auto Scaling group's desired count of active instances in your cluster during the rolling update.

**Note**

This method is not supported for worker node groups that were created with `eksctl`. If you created your cluster or worker node group with `eksctl`, see [Migrating to a New Worker Node Group \(p. 63\)](#).

### To update an existing worker node group

1. Determine your cluster's DNS provider.

```
kubectl get deployments -l k8s-app=kube-dns -n kube-system
```

Output (this cluster is using `kube-dns` for DNS resolution, but your cluster may return `coredns` instead):

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
<code>kube-dns</code>	1	1	1	1	31m

2. If your current deployment is running fewer than two replicas, scale out the deployment to two replicas. Substitute `coredns` for `kube-dns` if your previous command output returned that instead.

```
kubectl scale deployments/kube-dns --replicas=2 -n kube-system
```

3. (Optional) If you are using the Kubernetes [Cluster Autoscaler](#), scale the deployment down to zero replicas to avoid conflicting scaling actions.

```
kubectl scale deployments/cluster-autoscaler --replicas=0 -n kube-system
```

4. Determine the instance type and desired instance count of your current worker node group. You will enter these values later when you update the AWS CloudFormation template for the group.
  - a. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
  - b. Choose **Launch Configurations** in the left navigation, and note the instance type for your existing worker node launch configuration.
  - c. Choose **Auto Scaling Groups** in the left navigation and note the **Desired** instance count for your existing worker node Auto Scaling group.
5. Open the AWS CloudFormation console at <https://console.aws.amazon.com/cloudformation>.
6. Select your worker node group stack, and then choose **Actions, Update stack**.
7. For **Choose a template**, select **Specify an Amazon S3 template URL**.
8. Paste the following URL into the text area to ensure that you are using the latest version of the worker node AWS CloudFormation template, and then choose **Next**:

```
https://amazon-eks.s3-us-west-2.amazonaws.com/cloudformation/2019-02-11/amazon-eks-nodegroup.yaml
```

9. On the **Specify Details** page, fill out the following parameters, and choose **Next**:
  - **NodeAutoScalingGroupDesiredCapacity** – Enter the desired instance count that you recorded in [Step 4 \(p. 69\)](#), or enter a new desired number of nodes to scale to when your stack is updated.
  - **NodeAutoScalingGroupMaxSize** – Enter the maximum number of nodes to which your worker node Auto Scaling group can scale out. **This value must be at least one node greater than your desired capacity so that you can perform a rolling update of your worker nodes without reducing your node count during the update.**
  - **NodeInstanceType** – Choose the instance type you recorded in [Step 4 \(p. 69\)](#), or choose a different instance type for your worker nodes.

**Note**

The supported instance types for the latest version of the [Amazon VPC CNI plugin for Kubernetes](#) are shown [here](#). You may need to update your CNI version to take advantage of the latest supported instance types. For more information, see [Amazon VPC CNI Plugin for Kubernetes Upgrades \(p. 97\)](#).

**Important**

Some instance types might not be available in all regions.

- **NodeImageId** – Enter the current Amazon EKS worker node AMI ID for your Region. The AMI IDs for the latest Amazon EKS-optimized AMI (with and without [GPU support \(p. 52\)](#)) are shown in the following table.

**Note**

The Amazon EKS-optimized AMI with GPU support only supports P2 and P3 instance types. Be sure to specify these instance types in your worker node AWS CloudFormation template. By using the Amazon EKS-optimized AMI with GPU support, you agree to [NVIDIA's end user license agreement \(EULA\)](#).

Kubernetes version 1.13.7

Region	Amazon EKS-optimized AMI	with GPU support
US East (Ohio) (us-east-2)	ami-0485258c2d1c3608f	ami-0ccac9d9b57864000
US East (N. Virginia) (us-east-1)	ami-0f2e8e5663e16b436	ami-0017d945a10387606
US West (Oregon) (us-west-2)	ami-03a55127c613349a7	ami-08335952e837d087b
Asia Pacific (Hong Kong) (ap-east-1)	ami-032850771ac6f8ae2	N/A*
Asia Pacific (Mumbai) (ap-south-1)	ami-0a9b1c1807b1a40ab	ami-005b754faac73f0cc
Asia Pacific (Tokyo) (ap-northeast-1)	ami-0fde798d17145fae1	ami-04cf69bbd6c0fae0b
Asia Pacific (Seoul) (ap-northeast-2)	ami-07fd7609df6c8e39b	ami-0730e699ed0118737
Asia Pacific (Singapore) (ap-southeast-1)	ami-0361e14efd56a71c7	ami-07be5e97a529cd146
Asia Pacific (Sydney) (ap-southeast-2)	ami-0237d87bc27daba65	ami-0a2f4c3aeb596aa7e
EU (Frankfurt) (eu-central-1)	ami-0b7127e7a2a38802a	ami-0fbbd205f797eccccd
EU (Ireland) (eu-west-1)	ami-00ac2e6b3cb38a9b9	ami-0f9571a3e65dc4e20
EU (London) (eu-west-2)	ami-0147919d2ff9a6ad5	ami-032348bd69c5dd665
EU (Paris) (eu-west-3)	ami-0537ee9329c1628a2	ami-053962359d6859fec
EU (Stockholm) (eu-north-1)	ami-0fd05922165907b85	ami-0641def7f02a4cac5

Kubernetes version 1.12.7

Region	Amazon EKS-optimized AMI	with GPU support
US East (Ohio) (us-east-2)	ami-0fe61ae4c397e710d	ami-067d88fb64d3d7990
US East (N. Virginia) (us-east-1)	ami-0e380e0a62d368837	ami-06e46a15650294dfa
US West (Oregon) (us-west-2)	ami-0355c210cb3f58aa2	ami-084e8e620163aa50e
Asia Pacific (Hong Kong) (ap-east-1)	ami-03885e0a7493b073f	N/A*

Region	Amazon EKS-optimized AMI	with GPU support
Asia Pacific (Mumbai) (ap-south-1)	ami-01b6a163133c31994	ami-09ad3a49fb13389a0
Asia Pacific (Tokyo) (ap-northeast-1)	ami-0a9b3f8b4b65b402b	ami-0cd09d7293f31df8a
Asia Pacific (Seoul) (ap-northeast-2)	ami-069f6a654a8795f72	ami-006549812c03748cb
Asia Pacific (Singapore) (ap-southeast-1)	ami-03737a1ac334a5767	ami-01be8fddd9b16320c
Asia Pacific (Sydney) (ap-southeast-2)	ami-07580768e8538626f	ami-0a1bf783357dd8492
EU (Frankfurt) (eu-central-1)	ami-0ee5ca4231511cafc	ami-0ae5976723472b6d4
EU (Ireland) (eu-west-1)	ami-0404d23c7e8188740	ami-042f9abf2f96a0097
EU (London) (eu-west-2)	ami-07346d8553f83f9d6	ami-0b87e9246afd42760
EU (Paris) (eu-west-3)	ami-038cb36289174bac4	ami-0d9405868a6e9ee11
EU (Stockholm) (eu-north-1)	ami-03e60b5a990893129	ami-0122b7e2a6736e3c5

Kubernetes version 1.11.9

Region	Amazon EKS-optimized AMI	with GPU support
US East (Ohio) (us-east-2)	ami-03c6648b74285020f	ami-0b87186dda80931ee
US East (N. Virginia) (us-east-1)	ami-0a5f5d5b0f6f58199	ami-07207754196c1a8fc
US West (Oregon) (us-west-2)	ami-057d1c0dc254a878	ami-052da6a4e0ae156ad
Asia Pacific (Hong Kong) (ap-east-1)	ami-006e7a52666ab7494	N/A*
Asia Pacific (Mumbai) (ap-south-1)	ami-00f1adebe5ab9a431	ami-04645af6384529c5d
Asia Pacific (Tokyo) (ap-northeast-1)	ami-0a0b6606652f9b3b9	ami-0a8f4e1f9bf09a81f
Asia Pacific (Seoul) (ap-northeast-2)	ami-0c84b3f055cda1afb	ami-01db6bb089f6adfcf
Asia Pacific (Singapore) (ap-southeast-1)	ami-05e92412054db3f87	ami-0e001196bd450aa0c
Asia Pacific (Sydney) (ap-southeast-2)	ami-07eb76498b1ba6cd6	ami-0c7132a332aa55aa6

Region	Amazon EKS-optimized AMI	with GPU support
EU (Frankfurt) (eu-central-1)	ami-0234bc9c2b341aa02	ami-05cb4f6e8be8b83f1
EU (Ireland) (eu-west-1)	ami-06902949103360023	ami-02f337476a5c33f1b
EU (London) (eu-west-2)	ami-0db100ad46c7966d2	ami-0aa2208dbb9bb7cc5
EU (Paris) (eu-west-3)	ami-052046d313576d0ba	ami-0f6ea479cb4e7a4d2
EU (Stockholm) (eu-north-1)	ami-02ebf24da505128f9	ami-078c260b9a737fc35

\* GPU instance types are not available in the Asia Pacific (Hong Kong) (ap-east-1) region, so Amazon EKS does not publish the Amazon EKS-optimized AMI with GPU support in that region.

**Note**

The Amazon EKS worker node AMI is based on Amazon Linux 2. You can track security or privacy events for Amazon Linux 2 at the [Amazon Linux Security Center](#) or subscribe to the associated [RSS feed](#). Security and privacy events include an overview of the issue, what packages are affected, and how to update your instances to correct the issue.

10. (Optional) On the **Options** page, tag your stack resources. Choose **Next**.
11. On the **Review** page, review your information, acknowledge that the stack might create IAM resources, and then choose **Update**.

**Note**

Wait for the update to complete before performing the next steps.

12. If your cluster's DNS provider is kube-dns, scale in the kube-dns deployment to one replica.

```
kubectl scale deployments/kube-dns --replicas=1 -n kube-system
```

13. (Optional) If you are using the Kubernetes [Cluster Autoscaler](#), scale the deployment back to one replica.

```
kubectl scale deployments/cluster-autoscaler --replicas=1 -n kube-system
```

14. (Optional) Verify that you are using the latest version of the [Amazon VPC CNI plugin for Kubernetes](#). You may need to update your CNI version to take advantage of the latest supported instance types. For more information, see [Amazon VPC CNI Plugin for Kubernetes Upgrades](#) (p. 97).

# Storage Classes

Amazon EKS clusters that were created prior to Kubernetes version 1.11 were not created with any storage classes. You must define storage classes for your cluster to use and you should define a default storage class for your persistent volume claims. For more information, see [Storage Classes](#) in the Kubernetes documentation.

## To create an AWS storage class for your Amazon EKS cluster

1. Create an AWS storage class manifest file for your storage class. The `gp2-storage-class.yaml` example below defines a storage class called `gp2` that uses the Amazon EBS `gp2` volume type.

For more information about the options available for AWS storage classes, see [AWS EBS](#) in the Kubernetes documentation.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
  name: gp2
  annotations:
    storageclass.kubernetes.io/is-default-class: "true"
provisioner: kubernetes.io/aws-ebs
parameters:
  type: gp2
  fsType: ext4
```

2. Use **kubectl** to create the storage class from the manifest file.

```
kubectl create -f gp2-storage-class.yaml
```

Output:

```
storageclass "gp2" created
```

## To define a default storage class

1. List the existing storage classes for your cluster. A storage class must be defined before you can set it as a default.

```
kubectl get storageclass
```

Output:

NAME	PROVISIONER	AGE
gp2	kubernetes.io/aws-ebs	8m

2. Choose a storage class and set it as your default by setting the `storageclass.kubernetes.io/is-default-class=true` annotation.

```
kubectl patch storageclass gp2 -p '{"metadata": {"annotations": {"storageclass.kubernetes.io/is-default-class": "true"}}}'
```

Output:

```
storageclass "gp2" patched
```

3. Verify that the storage class is now set as default.

```
kubectl get storageclass
```

Output:

```
gp2 (default)   kubernetes.io/aws-ebs   12m
```

# Load Balancing and Ingress

This chapter covers common load balancing and Ingress configuration for Amazon EKS clusters.

## Topics

- [Load Balancing \(p. 75\)](#)
- [ALB Ingress Controller on Amazon EKS \(p. 76\)](#)

## Load Balancing

Amazon EKS supports the Network Load Balancer and the Classic Load Balancer through the Kubernetes service of type `LoadBalancer`. The configuration of your load balancer is controlled by annotations that are added to the manifest for your service.

By default, Classic Load Balancers are used for `LoadBalancer` type services. To use the Network Load Balancer instead, apply the following annotation to your service:

```
service.beta.kubernetes.io/aws-load-balancer-type: nlb
```

For more information about using Network Load Balancer with Kubernetes, see [Network Load Balancer support on AWS](#) in the Kubernetes documentation.

By default, services of type `LoadBalancer` create public-facing load balancers. To use an internal load balancer, apply the following annotation to your service:

```
service.beta.kubernetes.io/aws-load-balancer-internal: 0.0.0.0/0
```

For internal load balancers, your Amazon EKS cluster must be configured to use at least one private subnet in your VPC. Kubernetes examines the route table for your subnets to identify whether they are public or private. Public subnets have a route directly to the internet using an internet gateway, but private subnets do not.

## Subnet Tagging for Load Balancers

Public subnets in your VPC may be tagged accordingly so that Kubernetes knows to use only those subnets for external load balancers, instead of choosing a public subnet in each Availability Zone (in lexicographical order by subnet ID):

Key	Value
<code>kubernetes.io/role/elb</code>	1

Private subnets in your VPC should be tagged accordingly so that Kubernetes knows that it can use them for internal load balancers:



Key	Value
kubernetes.io/role/internal-elb	1

## ALB Ingress Controller on Amazon EKS

The [AWS ALB Ingress Controller for Kubernetes](#) is a controller that triggers the creation of an Application Load Balancer and the necessary supporting AWS resources whenever an Ingress resource is created on the cluster with the `kubernetes.io/ingress.class: alb` annotation. The Ingress resource uses the ALB to route HTTP or HTTPS traffic to different endpoints within the cluster.

To ensure that your Ingress objects use the ALB Ingress Controller, add the following annotation to your Ingress specification. For more information, see [Ingress specification](#) in the documentation.

```
annotations:
  kubernetes.io/ingress.class: alb
```

For other available annotations supported by the ALB Ingress Controller, see [Ingress annotations](#).

This topic shows you how to configure the ALB Ingress Controller to work with your Amazon EKS cluster.

### To deploy the ALB Ingress Controller to an Amazon EKS cluster

1. Tag the subnets in your VPC that you want to use for your load balancers so that the ALB Ingress Controller knows that it can use them.
  - Public subnets in your VPC should be tagged accordingly so that Kubernetes knows to use only those subnets for external load balancers.

Key	Value
kubernetes.io/role/elb	1

- Private subnets in your VPC should be tagged accordingly so that Kubernetes knows that it can use them for internal load balancers:

Key	Value
kubernetes.io/role/internal-elb	1

2. Create an IAM policy called `ALBIngressControllerIAMPolicy` for your worker node instance profile that allows the ALB Ingress Controller to make calls to AWS APIs on your behalf. Use the following AWS CLI command to create the IAM policy in your AWS account. You can view the policy document [on GitHub](#).

```
aws iam create-policy \
  --policy-name ALBIngressControllerIAMPolicy \
  --policy-document https://raw.githubusercontent.com/kubernetes-sigs/aws-alb-ingress-
  controller/v1.1.2/docs/examples/iam-policy.json
```

Take note of the policy ARN that is returned.

3. Get the IAM role name for your worker nodes. Use the following command to print the `aws-auth` configmap.

```
kubectl -n kube-system describe configmap aws-auth
```

Output:

```
Name:          aws-auth
Namespace:     kube-system
Labels:        <none>
Annotations:   <none>

Data
====
mapRoles:
-----
- groups:
  - system:bootstrappers
  - system:nodes
  rolearn: arn:aws:iam::111122223333:role/eksctl-alb-nodegroup-ng-b1f603c5-NodeInstanceRole-GKNS581EASPU
  username: system:node:{{EC2PrivateDNSName}}

Events:   <none>
```

Record the role name for any `rolearn` values that have the `system:nodes` group assigned to them. In the above example output, the role name is **`eksctl-alb-nodegroup-ng-b1f603c5-NodeInstanceRole-GKNS581EASPU`**. You should have one value for each node group in your cluster.

4. Attach the new `ALBIngressControllerIAMPolicy` IAM policy to each of the worker node IAM roles you identified earlier with the following command, substituting the red text with your own AWS account number and worker node IAM role name.

```
aws iam attach-role-policy \
--policy-arn arn:aws:iam::111122223333:policy/ALBIngressControllerIAMPolicy \
--role-name eksctl-alb-nodegroup-ng-b1f603c5-NodeInstanceRole-GKNS581EASPU
```

5. Create a service account, cluster role, and cluster role binding for the ALB Ingress Controller to use with the following command.

```
kubectl apply -f https://raw.githubusercontent.com/kubernetes-sigs/aws-alb-ingress-controller/v1.1.2/docs/examples/rbac-role.yaml
```

6. Deploy the ALB Ingress Controller with the following command.

```
kubectl apply -f https://raw.githubusercontent.com/kubernetes-sigs/aws-alb-ingress-controller/v1.1.2/docs/examples/alb-ingress-controller.yaml
```

7. Open the ALB Ingress Controller deployment manifest for editing with the following command.

```
kubectl edit deployment.apps/alb-ingress-controller -n kube-system
```

8. Add the cluster name, VPC ID, and AWS Region name for your cluster after the `--ingress-class=alb` line and then save and close the file.

```
spec:
  containers:
  - args:
    - --ingress-class=alb
    - --cluster-name=my_cluster
```

```
- --aws-vpc-id=vpc-03468a8157edca5bd  
- --aws-region=us-west-2
```

## To deploy a sample application

1. Deploy a sample application to verify that the ALB Ingress Controller creates an Application Load Balancer as a result of the Ingress object. Use the following commands to deploy the game [2048](#) as a sample application.

```
kubectl apply -f https://raw.githubusercontent.com/kubernetes-sigs/aws-alb-ingress-controller/v1.1.2/docs/examples/2048/2048-namespace.yaml  
kubectl apply -f https://raw.githubusercontent.com/kubernetes-sigs/aws-alb-ingress-controller/v1.1.2/docs/examples/2048/2048-deployment.yaml  
kubectl apply -f https://raw.githubusercontent.com/kubernetes-sigs/aws-alb-ingress-controller/v1.1.2/docs/examples/2048/2048-service.yaml  
kubectl apply -f https://raw.githubusercontent.com/kubernetes-sigs/aws-alb-ingress-controller/v1.1.2/docs/examples/2048/2048-ingress.yaml
```

2. After a few minutes, verify that the Ingress resource was created with the following command.

```
kubectl get ingress/2048-ingress -n 2048-game
```

Output:

NAME	HOSTS	ADDRESS
PORTS	AGE	
2048-ingress	*	<b>example-2048game-2048ingr-6fa0-352729433.us-west-2.elb.amazonaws.com</b>
	80	24h

3. Open a browser and navigate to the ADDRESS URL from the previous command output to see the sample application.
4. When you finish experimenting with your sample application, delete it with the following commands.

```
kubectl delete -f https://raw.githubusercontent.com/kubernetes-sigs/aws-alb-ingress-controller/v1.1.2/docs/examples/2048/2048-ingress.yaml  
kubectl delete -f https://raw.githubusercontent.com/kubernetes-sigs/aws-alb-ingress-controller/v1.1.2/docs/examples/2048/2048-service.yaml  
kubectl delete -f https://raw.githubusercontent.com/kubernetes-sigs/aws-alb-ingress-controller/v1.1.2/docs/examples/2048/2048-deployment.yaml  
kubectl delete -f https://raw.githubusercontent.com/kubernetes-sigs/aws-alb-ingress-controller/v1.1.2/docs/examples/2048/2048-namespace.yaml
```

# Amazon EKS Networking

This chapter covers networking considerations for running Kubernetes on Amazon EKS.

## Topics

- [Creating a VPC for Your Amazon EKS Cluster \(p. 79\)](#)
- [Cluster VPC Considerations \(p. 81\)](#)
- [Cluster Security Group Considerations \(p. 83\)](#)
- [Pod Networking \(p. 85\)](#)
- [CNI Configuration Variables \(p. 87\)](#)
- [Installing CoreDNS \(p. 88\)](#)
- [External Source Network Address Translation \(SNAT\) \(p. 91\)](#)
- [CNI Custom Networking \(p. 94\)](#)
- [Amazon VPC CNI Plugin for Kubernetes Upgrades \(p. 97\)](#)
- [Installing Calico on Amazon EKS \(p. 97\)](#)

## Creating a VPC for Your Amazon EKS Cluster

This topic guides you through creating a VPC for your cluster with either 3 public subnets, or two public subnets and two private subnets, which are provided with internet access through a NAT gateway. You can use this VPC for your Amazon EKS cluster. We recommend a network architecture that uses private subnets for your worker nodes, and public subnets for Kubernetes to create public load balancers within.

Choose the tab below that represents your desired VPC configuration.

Only public subnets

### To create your cluster VPC with only public subnets

1. Open the AWS CloudFormation console at <https://console.aws.amazon.com/cloudformation>.
2. From the navigation bar, select a Region that supports Amazon EKS.
3. Choose **Create stack**.
4. For **Choose a template**, select **Specify an Amazon S3 template URL**.
5. Paste the following URL into the text area and choose **Next**:

```
https://amazon-eks.s3-us-west-2.amazonaws.com/cloudformation/2019-02-11/amazon-eks-vpc-sample.yaml
```

6. On the **Specify Details** page, fill out the parameters accordingly, and then choose **Next**.
  - **Stack name:** Choose a stack name for your AWS CloudFormation stack. For example, you can call it **eks-vpc**.
  - **VpcBlock:** Choose a CIDR range for your VPC. You can keep the default value.
  - **Subnet01Block:** Specify a CIDR range for subnet 1. We recommend that you keep the default value so that you have plenty of IP addresses for pods to use.
  - **Subnet02Block:** Specify a CIDR range for subnet 2. We recommend that you keep the default value so that you have plenty of IP addresses for pods to use.

- **Subnet03Block:** Specify a CIDR range for subnet 3. We recommend that you keep the default value so that you have plenty of IP addresses for pods to use.
7. (Optional) On the **Options** page, tag your stack resources. Choose **Next**.
  8. On the **Review** page, choose **Create**.
  9. When your stack is created, select it in the console and choose **Outputs**.
  10. Record the **SecurityGroups** value for the security group that was created. You need this when you create your EKS cluster; this security group is applied to the cross-account elastic network interfaces that are created in your subnets that allow the Amazon EKS control plane to communicate with your worker nodes.
  11. Record the **VpcId** for the VPC that was created. You need this when you launch your worker node group template.
  12. Record the **SubnetIds** for the subnets that were created. You need this when you create your EKS cluster; these are the subnets that your worker nodes are launched into.

## Public and private subnets

### To create your cluster VPC with public and private subnets

1. Open the AWS CloudFormation console at <https://console.aws.amazon.com/cloudformation>.
2. From the navigation bar, select a Region that supports Amazon EKS.
3. Choose **Create stack**.
4. For **Choose a template**, select **Specify an Amazon S3 template URL**.
5. Paste the following URL into the text area and choose **Next**:

```
https://amazon-eks.s3-us-west-2.amazonaws.com/cloudformation/2019-02-11/amazon-eks-vpc-private-subnets.yaml
```

6. On the **Specify Details** page, fill out the parameters accordingly, and then choose **Next**.
  - **Stack name:** Choose a stack name for your AWS CloudFormation stack. For example, you can call it **eks-vpc**.
  - **VpcBlock:** Choose a CIDR range for your VPC. You can keep the default value.
  - **PublicSubnet01Block:** Specify a CIDR range for public subnet 1. We recommend that you keep the default value so that you have plenty of IP addresses for pods to use.
  - **PublicSubnet02Block:** Specify a CIDR range for public subnet 2. We recommend that you keep the default value so that you have plenty of IP addresses for pods to use.
  - **PrivateSubnet01Block:** Specify a CIDR range for private subnet 1. We recommend that you keep the default value so that you have plenty of IP addresses for pods to use.
  - **PrivateSubnet02Block:** Specify a CIDR range for private subnet 2. We recommend that you keep the default value so that you have plenty of IP addresses for pods to use.
7. (Optional) On the **Options** page, tag your stack resources. Choose **Next**.
8. On the **Review** page, choose **Create**.
9. When your stack is created, select it in the console and choose **Outputs**.
10. Record the **SecurityGroups** value for the security group that was created. You need this when you create your EKS cluster; this security group is applied to the cross-account elastic network interfaces that are created in your subnets that allow the Amazon EKS control plane to communicate with your worker nodes.
11. Record the **VpcId** for the VPC that was created. You need this when you launch your worker node group template.
12. Record the **SubnetIds** for the subnets that were created. You need this when you create your EKS cluster; these are the subnets that your worker nodes are launched into.

13. Tag your private subnets so that Kubernetes knows that it can use them for internal load balancers.
  - a. Open the Amazon VPC console at <https://console.aws.amazon.com/vpc/>.
  - b. Choose **Subnets** in the left navigation.
  - c. Select one of the private subnets for your Amazon EKS cluster's VPC (you can filter them with the string `PrivateSubnet`), and choose the **Tags** tab, and then **Add/Edit Tags**.
  - d. Choose **Create Tag** and add the following key and value, and then choose **Save**.

Key	Value
kubernetes.io/role/internal-elb	1

- e. Repeat these substeps for each private subnet in your VPC.

## Next Steps

After you have created your VPC, you can try the [Getting Started with Amazon EKS \(p. 3\)](#) walkthrough, but you can skip the [Create your Amazon EKS Cluster VPC \(p. 9\)](#) section and use these subnets and security groups for your cluster.

## Cluster VPC Considerations

When you create an Amazon EKS cluster, you specify the Amazon VPC subnets for your cluster to use. Amazon EKS requires subnets in at least two Availability Zones. We recommend a network architecture that uses private subnets for your worker nodes and public subnets for Kubernetes to create internet-facing load balancers within.

When you create your cluster, specify all of the subnets that will host resources for your cluster (such as worker nodes and load balancers).

### Note

Internet-facing load balancers require a public subnet in your cluster. Worker nodes also require outbound internet access to the Amazon EKS APIs for cluster introspection and node registration at launch time. To pull container images, they require access to the Amazon S3 and Amazon ECR APIs (and any other container registries, such as DockerHub). For more information, see [Cluster Security Group Considerations \(p. 83\)](#) and [AWS IP Address Ranges](#) in the *AWS General Reference*.

The subnets that you pass when you create the cluster influence where Amazon EKS places elastic network interfaces that are used for the control plane to worker node communication.

It is possible to specify only public or private subnets when you create your cluster, but there are some limitations associated with these configurations:

- **Private-only:** Everything runs in a private subnet and Kubernetes cannot create internet-facing load balancers for your pods.
- **Public-only:** Everything runs in a public subnet, including your worker nodes.

Amazon EKS creates an elastic network interface in your private subnets to facilitate communication to your worker nodes. This communication channel supports Kubernetes functionality such as **kubectl exec** and **kubectl logs**. The security group that you specify when you create your cluster is applied to the elastic network interfaces that are created for your cluster control plane.

Your VPC must have DNS hostname and DNS resolution support. Otherwise, your worker nodes cannot register with your cluster. For more information, see [Using DNS with Your VPC](#) in the *Amazon VPC User Guide*.

## VPC IP Addressing

You can define both private (RFC 1918) and public (non-RFC 1918) CIDR ranges within the VPC used for your Amazon EKS cluster. For more information, see [VPCs and Subnets](#) and [IP Addressing in Your VPC](#) in the *Amazon VPC User Guide*.

### Important

Docker runs in the 172.17.0.0/16 CIDR range in Amazon EKS clusters. We recommend that your cluster's VPC subnets do not overlap this range. Otherwise, you will receive the following error:

```
Error: : error upgrading connection: error dialing backend: dial tcp
172.17.nn.nn:10250: getsockopt: no route to host
```

## VPC Tagging Requirement

When you create your Amazon EKS cluster, Amazon EKS tags the VPC containing the subnets you specify in the following way so that Kubernetes can discover it:

Key	Value
kubernetes.io/cluster/<cluster-name>	shared

- **Key:** The <cluster-name> value matches your Amazon EKS cluster's name.
- **Value:** The shared value allows more than one cluster to use this VPC.

## Subnet Tagging Requirement

When you create your Amazon EKS cluster, Amazon EKS tags the subnets you specify in the following way so that Kubernetes can discover them:

### Note

All subnets (public and private) that your cluster uses for resources should have this tag.

Key	Value
kubernetes.io/cluster/<cluster-name>	shared

- **Key:** The <cluster-name> value matches your Amazon EKS cluster.
- **Value:** The shared value allows more than one cluster to use this subnet.

## Private Subnet Tagging Requirement for Internal Load Balancers

Private subnets in your VPC should be tagged accordingly so that Kubernetes knows that it can use them for internal load balancers:

Key	Value
kubernetes.io/role/internal-elb	1

## Public Subnet Tagging Option for External Load Balancers

Public subnets in your VPC may be tagged accordingly so that Kubernetes knows to use only those subnets for external load balancers, instead of choosing a public subnet in each Availability Zone (in lexicographical order by subnet ID):

Key	Value
kubernetes.io/role/elb	1

## Cluster Security Group Considerations

If you create your VPC and worker node groups with the AWS CloudFormation templates provided in the [Getting Started with Amazon EKS \(p. 3\)](#) walkthrough, then your control plane and worker node security groups are configured with our recommended settings.

The security group for the worker nodes and the security group for the control plane communication to the worker nodes have been set up to prevent communication to privileged ports in the worker nodes. If your applications require added inbound or outbound access from the control plane or worker nodes, you must add these rules to the security groups associated with your cluster. For more information, see [Security Groups for Your VPC](#) in the *Amazon VPC User Guide*.

### Note

To allow proxy functionality on privileged ports or to run the CNCF conformance tests yourself, you must edit the security groups for your control plane and the worker nodes. The security group on the worker nodes' side needs to allow inbound access for ports 0-65535 from the control plane, and the control plane side needs to allow outbound access to the worker nodes on ports 0-65535.

The worker node AWS CloudFormation template modifies the cluster control plane security group when you [launch worker nodes \(p. 56\)](#). **Amazon EKS strongly recommends that you use a dedicated security group for each cluster control plane (one per cluster)**. If you share a cluster control plane security group with other Amazon EKS clusters or resources, you may block or disrupt connections to those resources.

The following tables show the minimum required and recommended security group settings for the control plane and worker node security groups for your cluster:

### Control Plane Security Group

	Protocol	Port Range	Source	Destination
Minimum inbound traffic	TCP	443	All worker node security groups  <b>When <a href="#">cluster endpoint private access (p. 34)</a> is enabled:</b> Any security groups	



	Protocol	Port Range	Source	Destination
			that generate API server client traffic (such as <code>kubectl</code> commands on a bastion host within your cluster's VPC)	
Recommended inbound traffic	TCP	443	All worker node security groups  <b>When <a href="#">cluster endpoint private access (p. 34)</a> is enabled:</b> Any security groups that generate API server client traffic (such as <code>kubectl</code> commands on a bastion host within your cluster's VPC)	
Minimum outbound traffic	TCP	10250		All worker node security groups
Recommended outbound traffic	TCP	1025-65535		All worker node security groups

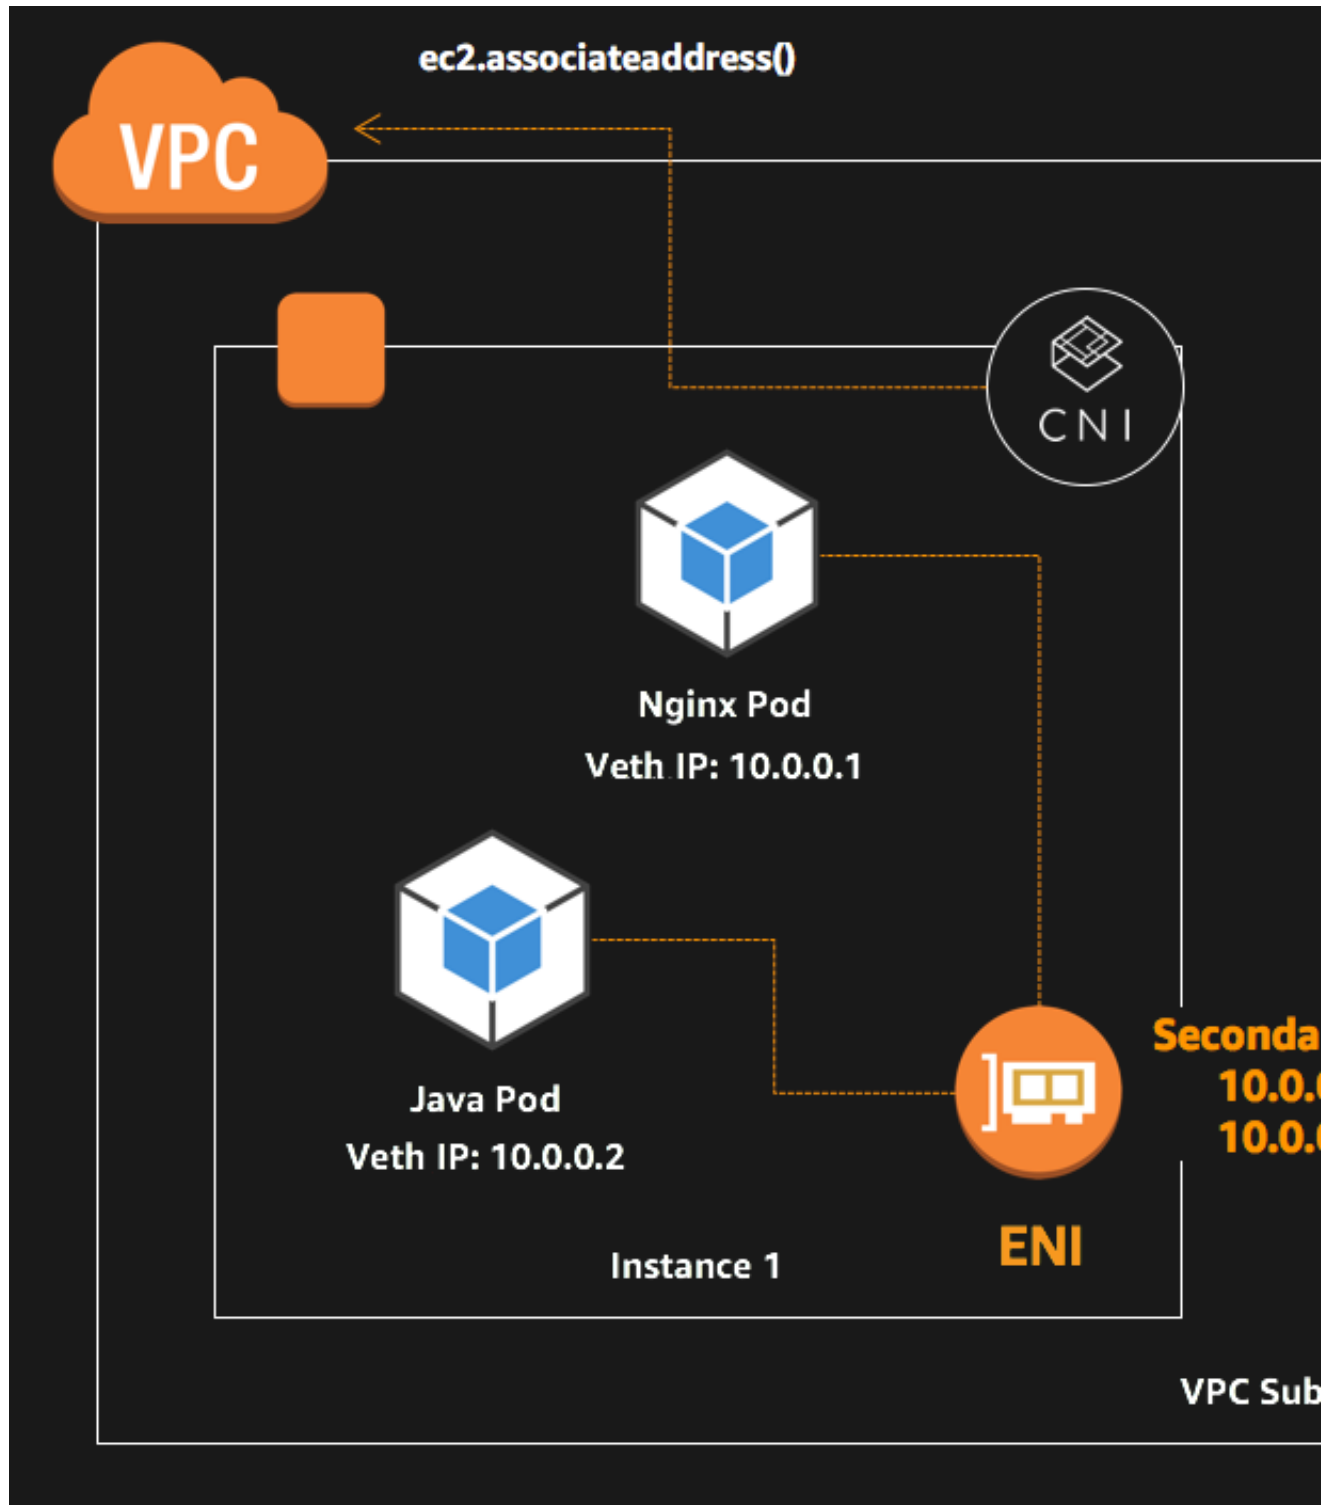
### Worker Node Security Groups

	Protocol	Port Range	Source	Destination
Minimum inbound traffic (from other worker nodes)	Any protocol you expect your worker nodes to use for inter-worker communication	Any ports you expect your worker nodes to use for inter-worker communication	All worker node security groups	
Minimum inbound traffic (from control plane)	TCP	10250	Control plane security group	
Recommended inbound traffic	All TCP	All 443, 1025-65535	All worker node security groups  Control plane security group	
Minimum outbound traffic*	TCP	443		Control plane security group
Recommended outbound traffic	All	All		0.0.0.0/0

\* Worker nodes also require outbound internet access to the Amazon EKS APIs for cluster introspection and node registration at launch time. To pull container images, they require access to the Amazon S3 and Amazon ECR APIs (and any other container registries, such as DockerHub). For more information, see [AWS IP Address Ranges](#) in the *AWS General Reference*.

## Pod Networking

Amazon EKS supports native VPC networking via the Amazon VPC CNI plugin for Kubernetes. Using this CNI plugin allows Kubernetes pods to have the same IP address inside the pod as they do on the VPC network. This CNI plugin is an open-source project that is maintained on [GitHub](#).



The CNI plugin is responsible for allocating VPC IP addresses to Kubernetes nodes and configuring the necessary networking for pods on each node. The plugin consists of two primary components:

- The L-IPAM daemon is responsible for attaching elastic network interfaces to instances, assigning secondary IP addresses to elastic network interfaces, and maintaining a "warm pool" of IP addresses on each node for assignment to Kubernetes pods when they are scheduled.
- The CNI plugin itself is responsible for wiring the host network (for example, configuring the interfaces and virtual Ethernet pairs) and adding the correct interface to the pod namespace.

For more information about the design and networking configuration, see [CNI plugin for Kubernetes networking over AWS VPC](#).

Elastic network interface and secondary IP address limitations by Amazon EC2 instance types are applicable. In general, larger instances can support more IP addresses. For more information, see [IP Addresses Per Network Interface Per Instance Type](#) in the *Amazon EC2 User Guide for Linux Instances*.

## CNI Configuration Variables

The Amazon VPC CNI plugin for Kubernetes supports a number of configuration options, which are set through environment variables. The following environment variables are available, and all of them are optional.

`AWS_VPC_CNI_NODE_PORT_SUPPORT`

Type: Boolean

Default: `true`

Specifies whether `NodePort` services are enabled on a worker node's primary network interface. This requires additional `iptables` rules and that the kernel's reverse path filter on the primary interface is set to `loose`.

`AWS_VPC_K8S_CNI_CUSTOM_NETWORK_CFG`

Type: Boolean

Default: `false`

Specifies that your pods may use subnets and security groups (within the same VPC as your control plane resources) that are independent of your cluster's `resourcesVpcConfig`. By default, pods share the same subnet and security groups as the worker node's primary interface. Setting this variable to `true` causes `ipamD` to use the security groups and subnets in a worker node's `ENIConfig` for elastic network interface allocation. You must create an `ENIConfig` custom resource definition for each subnet that your pods will reside in, and then annotate each worker node to use a specific `ENIConfig` (multiple worker nodes can be annotated with the same `ENIConfig`). Worker nodes can only be annotated with a single `ENIConfig` at a time, and the subnet in the `ENIConfig` must belong to the same Availability Zone that the worker node resides in. For more information, see [CNI Custom Networking \(p. 94\)](#).

`AWS_VPC_K8S_CNI_EXTERNALSNAT`

Type: Boolean

Default: `false`

Specifies whether an external NAT gateway should be used to provide SNAT of secondary ENI IP addresses. If set to `true`, the SNAT `iptables` rule and off-VPC IP rule are not applied, and these rules are removed if they have already been applied.

Disable SNAT if you need to allow inbound communication to your pods from external VPNs, direct connections, and external VPCs, and your pods do not need to access the Internet directly via an

Internet Gateway. However, your nodes must be running in a private subnet and connected to the internet through an AWS NAT Gateway or another external NAT device.

For more information, see [External Source Network Address Translation \(SNAT\) \(p. 91\)](#).

`WARM_ENI_TARGET`

Type: Integer

Default: 1

Specifies the number of free elastic network interfaces (and all of their available IP addresses) that the `ipamD` daemon should attempt to keep available for pod assignment on the node. By default, `ipamD` attempts to keep 1 elastic network interface and all of its IP addresses available for pod assignment.

**Note**

The number of IP addresses per network interface varies by instance type. For more information, see [IP Addresses Per Network Interface Per Instance Type](#) in the *Amazon EC2 User Guide for Linux Instances*.

For example, an `m4.4xlarge` launches with 1 network interface and 30 IP addresses. If 5 pods are placed on the node and 5 free IP addresses are removed from the IP address warm pool, then `ipamD` attempts to allocate more interfaces until `WARM_ENI_TARGET` free interfaces are available on the node.

**Note**

If `WARM_IP_TARGET` is set, then this environment variable is ignored and the `WARM_IP_TARGET` behavior is used instead.

`WARM_IP_TARGET`

Type: Integer

Default: None

Specifies the number of free IP addresses that the `ipamD` daemon should attempt to keep available for pod assignment on the node. For example, if `WARM_IP_TARGET` is set to 10, then `ipamD` attempts to keep 10 free IP addresses available at all times. If the elastic network interfaces on the node are unable to provide these free addresses, `ipamD` attempts to allocate more interfaces until `WARM_IP_TARGET` free IP addresses are available.

**Note**

This environment variable overrides `WARM_ENI_TARGET` behavior.

## Installing CoreDNS

Clusters that were created with Kubernetes version 1.10 shipped with `kube-dns` as the default DNS and service discovery provider. If you have updated from a 1.10 cluster and you want to use CoreDNS for DNS and service discovery, you must install CoreDNS and remove `kube-dns`.

To check if your cluster is already running CoreDNS, use the following command.

```
kubectl get pod -n kube-system -l k8s-app=kube-dns
```

If the output shows `coredns` in the pod names, you're already running CoreDNS in your cluster. If not, use the following procedure to update your DNS and service discovery provider to CoreDNS.

### Note

The service for CoreDNS is still called `kube-dns` for backward compatibility.

Choose the tab below that corresponds to your desired CoreDNS installation method:

eksctl

### To install CoreDNS on an updated Amazon EKS cluster with eksctl

This procedure assumes that you have installed `eksctl`, and that your `eksctl` version is at least 0.1.37. You can check your version with the following command:

```
eksctl version
```

For more information on installing or upgrading `eksctl`, see [Installing or Upgrading eksctl](#) (p. 119).

1. Run the following command to install `coredns`, replacing the red text with your cluster name:

```
eksctl utils install-coredns --name dev --approve
```

Output:

```
[#] using region us-west-2
[#] created "kube-system:ServiceAccount/coredns"
[#] created "ClusterRole.rbac.authorization.k8s.io/system:coredns"
[#] created "ClusterRoleBinding.rbac.authorization.k8s.io/system:coredns"
[#] created "kube-system:ConfigMap/coredns"
[#] created "kube-system:Deployment.extensions/coredns"
[#] replaced "kube-system:Service/kube-dns"
[#] waiting for 2 of "coredns" pods to become ready
[#] deleted "kube-dns"
[#] "coredns" is now up-to-date
```

2. Check the current version of your cluster's `coredns` deployment.

```
kubectl describe deployment coredns --namespace kube-system | grep Image | cut -d
"/" -f 3
```

Output:

```
coredns:v1.1.3
```

The recommended `coredns` versions for their corresponding Kubernetes versions are as follows:

- **Kubernetes 1.13:** 1.2.6
- **Kubernetes 1.12:** 1.2.2
- **Kubernetes 1.11:** 1.1.3

If your current `coredns` version doesn't match the recommendation for your cluster version, update the `coredns` deployment to use the recommended image with the following command, replacing the red text with your cluster name:

```
eksctl utils update-coredns --name dev --approve
```

kubectl

### To install CoreDNS on an updated Amazon EKS cluster with kubectl

1. Add the `{"eks.amazonaws.com/component": "kube-dns"}` selector to the kube-dns deployment for your cluster. This prevents the two DNS deployments from competing for control of the same set of labels.

```
kubectl patch -n kube-system deployment/kube-dns --patch \
'{"spec":{"selector":{"matchLabels":{"eks.amazonaws.com/component":"kube-dns"}}}}'
```

2. Deploy CoreDNS to your cluster.
  - a. Set your cluster's DNS IP address to the `DNS_CLUSTER_IP` environment variable.

```
export DNS_CLUSTER_IP=$(kubectl get svc -n kube-system kube-dns -o
jsonpath='{.spec.clusterIP}')
```

- b. Set your cluster's AWS Region to the `REGION` environment variable.

```
export REGION="us-west-2"
```

- c. Download the CoreDNS manifest from the Amazon EKS resource bucket.

```
curl -o dns.yaml https://amazon-eks.s3-us-west-2.amazonaws.com/
cloudformation/2019-02-11/dns.yaml
```

- d. Replace the variable placeholders in the `dns.yaml` file with your environment variable values and apply the updated manifest to your cluster. The following command completes this in one step.

```
cat dns.yaml | sed -e "s/REGION/$REGION/g" | sed -e "s/DNS_CLUSTER_IP/
$DNS_CLUSTER_IP/g" | kubectl apply -f -
```

- e. Fetch the coredns pod name from your cluster.

```
COREDNS_POD=$(kubectl get pod -n kube-system -l eks.amazonaws.com/
component=coredns \
-o jsonpath='{.items[0].metadata.name}')
```

- f. Query the coredns pod to ensure that it's receiving requests.

```
kubectl get --raw /api/v1/namespaces/kube-system/pods/$COREDNS_POD:9153/proxy/
metrics \
| grep 'coredns_dns_request_count_total'
```

#### Note

It might take several minutes for the expected output to return properly, depending on the rate of DNS requests in your cluster.

Expected output (the number in red is the DNS request count total):

```
# HELP coredns_dns_request_count_total Counter of DNS requests made per zone,
protocol and family.
# TYPE coredns_dns_request_count_total counter
coredns_dns_request_count_total{family="1",proto="udp",server="dns://:53",zone="."} 23
```

3. Scale down the kube-dns deployment to zero replicas.

```
kubectl scale -n kube-system deployment/kube-dns --replicas=0
```

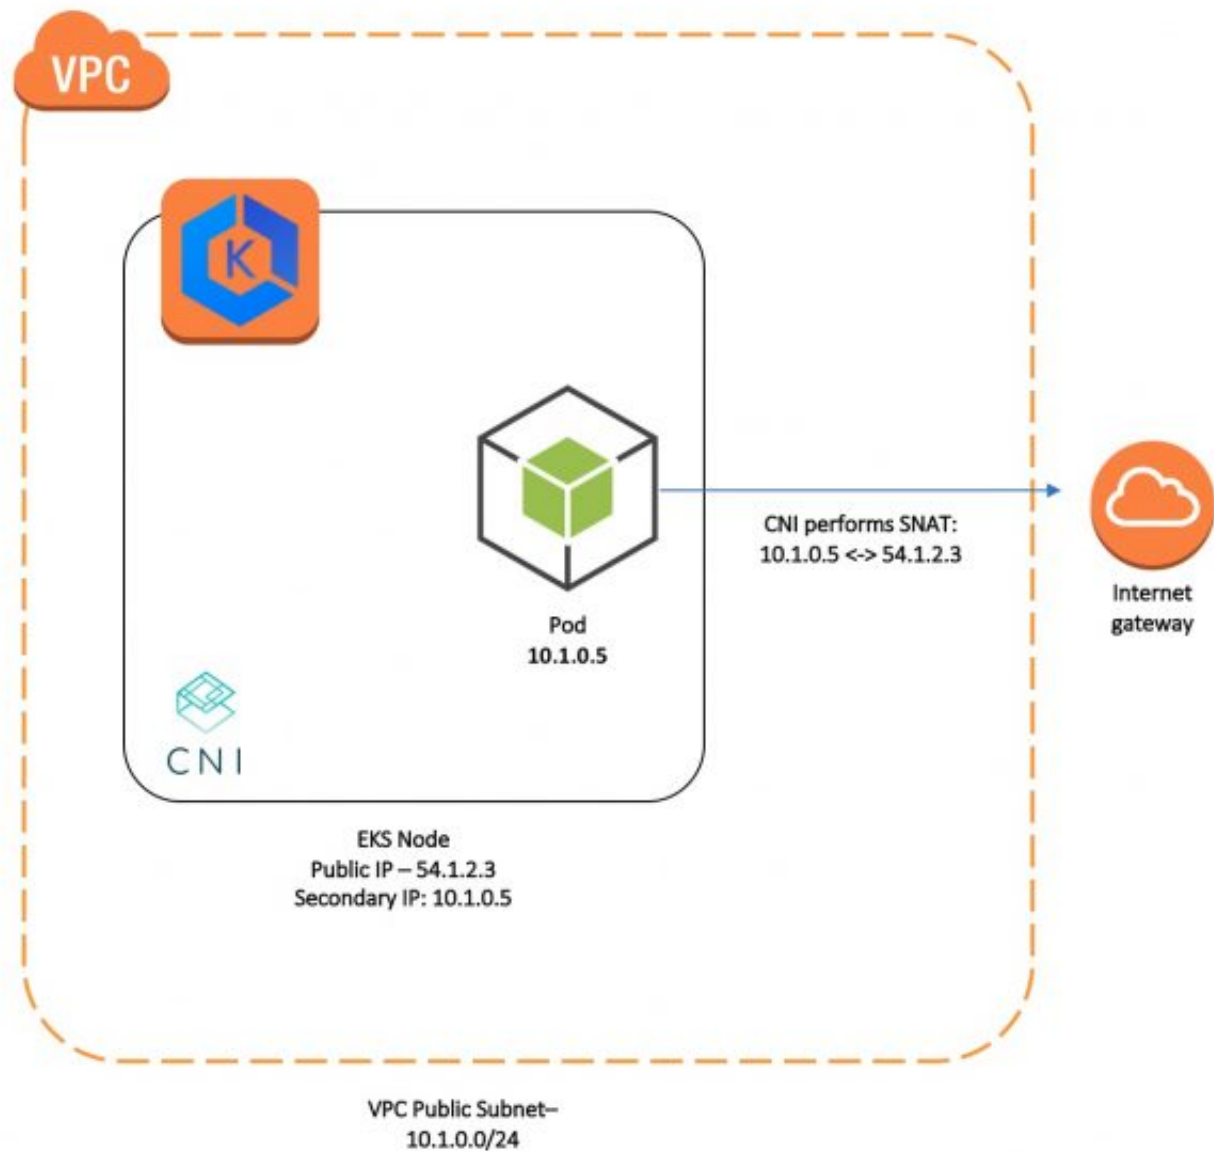
4. Clean up the old kube-dns resources.

```
kubectl delete -n kube-system deployment/kube-dns serviceaccount/kube-dns  
configmap/kube-dns
```

## External Source Network Address Translation (SNAT)

By default, the [Amazon VPC CNI plugin for Kubernetes](#) configures pods with source network address translation (SNAT) enabled. This sets the return address for a packet to the primary public IP of the instance and allows for communication with the internet. In this default configuration, when you use an internet gateway and a public address, the return packet is routed to the correct Amazon EC2 instance.

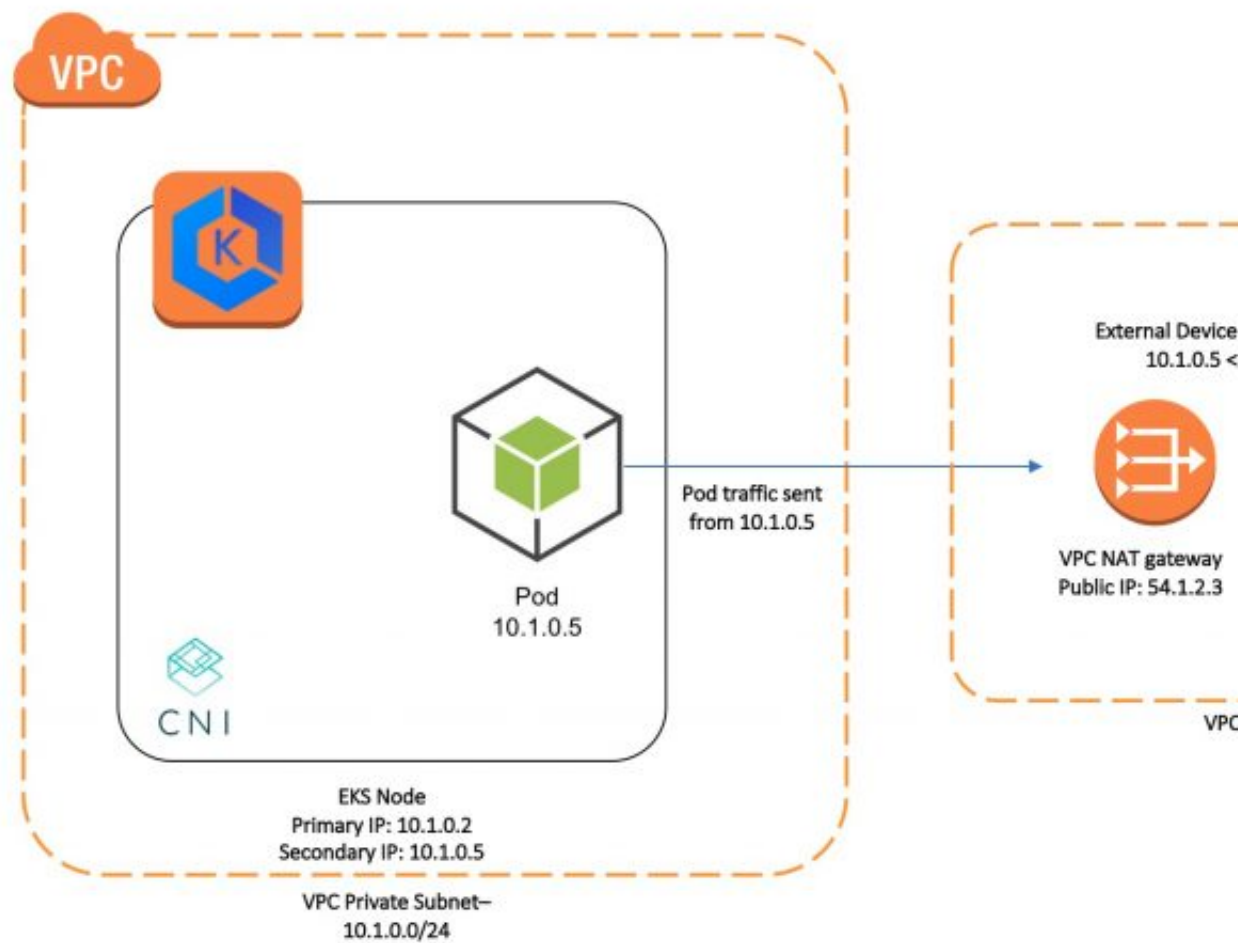




However, SNAT can cause issues if traffic from another private IP space (for example, [VPC peering](#), [Transit VPC](#), or [Direct Connect](#)) attempts to communicate directly to a pod that is not attached to the primary elastic network interface of the Amazon EC2 instance. To specify that NAT be handled by an external device (such as a NAT gateway, and not on the instance itself), you can disable SNAT on the instance by setting the `AWS_VPC_K8S_CNI_EXTERNALSNAT` environment variable to `true`. Disable SNAT to allow inbound communication to your pods from external VPNs, direct connections, and external VPCs, and your pods do not need to access the internet directly via an internet gateway.

**Note**

SNAT is required for nodes that reside in a public subnet. To use external SNAT, your nodes must reside in a private subnet and connect to the internet through a NAT gateway or another external NAT device.



### To disable SNAT on your worker nodes

1. Edit the `aws-node` daemonset:

```
kubectl edit daemonset -n kube-system aws-node
```

2. Add the `AWS_VPC_K8S_CNI_EXTERNALSNAT` environment variable to the node container spec and set it to `true`:

```
...
spec:
  containers:
  - env:
    - name: AWS_VPC_K8S_CNI_EXTERNALSNAT
      value: "true"
```

```
- name: AWS_VPC_K8S_CNI_LOGLEVEL
  value: DEBUG
- name: MY_NODE_NAME
...

```

3. Save the file and exit your text editor.

## CNI Custom Networking

By default, when new network interfaces are allocated for pods, [ipamD](#) uses the worker node's primary elastic network interface's security groups and subnet. However, there are use cases where your pod network interfaces should use a different security group or subnet, within the same VPC as your control plane security group. For example:

- There are a limited number of IP addresses available in a subnet. This limits the number of pods can be created in the cluster. Using different subnets for pod groups allows you to increase the number of available IP addresses.
- For security reasons, your pods must use different security groups or subnets than the node's primary network interface.
- The worker nodes are configured in public subnets and you want the pods to be placed in private subnets using a NAT Gateway. For more information, see [External Source Network Address Translation \(SNAT\)](#) (p. 91).

### Note

The use cases discussed in this topic require [Amazon VPC CNI plugin for Kubernetes](#) version 1.4.0 or later. To check your CNI version, and upgrade if necessary, see [Amazon VPC CNI Plugin for Kubernetes Upgrades](#) (p. 97).

Enabling this feature effectively removes an available elastic network interface (and all of its available IP addresses for pods) from each worker node that uses it. The primary network interface for the worker node is not used for pod placement when this feature is enabled. You should choose larger instance types with more available elastic network interfaces if you choose to enable this feature.

### To configure CNI custom networking

1. Associate a secondary CIDR block to your cluster's VPC. For more information, see [Associating a Secondary IPv4 CIDR Block with Your VPC](#) in the *Amazon VPC User Guide*.
2. Create a subnet in your VPC for each Availability Zone, using your secondary CIDR block. Your custom subnets must be from a different VPC CIDR block than the subnet that your worker nodes were launched into. For more information, see [Creating a Subnet in Your VPC](#) in the *Amazon VPC User Guide*.
3. Edit the `aws-node` daemonset for your cluster:

```
kubectl edit daemonset -n kube-system aws-node
```

4. Add the `AWS_VPC_K8S_CNI_CUSTOM_NETWORK_CFG` environment variable to the node container spec and set it to `true`:

```
...
spec:
  containers:
  - env:
    - name: AWS_VPC_K8S_CNI_CUSTOM_NETWORK_CFG
      value: "true"
    - name: AWS_VPC_K8S_CNI_LOGLEVEL
```

```

        value: DEBUG
      - name: MY_NODE_NAME
    ...

```

5. Save the file and exit your text editor.
6. Define a new `ENIConfig` custom resource for your cluster.
  - a. Create a file called `ENIConfig.yaml` and paste the following content into it:

```

apiVersion: apiextensions.k8s.io/v1beta1
kind: CustomResourceDefinition
metadata:
  name: eniconfigs.crd.k8s.amazonaws.com
spec:
  scope: Cluster
  group: crd.k8s.amazonaws.com
  version: v1alpha1
  names:
    plural: eniconfigs
    singular: eniconfig
    kind: ENIConfig

```

- b. Apply the file to your cluster with the following command:

```
kubectl apply -f ENIConfig.yaml
```

7. Create an `ENIConfig` custom resource for each subnet that you want to schedule pods in.
  - a. Create a unique file for each elastic network interface configuration to use with the following information. Replacing the subnet and security group IDs with your own values. If you don't have a specific security group that you want to attach for your pods, you can leave that value empty for now. Later, you will specify the worker node security group in the `ENIConfig`.

For this example, the file is called `custom-pod-netconfig.yaml`.

#### Note

Each subnet and security group combination requires its own custom resource.

```

apiVersion: crd.k8s.amazonaws.com/v1alpha1
kind: ENIConfig
metadata:
  name: custom-pod-netconfig
spec:
  securityGroups:
    - sg-0dff363a7d37c3c61
  subnet: subnet-017b472c2f79fdf96

```

- b. Apply each custom resource file that you created earlier to your cluster with the following command:

```
kubectl apply -f custom-pod-netconfig.yaml
```

8. Create a new worker node group for each `ENIConfig` that you configured, and limit the Auto Scaling group to the same Availability Zone as the `ENIConfig`.

Follow the steps in [Launching Amazon EKS Worker Nodes \(p. 56\)](#) to create each new worker node group. When you create each group, apply the `k8s.amazonaws.com/eniConfig` label to the node group, and set the value to the name of the `ENIConfig` to use for that worker node group.

- If you use `eksctl` to create your worker node groups, add the following flag to your create cluster command:

```
--node-labels k8s.amazonaws.com/eniConfig=custom-pod-netconfig
```

- If you use the Amazon EKS-provided AWS CloudFormation templates to create your worker node groups, add the following option to the **BootstrapArguments** field in the AWS CloudFormation console:

```
--kubelet-extra-args '--node-labels=k8s.amazonaws.com/eniConfig=custom-pod-netconfig'
```

9. After your worker node groups are created, record the security group that was created for each worker node group and apply it to its associated ENIConfig. Edit each ENIConfig with the following command, replacing the red text with your value):

```
kubectl edit eniconfig.crd.k8s.amazonaws.com/custom-pod-netconfig
```

The spec section should look like this:

```
spec:
  securityGroups:
  - sg-08052d900a2c7fb0a
  subnet: subnet-017b472c2f79fdf96
```

10. If you have any worker nodes in your cluster that had pods placed on them before you completed this procedure, you should terminate them. Only new nodes that are registered with the k8s.amazonaws.com/eniConfig label will use the new custom networking feature.

### To automatically apply an ENIConfig to a node based on its Availability Zone

- By default, Kubernetes applies the availability zone of a node to the failure-domain.beta.kubernetes.io/zone label. You can name your ENIConfig custom resources after each Availability Zone in your VPC, and then specify this label as the value of the ENI\_CONFIG\_LABEL\_DEF environment variable in the aws-node container spec for your worker nodes.

```
...
spec:
  containers:
  - env:
    - name: AWS_VPC_K8S_CNI_CUSTOM_NETWORK_CFG
      value: "true"
    - name: ENI_CONFIG_LABEL_DEF
      value: failure-domain.beta.kubernetes.io/zone
    - name: AWS_VPC_K8S_CNI_LOGLEVEL
      value: DEBUG
    - name: MY_NODE_NAME
  ...
```

For example, if subnet-0c4678ec01ce68b24 is in the us-east-1a Availability Zone, you could use the following ENIConfig for that Availability Zone by naming it us-east-1a:

```
apiVersion: crd.k8s.amazonaws.com/v1alpha1
kind: ENIConfig
metadata:
  name: us-east-1a
spec:
  securityGroups:
  - sg-08052d900a2c7fb0a
```

```
subnet: subnet-0c4678ec01ce68b24
```

## Amazon VPC CNI Plugin for Kubernetes Upgrades

When you launch an Amazon EKS cluster, we apply a recent version of the [Amazon VPC CNI plugin for Kubernetes](#) to your cluster (the absolute latest version of the plugin is available [on GitHub](#) for a short grace period before new clusters are switched over to use it). However, Amazon EKS does not automatically upgrade the CNI plugin on your cluster when new versions are released. You must upgrade the CNI plugin manually to get the latest version on existing clusters.

The current default CNI version for new clusters is 1.5.2.

The latest CNI version available [on GitHub](#) is 1.5.2. You can view the different releases available for the plugin, and read the release notes for each version [on GitHub](#).

Use the following procedures to check your CNI version and upgrade to the latest version.

### To check your Amazon VPC CNI Plugin for Kubernetes version

- Use the following command to print your cluster's CNI version:

```
kubectl describe daemonset aws-node --namespace kube-system | grep Image | cut -d "/" -f 2
```

Output:

```
amazon-k8s-cni:1.4.1
```

In this example output, the CNI version is 1.4.1, which is earlier than the current version, 1.5.2. Use the following procedure to upgrade the CNI.

### To upgrade the Amazon VPC CNI Plugin for Kubernetes

- Use the following command to upgrade your CNI version to the latest version:
  - For Kubernetes 1.10 clusters:

```
kubectl apply -f https://raw.githubusercontent.com/aws/amazon-vpc-cni-k8s/release-1.5/config/v1.5/aws-k8s-cni-1.10.yaml
```

- For all other Kubernetes versions:

```
kubectl apply -f https://raw.githubusercontent.com/aws/amazon-vpc-cni-k8s/release-1.5/config/v1.5/aws-k8s-cni.yaml
```

## Installing Calico on Amazon EKS

[Project Calico](#) is a network policy engine for Kubernetes. With Calico network policy enforcement, you can implement network segmentation and tenant isolation. This is useful in multi-tenant environments where you must isolate tenants from each other or when you want to create separate environments for development, staging, and production. Network policies are similar to AWS security groups in that you can create network ingress and egress rules. Instead of assigning instances to a security group, you

assign network policies to pods using pod selectors and labels. The following procedure shows you how to install Calico on your Amazon EKS cluster.

### To install Calico on your Amazon EKS cluster

1. Apply the Calico manifest from the [aws/amazon-vpc-cni-k8s GitHub project](https://github.com/aws/amazon-vpc-cni-k8s). This manifest creates DaemonSets in the kube-system namespace.

```
kubectl apply -f https://raw.githubusercontent.com/aws/amazon-vpc-cni-k8s/release-1.5/config/v1.5/calico.yaml
```

2. Watch the kube-system DaemonSets and wait for the calico-node DaemonSet to have the DESIRED number of pods in the READY state. When this happens, Calico is working.

```
kubectl get daemonset calico-node --namespace kube-system
```

Output:

NAME	DESIRED	CURRENT	READY	UP-TO-DATE	AVAILABLE	NODE SELECTOR
calico-node	3	3	3	3	3	<none>
AGE						
38s						

### To delete Calico from your Amazon EKS cluster

- If you are done using Calico in your Amazon EKS cluster, you can delete the DaemonSet with the following command:

```
kubectl delete -f https://raw.githubusercontent.com/aws/amazon-vpc-cni-k8s/release-1.5/config/v1.5/calico.yaml
```

## Stars Policy Demo

This section walks through the [Stars Policy Demo](#) provided by the Project Calico documentation. The demo creates a frontend, backend, and client service on your Amazon EKS cluster. The demo also creates a management GUI that shows the available ingress and egress paths between each service.

Before you create any network policies, all services can communicate bidirectionally. After you apply the network policies, you can see that the client can only communicate with the frontend service, and the backend can only communicate with the frontend.

### To run the Stars Policy demo

1. Apply the frontend, backend, client, and management UI services:

```
kubectl apply -f https://docs.projectcalico.org/v3.3/getting-started/kubernetes/tutorials/stars-policy/manifests/00-namespace.yaml
kubectl apply -f https://docs.projectcalico.org/v3.3/getting-started/kubernetes/tutorials/stars-policy/manifests/01-management-ui.yaml
kubectl apply -f https://docs.projectcalico.org/v3.3/getting-started/kubernetes/tutorials/stars-policy/manifests/02-backend.yaml
kubectl apply -f https://docs.projectcalico.org/v3.3/getting-started/kubernetes/tutorials/stars-policy/manifests/03-frontend.yaml
kubectl apply -f https://docs.projectcalico.org/v3.3/getting-started/kubernetes/tutorials/stars-policy/manifests/04-client.yaml
```


2. Wait for all of the pods to reach the Running status:

```
kubectl get pods --all-namespaces --watch
```

3. To connect to the management UI, forward your local port 9001 to the management-ui service running on your cluster:

```
kubectl port-forward service/management-ui -n management-ui 9001
```

4. Open a browser on your local system and point it to <http://localhost:9001/>. You should see the management UI. The **C** node is the client service, the **F** node is the frontend service, and the **B** node is the backend service. Each node has full communication access to all other nodes (as indicated by the bold, colored lines).





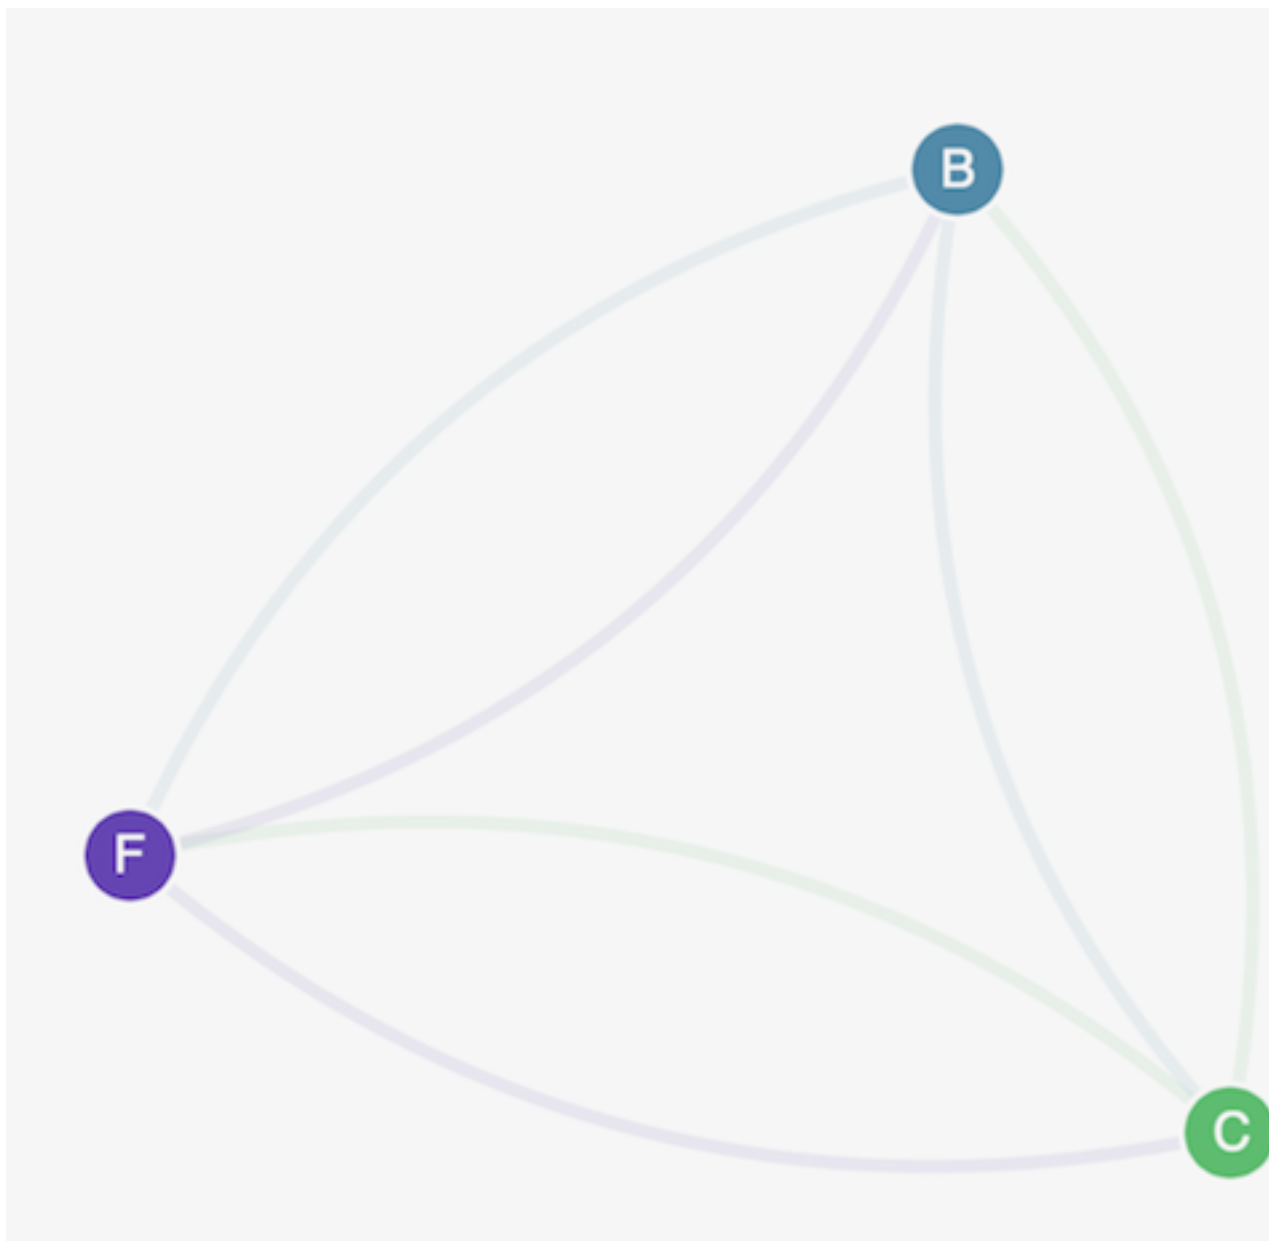
5. Apply the following network policies to isolate the services from each other:

```
kubectl apply -n stars -f https://docs.projectcalico.org/v3.3/getting-started/  
kubernetes/tutorials/stars-policy/policies/default-deny.yaml  
kubectl apply -n client -f https://docs.projectcalico.org/v3.3/getting-started/  
kubernetes/tutorials/stars-policy/policies/default-deny.yaml
```

6. Refresh your browser. You see that the management UI can no longer reach any of the nodes, so they don't show up in the UI.
7. Apply the following network policies to allow the management UI to access the services:

```
kubectl apply -f https://docs.projectcalico.org/v3.3/getting-started/kubernetes/  
tutorials/stars-policy/policies/allow-ui.yaml  
kubectl apply -f https://docs.projectcalico.org/v3.3/getting-started/kubernetes/  
tutorials/stars-policy/policies/allow-ui-client.yaml
```

8. Refresh your browser. You see that the management UI can reach the nodes again, but the nodes cannot communicate with each other.




9. Apply the following network policy to allow traffic from the frontend service to the backend service:

```
kubectl apply -f https://docs.projectcalico.org/v3.3/getting-started/kubernetes/tutorials/stars-policy/policies/backend-policy.yaml
```

10. Apply the following network policy to allow traffic from the `client` namespace to the frontend service:

```
kubectl apply -f https://docs.projectcalico.org/v3.3/getting-started/kubernetes/tutorials/stars-policy/policies/frontend-policy.yaml
```




11. (Optional) When you are done with the demo, you can delete its resources with the following commands:

```
kubectl delete -f https://docs.projectcalico.org/v3.3/getting-started/kubernetes/tutorials/stars-policy/manifests/04-client.yaml
kubectl delete -f https://docs.projectcalico.org/v3.3/getting-started/kubernetes/tutorials/stars-policy/manifests/03-frontend.yaml
kubectl delete -f https://docs.projectcalico.org/v3.3/getting-started/kubernetes/tutorials/stars-policy/manifests/02-backend.yaml
kubectl delete -f https://docs.projectcalico.org/v3.3/getting-started/kubernetes/tutorials/stars-policy/manifests/01-management-ui.yaml
```

```
kubectl delete -f https://docs.projectcalico.org/v3.3/getting-started/kubernetes/  
tutorials/stars-policy/manifests/00-namespace.yaml
```

# Managing Cluster Authentication

Amazon EKS uses IAM to provide authentication to your Kubernetes cluster (through the `aws eks get-token` command, available in version 1.16.156 or greater of the AWS CLI, or the [AWS IAM Authenticator for Kubernetes](#)), but it still relies on native Kubernetes [Role Based Access Control](#) (RBAC) for authorization. This means that IAM is only used for authentication of valid IAM entities. All permissions for interacting with your Amazon EKS cluster's Kubernetes API is managed through the native Kubernetes RBAC system.



## Topics

- [Installing `kubectl` \(p. 104\)](#)
- [Installing `aws-iam-authenticator` \(p. 108\)](#)
- [Create a `kubeconfig` for Amazon EKS \(p. 111\)](#)
- [Managing Users or IAM Roles for your Cluster \(p. 115\)](#)

## Installing `kubectl`

Kubernetes uses a command line utility called `kubectl` for communicating with the cluster API server. The `kubectl` binary is available in many operating system package managers, and this option is often much easier than a manual download and install process. You can follow the instructions for your specific operating system or package manager in the [Kubernetes documentation](#) to install.

This topic helps you to download and install the Amazon EKS-vended **kubectl** binaries for macOS, Linux, and Windows operating systems. These binaries are identical to the upstream community versions, and are not unique to Amazon EKS or AWS.

### Note

You must use a **kubectl** version that is within one minor version difference of your Amazon EKS cluster control plane. For example, a 1.12 **kubectl** client should work with Kubernetes 1.11, 1.12, and 1.13 clusters.

macOS

### To install kubectl on macOS

1. Download the Amazon EKS-vended **kubectl** binary for your cluster's Kubernetes version from Amazon S3:

- **Kubernetes 1.13:**

```
curl -o kubectl https://amazon-eks.s3-us-west-2.amazonaws.com/1.13.7/2019-06-11/bin/darwin/amd64/kubectl
```

- **Kubernetes 1.12:**

```
curl -o kubectl https://amazon-eks.s3-us-west-2.amazonaws.com/1.12.9/2019-06-21/bin/darwin/amd64/kubectl
```

- **Kubernetes 1.11:**

```
curl -o kubectl https://amazon-eks.s3-us-west-2.amazonaws.com/1.11.10/2019-06-21/bin/darwin/amd64/kubectl
```

2. (Optional) Verify the downloaded binary with the SHA-256 sum for your binary.
  - a. Download the SHA-256 sum for your cluster's Kubernetes version for macOS:

- **Kubernetes 1.13:**

```
curl -o kubectl.sha256 https://amazon-eks.s3-us-west-2.amazonaws.com/1.13.7/2019-06-11/bin/darwin/amd64/kubectl.sha256
```

- **Kubernetes 1.12:**

```
curl -o kubectl.sha256 https://amazon-eks.s3-us-west-2.amazonaws.com/1.12.9/2019-06-21/bin/darwin/amd64/kubectl.sha256
```

- **Kubernetes 1.11:**

```
curl -o kubectl.sha256 https://amazon-eks.s3-us-west-2.amazonaws.com/1.11.10/2019-06-21/bin/darwin/amd64/kubectl.sha256
```

- b. Check the SHA-256 sum for your downloaded binary.

```
openssl sha1 -sha256 kubectl
```

- c. Compare the generated SHA-256 sum in the command output against your downloaded SHA-256 file. The two should match.
3. Apply execute permissions to the binary.

```
chmod +x ./kubectl
```

4. Copy the binary to a folder in your `PATH`. If you have already installed a version of **kubectl**, then we recommend creating a `$HOME/bin/kubectl` and ensuring that `$HOME/bin` comes first in your `$PATH`.

```
mkdir -p $HOME/bin && cp ./kubectl $HOME/bin/kubectl && export PATH=$HOME/bin:$PATH
```

5. (Optional) Add the `$HOME/bin` path to your shell initialization file so that it is configured when you open a shell.

```
echo 'export PATH=$HOME/bin:$PATH' >> ~/.bash_profile
```

6. After you install **kubectl**, you can verify its version with the following command:

```
kubectl version --short --client
```

## Linux

### To install kubectl on Linux

1. Download the Amazon EKS-vended **kubectl** binary for your cluster's Kubernetes version from Amazon S3:

- **Kubernetes 1.13:**

```
curl -o kubectl https://amazon-eks.s3-us-west-2.amazonaws.com/1.13.7/2019-06-11/bin/linux/amd64/kubectl
```

- **Kubernetes 1.12:**

```
curl -o kubectl https://amazon-eks.s3-us-west-2.amazonaws.com/1.12.9/2019-06-21/bin/linux/amd64/kubectl
```

- **Kubernetes 1.11:**

```
curl -o kubectl https://amazon-eks.s3-us-west-2.amazonaws.com/1.11.10/2019-06-21/bin/linux/amd64/kubectl
```

2. (Optional) Verify the downloaded binary with the SHA-256 sum for your binary.
  - a. Download the SHA-256 sum for your cluster's Kubernetes version for Linux:

- **Kubernetes 1.13:**

```
curl -o kubectl.sha256 https://amazon-eks.s3-us-west-2.amazonaws.com/1.13.7/2019-06-11/bin/linux/amd64/kubectl.sha256
```

- **Kubernetes 1.12:**

```
curl -o kubectl.sha256 https://amazon-eks.s3-us-west-2.amazonaws.com/1.12.9/2019-06-21/bin/linux/amd64/kubectl.sha256
```

- **Kubernetes 1.11:**

```
curl -o kubectl.sha256 https://amazon-eks.s3-us-west-2.amazonaws.com/1.11.10/2019-06-21/bin/linux/amd64/kubectl.sha256
```

- b. Check the SHA-256 sum for your downloaded binary.

```
openssl sha1 -sha256 kubectl
```

- c. Compare the generated SHA-256 sum in the command output against your downloaded SHA-256 file. The two should match.
3. Apply execute permissions to the binary.

```
chmod +x ./kubectl
```

4. Copy the binary to a folder in your `PATH`. If you have already installed a version of **kubectl**, then we recommend creating a `$HOME/bin/kubectl` and ensuring that `$HOME/bin` comes first in your `$PATH`.

```
mkdir -p $HOME/bin && cp ./kubectl $HOME/bin/kubectl && export PATH=$HOME/bin:$PATH
```

5. (Optional) Add the `$HOME/bin` path to your shell initialization file so that it is configured when you open a shell.

#### Note

This step assumes you are using the Bash shell; if you are using another shell, change the command to use your specific shell initialization file.

```
echo 'export PATH=$HOME/bin:$PATH' >> ~/.bashrc
```

6. After you install **kubectl**, you can verify its version with the following command:

```
kubectl version --short --client
```

## Windows

### To install kubectl on Windows

1. Open a PowerShell terminal.
2. Download the Amazon EKS-vended **kubectl** binary for your cluster's Kubernetes version from Amazon S3:

- **Kubernetes 1.13:**

```
curl -o kubectl.exe https://amazon-eks.s3-us-west-2.amazonaws.com/1.13.7/2019-06-11/bin/windows/amd64/kubectl.exe
```

- **Kubernetes 1.12:**

```
curl -o kubectl.exe https://amazon-eks.s3-us-west-2.amazonaws.com/1.12.9/2019-06-21/bin/windows/amd64/kubectl.exe
```

- **Kubernetes 1.11:**

```
curl -o kubectl.exe https://amazon-eks.s3-us-west-2.amazonaws.com/1.11.10/2019-06-21/bin/windows/amd64/kubectl.exe
```



3. (Optional) Verify the downloaded binary with the SHA-256 sum for your binary.
  - a. Download the SHA-256 sum for your cluster's Kubernetes version for Windows:

- **Kubernetes 1.13:**

```
curl -o kubectl.exe.sha256 https://amazon-eks.s3-us-west-2.amazonaws.com/1.13.7/2019-06-11/bin/windows/amd64/kubectl.exe.sha256
```

- **Kubernetes 1.12:**

```
curl -o kubectl.exe.sha256 https://amazon-eks.s3-us-west-2.amazonaws.com/1.12.9/2019-06-21/bin/windows/amd64/kubectl.exe.sha256
```

- **Kubernetes 1.11:**

```
curl -o kubectl.exe.sha256 https://amazon-eks.s3-us-west-2.amazonaws.com/1.11.10/2019-06-21/bin/windows/amd64/kubectl.exe.sha256
```

- b. Check the SHA-256 sum for your downloaded binary.

```
Get-FileHash kubectl.exe
```

- c. Compare the generated SHA-256 sum in the command output against your downloaded SHA-256 file. The two should match, although the PowerShell output will be uppercase.
4. Copy the binary to a folder in your `PATH`. If you have an existing directory in your `PATH` that you use for command line utilities, copy the binary to that directory. Otherwise, complete the following steps.
  - a. Create a new directory for your command line binaries, such as `C:\bin`.
  - b. Copy the `kubectl.exe` binary to your new directory.
  - c. Edit your user or system `PATH` environment variable to add the new directory to your `PATH`.
  - d. Close your PowerShell terminal and open a new one to pick up the new `PATH` variable.
5. After you install **kubectl**, you can verify its version with the following command:

```
kubectl version --short --client
```

## Installing aws-iam-authenticator

Amazon EKS uses IAM to provide authentication to your Kubernetes cluster through the [AWS IAM Authenticator for Kubernetes](#). You can configure the stock **kubectl** client to work with Amazon EKS by installing the AWS IAM Authenticator for Kubernetes and modifying your **kubectl** configuration file to use it for authentication.

macOS

### To install aws-iam-authenticator with Homebrew

The easiest way to install the **aws-iam-authenticator** is with [Homebrew](#).

1. If you do not already have [Homebrew](#) installed on your Mac, install it with the following command.

```
/usr/bin/ruby -e "$(curl -fsSL https://raw.githubusercontent.com/Homebrew/install/master/install)"
```

2. Install the aws-iam-authenticator with the following command.

```
brew install aws-iam-authenticator
```

3. Test that the aws-iam-authenticator binary works.

```
aws-iam-authenticator help
```

## To install aws-iam-authenticator on macOS

You can also install the AWS-vended version of the aws-iam-authenticator by following these steps.

1. Download the Amazon EKS-vended aws-iam-authenticator binary from Amazon S3:

```
curl -o aws-iam-authenticator https://amazon-eks.s3-us-west-2.amazonaws.com/1.13.7/2019-06-11/bin/darwin/amd64/aws-iam-authenticator
```

2. (Optional) Verify the downloaded binary with the SHA-256 sum provided in the same bucket prefix.

- a. Download the SHA-256 sum for your system.

```
curl -o aws-iam-authenticator.sha256 https://amazon-eks.s3-us-west-2.amazonaws.com/1.13.7/2019-06-11/bin/darwin/amd64/aws-iam-authenticator.sha256
```

- b. Check the SHA-256 sum for your downloaded binary.

```
openssl sha1 -sha256 aws-iam-authenticator
```

- c. Compare the generated SHA-256 sum in the command output against your downloaded aws-iam-authenticator.sha256 file. The two should match.

3. Apply execute permissions to the binary.

```
chmod +x ./aws-iam-authenticator
```

4. Copy the binary to a folder in your \$PATH. We recommend creating a \$HOME/bin/aws-iam-authenticator and ensuring that \$HOME/bin comes first in your \$PATH.

```
mkdir -p $HOME/bin && cp ./aws-iam-authenticator $HOME/bin/aws-iam-authenticator && export PATH=$HOME/bin:$PATH
```

5. Add \$HOME/bin to your PATH environment variable.

```
echo 'export PATH=$HOME/bin:$PATH' >> ~/.bash_profile
```

6. Test that the aws-iam-authenticator binary works.

```
aws-iam-authenticator help
```

## Linux

### To install aws-iam-authenticator on Linux

1. Download the Amazon EKS-vended aws-iam-authenticator binary from Amazon S3:

```
curl -o aws-iam-authenticator https://amazon-eks.s3-us-west-2.amazonaws.com/1.13.7/2019-06-11/bin/linux/amd64/aws-iam-authenticator
```

2. (Optional) Verify the downloaded binary with the SHA-256 sum provided in the same bucket prefix.

- a. Download the SHA-256 sum for your system.

```
curl -o aws-iam-authenticator.sha256 https://amazon-eks.s3-us-west-2.amazonaws.com/1.13.7/2019-06-11/bin/linux/amd64/aws-iam-authenticator.sha256
```

- b. Check the SHA-256 sum for your downloaded binary.

```
openssl sha1 -sha256 aws-iam-authenticator
```

- c. Compare the generated SHA-256 sum in the command output against your downloaded aws-iam-authenticator.sha256 file. The two should match.

3. Apply execute permissions to the binary.

```
chmod +x ./aws-iam-authenticator
```

4. Copy the binary to a folder in your \$PATH. We recommend creating a \$HOME/bin/aws-iam-authenticator and ensuring that \$HOME/bin comes first in your \$PATH.

```
mkdir -p $HOME/bin && cp ./aws-iam-authenticator $HOME/bin/aws-iam-authenticator && export PATH=$HOME/bin:$PATH
```

5. Add \$HOME/bin to your PATH environment variable.

```
echo 'export PATH=$HOME/bin:$PATH' >> ~/.bashrc
```

6. Test that the aws-iam-authenticator binary works.

```
aws-iam-authenticator help
```

## Windows

### To install aws-iam-authenticator on Windows with Chocolatey

1. If you do not already have Chocolatey installed on your Windows system, see [Installing Chocolatey](#).
2. Open a PowerShell terminal window and install the aws-iam-authenticator package with the following command:

```
choco install -y aws-iam-authenticator
```

3. Test that the aws-iam-authenticator binary works.

```
aws-iam-authenticator help
```

### To install aws-iam-authenticator on Windows

1. Open a PowerShell terminal window and download the Amazon EKS-vended aws-iam-authenticator binary from Amazon S3:

```
curl -o aws-iam-authenticator.exe https://amazon-eks.s3-us-west-2.amazonaws.com/1.13.7/2019-06-11/bin/windows/amd64/aws-iam-authenticator.exe
```

2. (Optional) Verify the downloaded binary with the SHA-256 sum provided in the same bucket prefix.
  - a. Download the SHA-256 sum for your system.

```
curl -o aws-iam-authenticator.sha256 https://amazon-eks.s3-us-west-2.amazonaws.com/1.13.7/2019-06-11/bin/windows/amd64/aws-iam-authenticator.exe.sha256
```

- b. Check the SHA-256 sum for your downloaded binary.

```
Get-FileHash aws-iam-authenticator.exe
```

- c. Compare the generated SHA-256 sum in the command output against your downloaded SHA-256 file. The two should match, although the PowerShell output will be uppercase.
3. Copy the binary to a folder in your PATH. If you have an existing directory in your PATH that you use for command line utilities, copy the binary to that directory. Otherwise, complete the following steps.
    - a. Create a new directory for your command line binaries, such as C:\bin.
    - b. Copy the aws-iam-authenticator.exe binary to your new directory.
    - c. Edit your user or system PATH environment variable to add the new directory to your PATH.
    - d. Close your PowerShell terminal and open a new one to pick up the new PATH variable.
  4. Test that the aws-iam-authenticator binary works.

```
aws-iam-authenticator help
```

If you have an existing Amazon EKS cluster, create a kubeconfig file for that cluster. For more information, see [Create a kubeconfig for Amazon EKS \(p. 111\)](#). Otherwise, see [Creating an Amazon EKS Cluster \(p. 20\)](#) to create a new Amazon EKS cluster.

## Create a kubeconfig for Amazon EKS

In this section, you create a kubeconfig file for your cluster (or update an existing one).

This section offers two procedures to create or update your kubeconfig. You can quickly create or update a kubeconfig with the AWS CLI **update-kubeconfig** command by using the first procedure, or you can create a kubeconfig manually with the second procedure.

Amazon EKS uses the **aws eks get-token** command, available in version 1.16.156 or greater of the AWS CLI or the [AWS IAM Authenticator for Kubernetes](#) with **kubectrl** for cluster authentication. If you have

installed the AWS CLI on your system, then by default the AWS IAM Authenticator for Kubernetes will use the same credentials that are returned with the following command:

```
aws sts get-caller-identity
```

For more information, see [Configuring the AWS CLI](#) in the *AWS Command Line Interface User Guide*.

### To create your kubeconfig file with the AWS CLI

1. Ensure that you have at least version 1.16.156 of the AWS CLI installed. To install or upgrade the AWS CLI, see [Installing the AWS Command Line Interface](#) in the *AWS Command Line Interface User Guide*.

#### Note

Your system's Python version must be 2.7.9 or greater. Otherwise, you receive `hostname doesn't match` errors with AWS CLI calls to Amazon EKS. For more information, see [What are "hostname doesn't match" errors?](#) in the Python Requests FAQ.

You can check your AWS CLI version with the following command:

```
aws --version
```

#### Important

Package managers such **yum**, **apt-get**, or Homebrew for macOS are often behind several versions of the AWS CLI. To ensure that you have the latest version, see [Installing the AWS Command Line Interface](#) in the *AWS Command Line Interface User Guide*.

2. Use the AWS CLI **update-kubeconfig** command to create or update your kubeconfig for your cluster.
  - By default, the resulting configuration file is created at the default kubeconfig path (`.kube/config`) in your home directory or merged with an existing kubeconfig at that location. You can specify another path with the `--kubeconfig` option.
  - You can specify an IAM role ARN with the `--role-arn` option to use for authentication when you issue **kubectl** commands. Otherwise, the IAM entity in your default AWS CLI or SDK credential chain is used. You can view your default AWS CLI or SDK identity by running the **aws sts get-caller-identity** command.
  - For more information, see the help page with the **aws eks update-kubeconfig help** command or see [update-kubeconfig](#) in the *AWS CLI Command Reference*.

```
aws eks --region region update-kubeconfig --name cluster_name
```

3. Test your configuration.

```
kubectl get svc
```

#### Note

If you receive the error `"aws-iam-authenticator": executable file not found in $PATH`, your **kubectl** isn't configured for Amazon EKS. For more information, see [Installing aws-iam-authenticator](#) (p. 108).

If you receive any other authorization or resource type errors, see [Unauthorized or Access Denied \(kubectl\)](#) (p. 179) in the troubleshooting section.

Output:

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
------	------	------------	-------------	---------	-----

svc/kubernetes	ClusterIP	10.100.0.1	<none>	443/TCP	1m
----------------	-----------	------------	--------	---------	----

## To create your kubeconfig file manually

1. Create the default ~/.kube directory if it does not already exist.

```
mkdir -p ~/.kube
```

2. Open your favorite text editor and copy one of the kubeconfig code blocks below into it, depending on your preferred client token method.
  - To use the AWS CLI **aws eks get-token** command (requires at least version 1.16.156 of the AWS CLI):

```
apiVersion: v1
clusters:
- cluster:
    server: <endpoint-url>
    certificate-authority-data: <base64-encoded-ca-cert>
    name: kubernetes
contexts:
- context:
    cluster: kubernetes
    user: aws
    name: aws
current-context: aws
kind: Config
preferences: {}
users:
- name: aws
  user:
    exec:
      apiVersion: client.authentication.k8s.io/v1alpha1
      command: aws
      args:
        - "eks"
        - "get-token"
        - "--cluster-name"
        - "<cluster-name>"
        # - "--role"
        # - "<role-arn>"
      # env:
      # - name: AWS_PROFILE
      #   value: "<aws-profile>"
```

- To use the [AWS IAM Authenticator for Kubernetes](#):

```
apiVersion: v1
clusters:
- cluster:
    server: <endpoint-url>
    certificate-authority-data: <base64-encoded-ca-cert>
    name: kubernetes
contexts:
- context:
    cluster: kubernetes
    user: aws
    name: aws
current-context: aws
kind: Config
preferences: {}
```

```
users:
- name: aws
  user:
    exec:
      apiVersion: client.authentication.k8s.io/v1alpha1
      command: aws-iam-authenticator
      args:
        - "token"
        - "-i"
        - "<cluster-name>"
        # - "-r"
        # - "<role-arn>"
      # env:
      # - name: AWS_PROFILE
      #   value: "<aws-profile>"
```

3. Replace the `<endpoint-url>` with the endpoint URL that was created for your cluster.
4. Replace the `<base64-encoded-ca-cert>` with the `certificateAuthority.data` that was created for your cluster.
5. Replace the `<cluster-name>` with your cluster name.
6. (Optional) To assume an IAM role to perform cluster operations instead of the default AWS credential provider chain, uncomment the `-r` or `--role` and `<role-arn>` lines and substitute an IAM role ARN to use with your user.
7. (Optional) To always use a specific named AWS credential profile (instead of the default AWS credential provider chain), uncomment the `env` lines and substitute `<aws-profile>` with the profile name to use.
8. Save the file to the default **kubectl** folder, with your cluster name in the file name. For example, if your cluster name is `devel`, save the file to `~/.kube/config-devel`.
9. Add that file path to your `KUBECONFIG` environment variable so that **kubectl** knows where to look for your cluster configuration.

```
export KUBECONFIG=$KUBECONFIG:~/.kube/config-devel
```

10. (Optional) Add the configuration to your shell initialization file so that it is configured when you open a shell.
  - For Bash shells on macOS:

```
echo 'export KUBECONFIG=$KUBECONFIG:~/.kube/config-devel' >> ~/.bash_profile
```

- For Bash shells on Linux:

```
echo 'export KUBECONFIG=$KUBECONFIG:~/.kube/config-devel' >> ~/.bashrc
```

11. Test your configuration.

```
kubectl get svc
```

#### Note

If you receive the error `"aws-iam-authenticator": executable file not found in $PATH`, your **kubectl** isn't configured for Amazon EKS. For more information, see [Installing aws-iam-authenticator \(p. 108\)](#).

If you receive any other authorization or resource type errors, see [Unauthorized or Access Denied \(kubectl\) \(p. 179\)](#) in the troubleshooting section.

Output:

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
svc/kubernetes	ClusterIP	10.100.0.1	<none>	443/TCP	1m

## Managing Users or IAM Roles for your Cluster

When you create an Amazon EKS cluster, the IAM entity user or role, such as a [federated user](#) that creates the cluster, is automatically granted `system:masters` permissions in the cluster's RBAC configuration. To grant additional AWS users or roles the ability to interact with your cluster, you must edit the `aws-auth` ConfigMap within Kubernetes.

### Note

For more information about different IAM identities, see [Identities \(Users, Groups, and Roles\)](#) in the *IAM User Guide*. For more information on Kubernetes RBAC configuration, see [Using RBAC Authorization](#).

The `aws-auth` ConfigMap is applied as part of the [Getting Started with Amazon EKS \(p. 3\)](#) guide which provides a complete end-to-end walkthrough from creating an Amazon EKS cluster to deploying a sample Kubernetes application. It is initially created to allow your worker nodes to join your cluster, but you also use this ConfigMap to add RBAC access to IAM users and roles. If you have not launched worker nodes and applied the `aws-auth` ConfigMap, you can do so with the following procedure.

### To apply the `aws-auth` ConfigMap to your cluster

1. Check to see if you have already applied the `aws-auth` ConfigMap.

```
kubectl describe configmap -n kube-system aws-auth
```

If you receive an error stating "Error from server (NotFound): configmaps "aws-auth" not found", then proceed with the following steps to apply the stock ConfigMap.

2. Download, edit, and apply the AWS authenticator configuration map.

- a. Download the configuration map:

```
curl -o aws-auth-cm.yaml https://amazon-eks.s3-us-west-2.amazonaws.com/cloudformation/2019-02-11/aws-auth-cm.yaml
```

- b. Open the file with your favorite text editor. Replace the `<ARN of instance role (not instance profile)>` snippet with the Amazon Resource Name (ARN) of the IAM role that is associated with your worker nodes, and save the file. You can inspect the AWS CloudFormation stack outputs for your worker node groups and look for the following values:

- **InstanceRoleARN** (for worker node groups that were created with `eksctl`)
- **NodeInstanceRole** (for worker node groups that were created with Amazon EKS-vended AWS CloudFormation templates in the AWS Management Console)

### Important

Do not modify any other lines in this file.

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: aws-auth
  namespace: kube-system
data:
```



```
mapRoles: |
  - rolearn: <ARN of instance role (not instance profile)>
    username: system:node:{{EC2PrivateDNSName}}
  groups:
    - system:bootstrappers
    - system:nodes
```

- c. Apply the configuration. This command may take a few minutes to finish.

```
kubectl apply -f aws-auth-cm.yaml
```

#### Note

If you receive the error "aws-iam-authenticator": executable file not found in \$PATH, your **kubectl** isn't configured for Amazon EKS. For more information, see [Installing aws-iam-authenticator \(p. 108\)](#).  
If you receive any other authorization or resource type errors, see [Unauthorized or Access Denied \(kubectl\) \(p. 179\)](#) in the troubleshooting section.

3. Watch the status of your nodes and wait for them to reach the Ready status.

```
kubectl get nodes --watch
```

### To add an IAM user or role to an Amazon EKS cluster

1. Ensure that the AWS credentials that **kubectl** is using are already authorized for your cluster. The IAM user that created the cluster has these permissions by default.
2. Open the aws-auth ConfigMap.

```
kubectl edit -n kube-system configmap/aws-auth
```

#### Note

If you receive an error stating "Error from server (NotFound): configmaps "aws-auth" not found", then use the previous procedure to apply the stock ConfigMap.

Example ConfigMap:

```
# Please edit the object below. Lines beginning with a '#' will be ignored,
# and an empty file will abort the edit. If an error occurs while saving this file will
# be
# reopened with the relevant failures.
#
apiVersion: v1
data:
  mapRoles: |
    - rolearn: arn:aws:iam::111122223333:role/doc-test-worker-nodes-NodeInstanceRole-
      WDO5P42N3ETB
      username: system:node:{{EC2PrivateDNSName}}
      groups:
        - system:bootstrappers
        - system:nodes
kind: ConfigMap
metadata:
  annotations:
    kubectl.kubernetes.io/last-applied-configuration: |
      {"apiVersion":"v1","data":{"mapRoles":"- rolearn: arn:aws:iam::111122223333:role/
      doc-test-worker-nodes-NodeInstanceRole-WDO5P42N3ETB\n username: system:node:
      {{EC2PrivateDNSName}}\n groups:\n - system:bootstrappers\n -
      system:nodes\n"},"kind":"ConfigMap","metadata":{"annotations":{},"name":"aws-
      auth","namespace":"kube-system"}}
```

```
creationTimestamp: 2018-04-04T18:49:10Z
name: aws-auth
namespace: kube-system
resourceVersion: "780"
selfLink: /api/v1/namespaces/kube-system/configmaps/aws-auth
uid: dcc31de5-3838-11e8-af26-02e00430057c
```

3. Add your IAM users, roles, or AWS accounts to the configMap.

- **To add an IAM user:** add the user details to the `mapUsers` section of the ConfigMap, under `data`. Add this section if it does not already exist in the file. Each entry supports the following parameters:
  - **userarn:** The ARN of the IAM user to add.
  - **username:** The user name within Kubernetes to map to the IAM user. By default, the user name is the ARN of the IAM user.
  - **groups:** A list of groups within Kubernetes to which the user is mapped to. For more information, see [Default Roles and Role Bindings](#) in the Kubernetes documentation.
- **To add an IAM role (for example, for federated users):** add the role details to the `mapRoles` section of the ConfigMap, under `data`. Add this section if it does not already exist in the file. Each entry supports the following parameters:
  - **rolearn:** The ARN of the IAM role to add.
  - **username:** The user name within Kubernetes to map to the IAM role. By default, the user name is the ARN of the IAM role.
  - **groups:** A list of groups within Kubernetes to which the role is mapped. For more information, see [Default Roles and Role Bindings](#) in the Kubernetes documentation.

For example, the block below contains:

- A `mapRoles` section that adds the worker node instance role so that worker nodes can register themselves with the cluster.
- A `mapUsers` section with the AWS users `admin` from the default AWS account, and `ops-user` from another AWS account. Both users are added to the `system:masters` group.

```
# Please edit the object below. Lines beginning with a '#' will be ignored,
# and an empty file will abort the edit. If an error occurs while saving this file will
# be
# reopened with the relevant failures.
#
apiVersion: v1
data:
  mapRoles: |
    - rolearn: arn:aws:iam::555555555555:role/devel-worker-nodes-
NodeInstanceRole-74RF4UBDUKL6
      username: system:node:{{EC2PrivateDNSName}}
      groups:
        - system:bootstrappers
        - system:nodes
  mapUsers: |
    - userarn: arn:aws:iam::555555555555:user/admin
      username: admin
      groups:
        - system:masters
    - userarn: arn:aws:iam::111122223333:user/ops-user
      username: ops-user
      groups:
        - system:masters
```

4. Save the file and exit your text editor.

# The eksctl Command Line Utility

This chapter covers `eksctl`, a simple command line utility for creating and managing Kubernetes clusters on Amazon EKS. The `eksctl` command line utility provides the fastest and easiest way to create a new cluster with worker nodes for Amazon EKS.

For more information and to see the official documentation, visit <https://eksctl.io/>.

## Installing or Upgrading eksctl

This section helps you to install or upgrade the `eksctl` command line utility.

Choose the tab below that best represents your client setup.

macOS

### To install or upgrade eksctl on macOS using Homebrew

The easiest way to get started with Amazon EKS and macOS is by installing `eksctl` with [Homebrew](#). The `eksctl` Homebrew recipe installs `eksctl` and any other dependencies that are required for Amazon EKS, such as `kubectl` and the `aws-iam-authenticator`.

1. If you do not already have Homebrew installed on macOS, install it with the following command.

```
/usr/bin/ruby -e "$(curl -fsSL https://raw.githubusercontent.com/Homebrew/install/master/install)"
```

2. Install the Weaveworks Homebrew tap.

```
brew tap weaveworks/tap
```

3. Install or upgrade `eksctl`.

- Install `eksctl` with the following command:

```
brew install weaveworks/tap/eksctl
```

- If `eksctl` is already installed, run the following command to upgrade:

```
brew upgrade eksctl && brew link --overwrite eksctl
```

4. Test that your installation was successful with the following command.

```
eksctl version
```

#### Note

The `GitTag` version should be at least `0.1.37`. If not, check your terminal output for any installation or upgrade errors.

## Linux

### To install or upgrade eksctl on Linux using curl

1. Download and extract the latest release of eksctl with the following command.

```
curl --silent --location "https://github.com/weaveworks/eksctl/releases/download/latest_release/eksctl_$(uname -s)_amd64.tar.gz" | tar xz -C /tmp
```

2. Move the extracted binary to /usr/local/bin.

```
sudo mv /tmp/eksctl /usr/local/bin
```

3. Test that your installation was successful with the following command.

```
eksctl version
```

#### Note

The GitTag version should be at least 0.1.37. If not, check your terminal output for any installation or upgrade errors.

## Windows

### To install or upgrade eksctl on Windows using Chocolatey

1. If you do not already have Chocolatey installed on your Windows system, see [Installing Chocolatey](#).
2. Install or upgrade eksctl and the aws-iam-authenticator.
  - Install the binaries with the following command:

```
chocolatey install -y eksctl aws-iam-authenticator
```

- If they are already installed, run the following command to upgrade:

```
chocolatey upgrade -y eksctl aws-iam-authenticator
```

3. Test that your installation was successful with the following command.

```
eksctl version
```

#### Note

The GitTag version should be at least 0.1.37. If not, check your terminal output for any installation or upgrade errors.

# Pod Security Policy

The Kubernetes pod security policy admission controller validates pod creation and update requests against a set of rules. By default, Amazon EKS clusters ship with a fully permissive security policy with no restrictions. For more information, see [Pod Security Policies](#) in the Kubernetes documentation.

## Note

The pod security policy admission controller is only enabled on Amazon EKS clusters running Kubernetes version 1.13 or later. You must update your cluster's Kubernetes version to at least 1.13 to use pod security policies. For more information, see [Updating an Amazon EKS Cluster Kubernetes Version \(p. 26\)](#).

## Amazon EKS Default Pod Security Policy

Amazon EKS clusters with Kubernetes version 1.13 and higher have a default pod security policy named `eks.privileged`. This policy has no restriction on what kind of pod can be accepted into the system, which is equivalent to running Kubernetes with the `PodSecurityPolicy` controller disabled.

## Note

This policy was created to maintain backwards compatibility with clusters that did not have the `PodSecurityPolicy` controller enabled. You can create more restrictive policies for your cluster and for individual namespaces and service accounts and then delete the default policy to enable the more restrictive policies.

You can view the default policy with the following command.

```
kubectl get psp eks.privileged
```

Output:

NAME	PRIV	CAPS	SELINUX	RUNASUSER	FSGROUP	SUPGROUP	READONLYROOTFS
VOLUMES							
eks.privileged	true	*	RunAsAny	RunAsAny	RunAsAny	RunAsAny	false
*							

For more details, you can describe the policy with the following command.

```
kubectl describe psp eks.privileged
```

Output:

```
Name:  eks.privileged

Settings:
  Allow Privileged:            true
  Allow Privilege Escalation:  0xc0004ce5f8
  Default Add Capabilities:    <none>
  Required Drop Capabilities:  <none>
  Allowed Capabilities:        *
  Allowed Volume Types:       *
  Allow Host Network:          true
  Allow Host Ports:            0-65535
  Allow Host PID:              true
```

Allow Host IPC:	true
Read Only Root Filesystem:	false
SELinux Context Strategy: RunAsAny	
User:	<none>
Role:	<none>
Type:	<none>
Level:	<none>
Run As User Strategy: RunAsAny	
Ranges:	<none>
FSGroup Strategy: RunAsAny	
Ranges:	<none>
Supplemental Groups Strategy: RunAsAny	
Ranges:	<none>

The following example shows the full YAML file for the `eks.privileged` pod security policy, its cluster role, and cluster role binding.

```
---
apiVersion: policy/v1beta1
kind: PodSecurityPolicy
metadata:
  name: eks.privileged
  annotations:
    kubernetes.io/description: 'privileged allows full unrestricted access to
      pod features, as if the PodSecurityPolicy controller was not enabled.'
    seccomp.security.alpha.kubernetes.io/allowedProfileNames: '*'
  labels:
    kubernetes.io/cluster-service: "true"
    eks.amazonaws.com/component: pod-security-policy
spec:
  privileged: true
  allowPrivilegeEscalation: true
  allowedCapabilities:
  - '*'
  volumes:
  - '*'
  hostNetwork: true
  hostPorts:
  - min: 0
    max: 65535
  hostIPC: true
  hostPID: true
  runAsUser:
    rule: 'RunAsAny'
  seLinux:
    rule: 'RunAsAny'
  supplementalGroups:
    rule: 'RunAsAny'
  fsGroup:
    rule: 'RunAsAny'
  readOnlyRootFilesystem: false

---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: eks:podsecuritypolicy:privileged
  labels:
    kubernetes.io/cluster-service: "true"
    eks.amazonaws.com/component: pod-security-policy
rules:
- apiGroups:
  - policy
  resourceNames:
```

```
- eks.privileged
resources:
- podsecuritypolicies
verbs:
- use

---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: eks:podsecuritypolicy:authenticated
  annotations:
    kubernetes.io/description: 'Allow all authenticated users to create privileged pods.'
  labels:
    kubernetes.io/cluster-service: "true"
    eks.amazonaws.com/component: pod-security-policy
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: eks:podsecuritypolicy:privileged
subjects:
- kind: Group
  apiGroup: rbac.authorization.k8s.io
  name: system:authenticated
```

### To delete the default pod security policy

After you create custom pod security policies for your cluster, you can delete the default Amazon EKS `eks.privileged` pod security policy to enable your custom policies.

1. Create a file called `privileged-podsecuritypolicy.yaml` and paste the full `eks.privileged` YAML file contents from the preceding example into it (this allows you to delete the pod security policy, the `ClusterRole`, and the `ClusterRoleBinding` associated with it).
2. Delete the YAML with the following command.

```
kubectl delete -f privileged-podsecuritypolicy.yaml
```

### To restore the default pod security policy

If you have modified or deleted the default Amazon EKS `eks.privileged` pod security policy, you can restore it with the following steps.

1. Create a file called `privileged-podsecuritypolicy.yaml` and paste the full `eks.privileged` YAML file contents from the preceding example into it.
2. Apply the YAML with the following command.

```
kubectl apply -f privileged-podsecuritypolicy.yaml
```



# Launch a Guest Book Application

In this topic, you create a sample guest book application to test your Amazon EKS cluster.

**Note**

For more information about setting up the guest book example, see <https://github.com/kubernetes/examples/blob/master/guestbook-go/README.md> in the Kubernetes documentation.

**To create your guest book application**

1. Create the Redis master replication controller.

```
kubectl apply -f https://raw.githubusercontent.com/kubernetes/examples/master/guestbook-go/redis-master-controller.json
```

**Note**

If you receive the error "aws-iam-authenticator": executable file not found in \$PATH, your **kubectl** isn't configured for Amazon EKS. For more information, see [Installing aws-iam-authenticator \(p. 108\)](#).

If you receive any other authorization or resource type errors, see [Unauthorized or Access Denied \(kubectl\) \(p. 179\)](#) in the troubleshooting section.

Output:

```
replicationcontroller "redis-master" created
```

2. Create the Redis master service.

```
kubectl apply -f https://raw.githubusercontent.com/kubernetes/examples/master/guestbook-go/redis-master-service.json
```

Output:

```
service "redis-master" created
```

3. Create the Redis slave replication controller.

```
kubectl apply -f https://raw.githubusercontent.com/kubernetes/examples/master/guestbook-go/redis-slave-controller.json
```

Output:

```
replicationcontroller "redis-slave" created
```

4. Create the Redis slave service.

```
kubectl apply -f https://raw.githubusercontent.com/kubernetes/examples/master/guestbook-go/redis-slave-service.json
```

Output:

```
service "redis-slave" created
```

5. Create the guestbook replication controller.

```
kubectl apply -f https://raw.githubusercontent.com/kubernetes/examples/master/guestbook-go/guestbook-controller.json
```

Output:

```
replicationcontroller "guestbook" created
```

6. Create the guestbook service.

```
kubectl apply -f https://raw.githubusercontent.com/kubernetes/examples/master/guestbook-go/guestbook-service.json
```

Output:

```
service "guestbook" created
```

7. Query the services in your cluster and wait until the **External IP** column for the guestbook service is populated.

**Note**

It might take several minutes before the IP address is available.

```
kubectl get services -o wide
```

8. After your external IP address is available, point a web browser to that address at port 3000 to view your guest book. For example, <http://a7a95c2b9e69711e7b1a3022f9dcfd2e-1985673473.us-west-2.elb.amazonaws.com:3000>

**Note**

It might take several minutes for DNS to propagate and for your guest book to show up.

# Guestbook

Rick

Morty

Bird Person

Sleepy Gary

Tophat Jones

## Important

If you are unable to connect to the external IP address with your browser, be sure that your corporate firewall is not blocking non-standards ports, like 3000. You can try switching to a guest network to verify.

## To clean up your guest book application

When you are finished experimenting with your guest book application, you should clean up the resources that you created for it.

- The following command deletes all of the services and replication controllers for the guest book application:

```
kubectl delete rc/redis-master rc/redis-slave rc/guestbook svc/redis-master svc/redis-slave svc/guestbook
```

## Note

If you receive the error "aws-iam-authenticator": executable file not found in \$PATH, your **kubectl** isn't configured for Amazon EKS. For more information, see [Installing aws-iam-authenticator \(p. 108\)](#).

If you receive any other authorization or resource type errors, see [Unauthorized or Access Denied \(kubectl\) \(p. 179\)](#) in the troubleshooting section.

If you are done with your Amazon EKS cluster, you should delete it and its resources so that you do not incur additional charges. For more information, see [Deleting a Cluster \(p. 41\)](#).

# Installing the Kubernetes Metrics Server

The Kubernetes metrics server is an aggregator of resource usage data in your cluster, and it is not deployed by default in Amazon EKS clusters. This topic explains how to deploy the Kubernetes metrics server on your Amazon EKS cluster.

## Note

The Kubernetes metrics server must be installed on your cluster to use the [Horizontal Pod Autoscaler](#).

## To install `metrics-server` from GitHub on an Amazon EKS cluster using `curl` and `jq`

If you have a macOS or Linux system with `curl`, `tar`, `gzip`, and the `jq` JSON parser installed, you can download, extract, and install the latest release with the following commands. Otherwise, use the next procedure to download the latest version using a web browser.

1. Open a terminal window and navigate to a directory where you would like to download the latest `metrics-server` release.
2. Copy and paste the commands below into your terminal window and type **Enter** to execute them. These commands download the latest release, extract it, and apply the version 1.8+ manifests to your cluster.

```
DOWNLOAD_URL=$(curl --silent "https://api.github.com/repos/kubernetes-incubator/metrics-server/releases/latest" | jq -r .tarball_url)
DOWNLOAD_VERSION=$(grep -o '^[^/]*#' <<< $DOWNLOAD_URL)
curl -Ls $DOWNLOAD_URL -o metrics-server-$DOWNLOAD_VERSION.tar.gz
mkdir metrics-server-$DOWNLOAD_VERSION
tar -xzf metrics-server-$DOWNLOAD_VERSION.tar.gz --directory metrics-server-$DOWNLOAD_VERSION --strip-components 1
kubectl apply -f metrics-server-$DOWNLOAD_VERSION/deploy/1.8+/
```

3. Verify that the `metrics-server` deployment is running the desired number of pods with the following command:

```
kubectl get deployment metrics-server -n kube-system
```

Output:

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
metrics-server	1	1	1	1	56m

## To install `metrics-server` from GitHub on an Amazon EKS cluster using a web browser

1. Download and extract the latest version of the metrics server code from GitHub.
  - a. Navigate to the latest release page of the `metrics-server` project on GitHub (<https://github.com/kubernetes-incubator/metrics-server/releases/latest>), then choose a source code archive for the latest release to download it.
  - b. Navigate to your downloads folder and extract the source code archive. For example, if you downloaded the `.tar.gz` archive on a macOS or Linux system, use the following command to extract (substituting your release version).

```
tar -xzf metrics-server-0.3.1.tar.gz
```

2. Apply all of the YAML manifests in the `metrics-server-0.3.1/deploy/1.8+` directory (substituting your release version).

```
kubectl apply -f metrics-server-0.3.1/deploy/1.8+/
```

3. Verify that the `metrics-server` deployment is running the desired number of pods with the following command:

```
kubectl get deployment metrics-server -n kube-system
```

Output:

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
metrics-server	1	1	1	1	56m

# Control Plane Metrics with Prometheus

The Kubernetes API server exposes a number of metrics that are useful for monitoring and analysis. These metrics are exposed internally through a metrics endpoint that refers to the `/metrics` HTTP API. Like other endpoints, this endpoint is exposed on the Amazon EKS control plane. This topic explains some of the ways you can use this endpoint to view and analyze what your cluster is doing.

## Viewing the Raw Metrics

To view the raw metrics output, use `kubectl` with the `--raw` flag. This command allows you to pass any HTTP path and returns the raw response.

```
kubectl get --raw /metrics
```

Example output:

```
...
# HELP rest_client_requests_total Number of HTTP requests, partitioned by status code,
# method, and host.
# TYPE rest_client_requests_total counter
rest_client_requests_total{code="200",host="127.0.0.1:21362",method="POST"} 4994
rest_client_requests_total{code="200",host="127.0.0.1:443",method="DELETE"} 1
rest_client_requests_total{code="200",host="127.0.0.1:443",method="GET"} 1.326086e+06
rest_client_requests_total{code="200",host="127.0.0.1:443",method="PUT"} 862173
rest_client_requests_total{code="404",host="127.0.0.1:443",method="GET"} 2
rest_client_requests_total{code="409",host="127.0.0.1:443",method="POST"} 3
rest_client_requests_total{code="409",host="127.0.0.1:443",method="PUT"} 8
# HELP ssh_tunnel_open_count Counter of ssh tunnel total open attempts
# TYPE ssh_tunnel_open_count counter
ssh_tunnel_open_count 0
# HELP ssh_tunnel_open_fail_count Counter of ssh tunnel failed open attempts
# TYPE ssh_tunnel_open_fail_count counter
ssh_tunnel_open_fail_count 0
```

This raw output returns verbatim what the API server exposes. These metrics are represented in a [Prometheus format](#). This format allows the API server to expose different metrics broken down by line. Each line includes a metric name, tags, and a value.

```
metric_name{"tag"="value"[,...]} value
```

While this endpoint is useful if you are looking for a specific metric, you typically want to analyze these metrics over time. To do this, you can deploy [Prometheus](#) into your cluster. Prometheus is a monitoring and time series database that scrapes exposed endpoints and aggregates data, allowing you to filter, graph, and query the results.

## Deploying Prometheus

This topic helps you deploy Prometheus into your cluster with Helm. Helm is a package manager for Kubernetes clusters. For more information, see [Using Helm with Amazon EKS \(p. 134\)](#).

After you configure Helm for your Amazon EKS cluster, you can use it to deploy Prometheus with the following steps.

### To deploy Prometheus using Helm

1. Follow the steps in [Using Helm with Amazon EKS \(p. 134\)](#) to get working `helm` and `tiller` terminal windows, so that you can install Helm charts.
2. In the Helm terminal window, run the following commands to deploy Prometheus.
  - a. Create a Prometheus namespace.

```
kubectl create namespace prometheus
```

- b. Deploy Prometheus.

```
helm install stable/prometheus \
--name prometheus \
--namespace prometheus \
--set
  alertmanager.persistentVolume.storageClass="gp2",server.persistentVolume.storageClass="gp2"
```

3. Verify that all of the pods in the `prometheus` namespace are in the `READY` state.

```
kubectl get pods -n prometheus
```

Output:

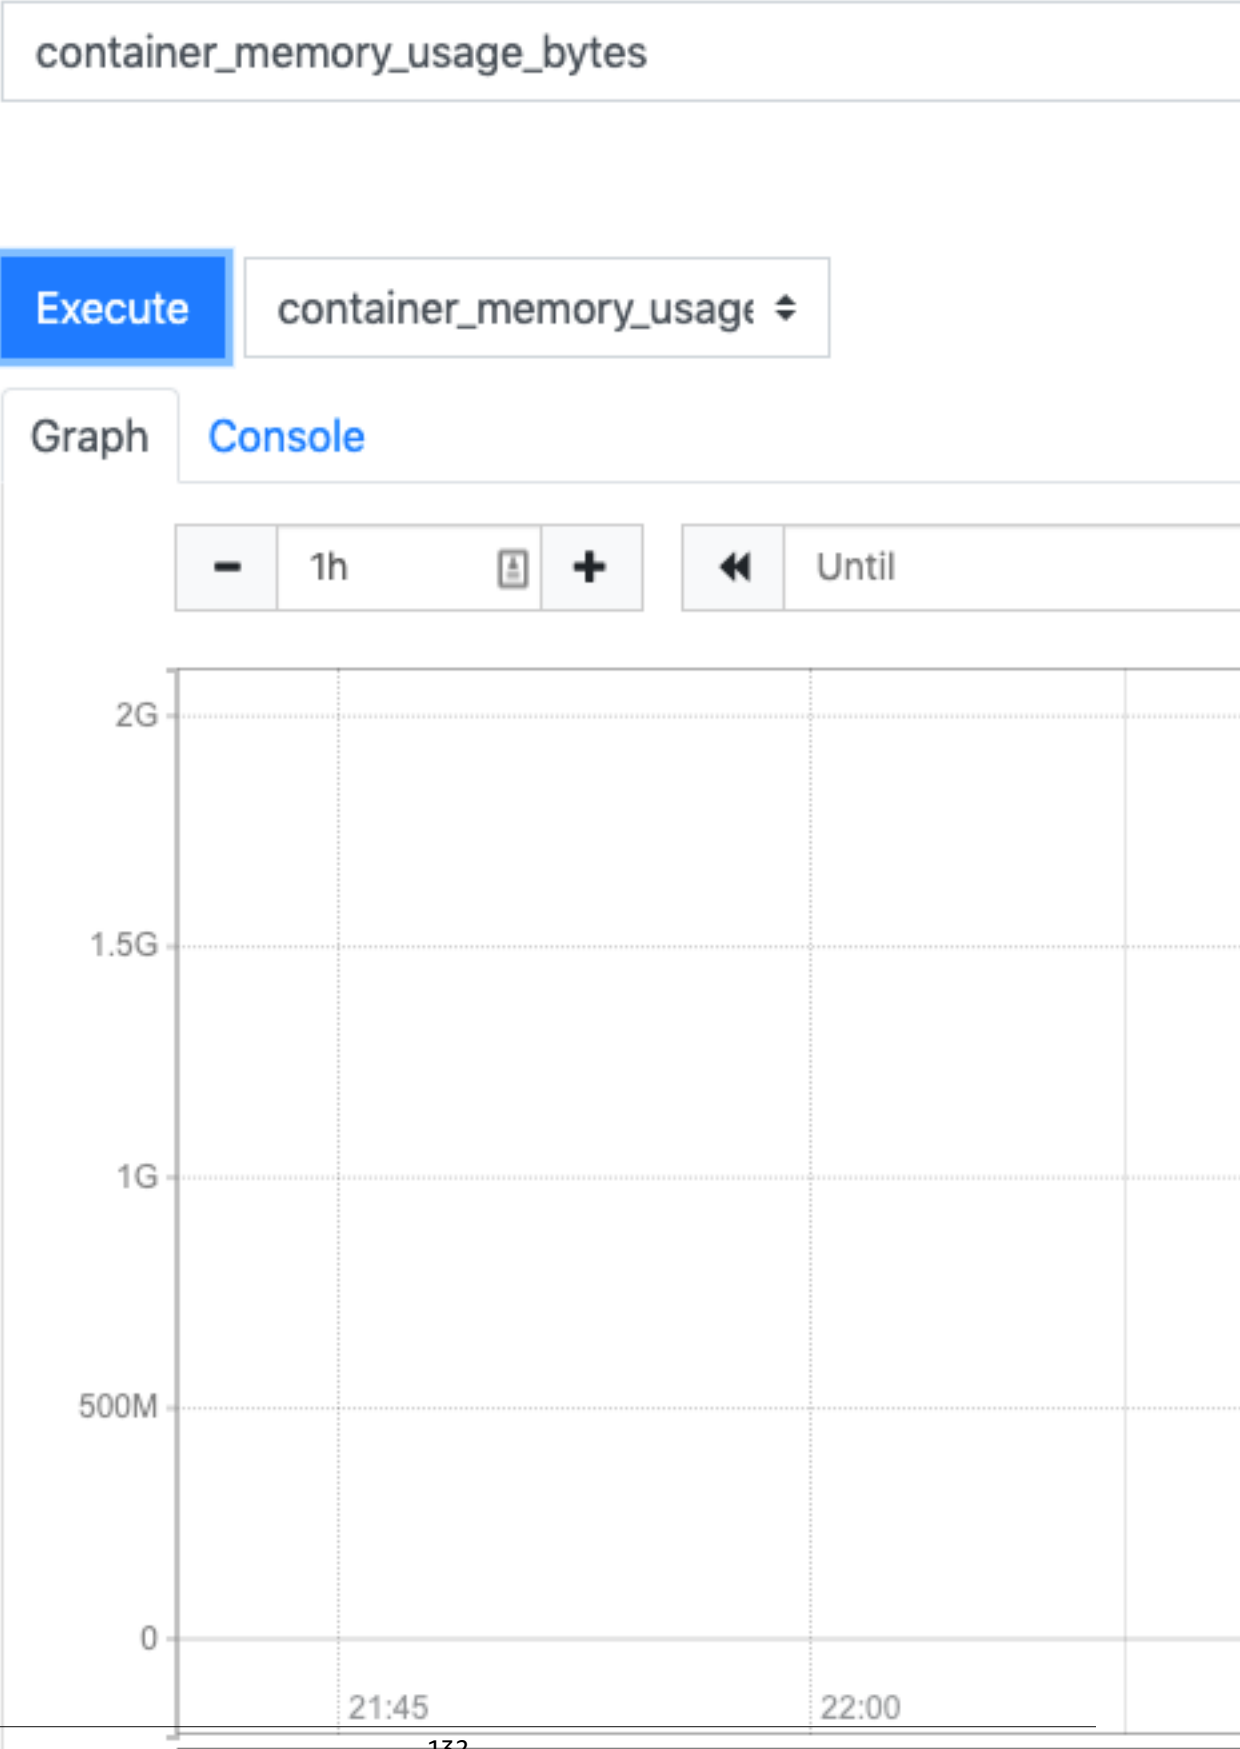
NAME	READY	STATUS	RESTARTS	AGE
prometheus-alertmanager-848fb754f5-2wpbm	2/2	Running	0	85s
prometheus-kube-state-metrics-86cbcf9b6f-drnfq	1/1	Running	0	85s
prometheus-node-exporter-8qpcl	1/1	Running	0	85s
prometheus-node-exporter-czz9g	1/1	Running	0	85s
prometheus-node-exporter-ffs19	1/1	Running	0	85s
prometheus-pushgateway-564f65fcc8-hmzp6	1/1	Running	0	85s
prometheus-server-5b65bd569b-6wgwx	2/2	Running	0	85s

4. Use `kubectl` to port forward the Prometheus console to your local machine.

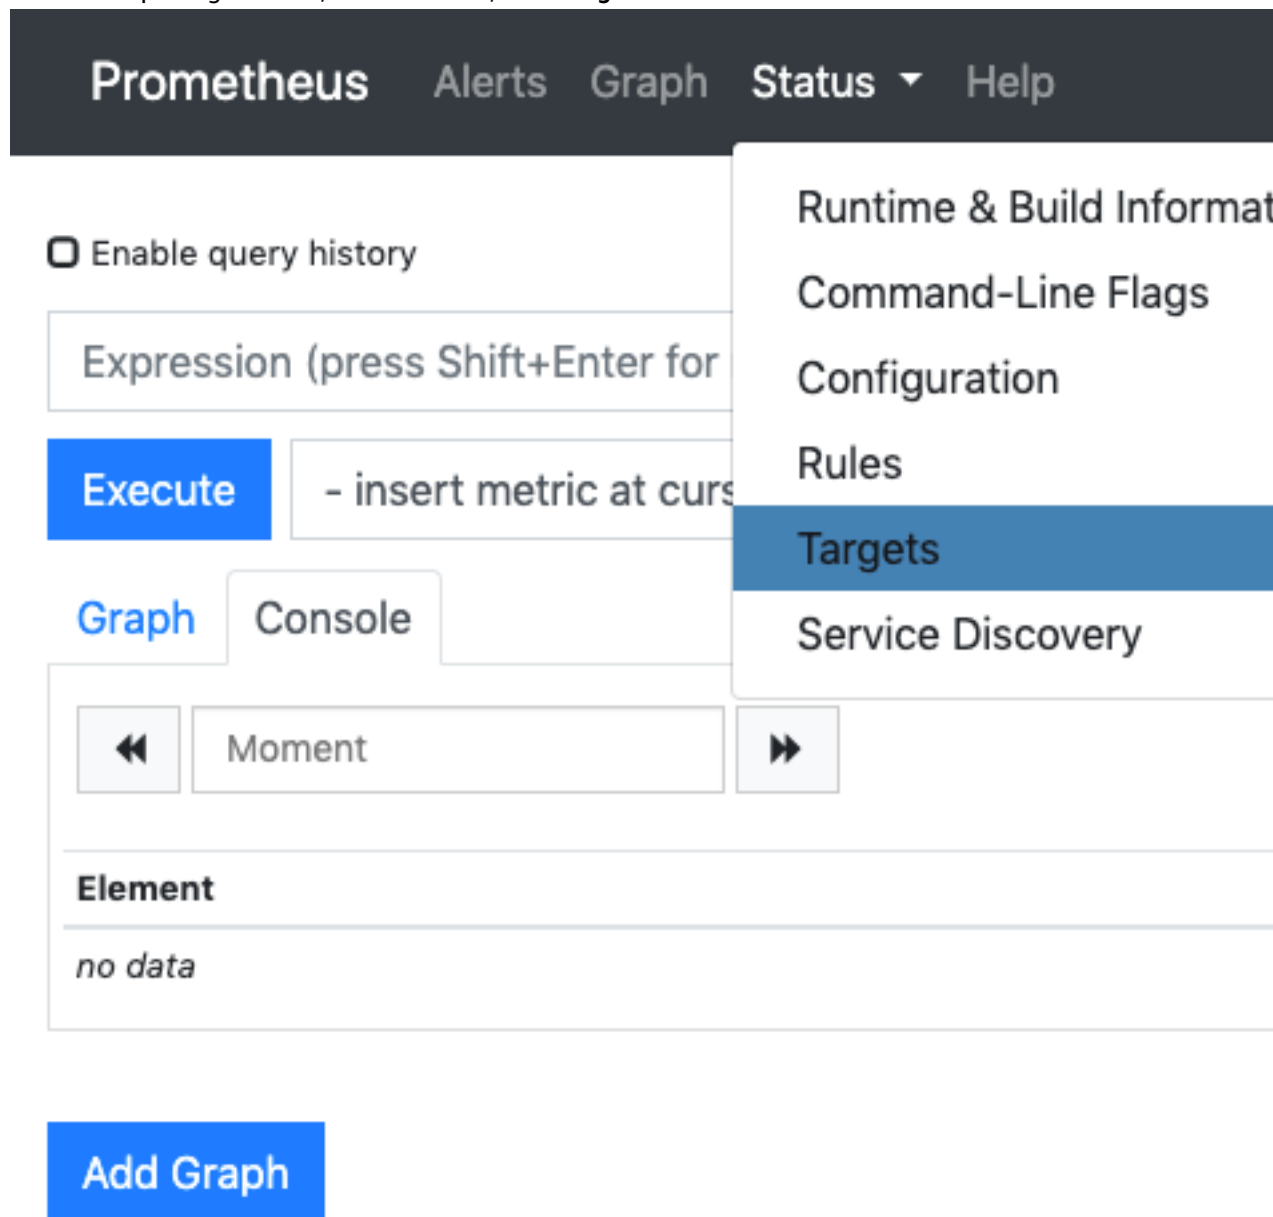
```
kubectl --namespace=prometheus port-forward deploy/prometheus-server 9090
```

5. Point a web browser to [localhost:9090](http://localhost:9090) to view the Prometheus console.
6. Choose a metric from the - **insert metric at cursor** menu, then choose **Execute**. Choose the **Graph** tab to show the metric over time. The following image shows `container_memory_usage_bytes` over time.





7. From the top navigation bar, choose **Status**, then **Targets**.



All of the Kubernetes endpoints that are connected to Prometheus using service discovery are displayed.

# Using Helm with Amazon EKS

The `helm` package manager for Kubernetes helps you install and manage applications on your Kubernetes cluster. For more information, see the [Helm documentation](#). This topic helps you install and run the `helm` and `tiller` binaries locally so that you can install and manage charts using the `helm` CLI on your local system.

Although you can run the server-side `tiller` component in your cluster (and many public Helm installation articles offer only this option), running `tiller` locally in its own namespace as described in this topic reduces the risk of exploit for your cluster in the following ways:

- When you run the `tiller` server on your cluster, it gets its own Kubernetes Identity and associated permission set, often with full Kubernetes administrator permissions. This opens up the possibility for a privilege escalation, where an unprivileged Kubernetes user who has network access to the `tiller` server can gain additional Kubernetes permissions by way of installing a chart.
- When you run the `tiller` server on your local machine, users don't inherit the `tiller` server permissions on the cluster (likely full-admin), but instead `tiller` inherits the Kubernetes permissions of the end-user.
- Running `tiller` in its own namespace allows you to control access to the Kubernetes secrets that the `tiller` server stores by controlling access to that namespace.

## Important

Before you can install Helm charts on your Amazon EKS cluster, you must configure `kubectrl` to work for Amazon EKS. If you have not already done this, see [Installing aws-iam-authenticator \(p. 108\)](#) and [Create a kubeconfig for Amazon EKS \(p. 111\)](#) before proceeding. If the following command succeeds for your cluster, you're properly configured.

```
kubectrl get svc
```

## To install the `helm` and `tiller` binaries on your local system

1. • If you're using macOS with [Homebrew](#), install the binaries with the following command.

```
brew install kubernetes-helm
```

- If you're using Windows with [Chocolatey](#), install the binaries with the following command.

```
choco install kubernetes-helm
```

- Otherwise, see [Installing the Helm Client](#) in the Helm documentation.

## Important

Don't proceed to install the `tiller` server-side component with the Helm documentation (stop before you reach [Installing Tiller](#)). This topic explains how to run `tiller` locally in its own namespace, which reduces the risk of exploit for your cluster.

2. To pick up the new binaries in your `PATH`, Close your current terminal window and open a new one.

## To run `helm` and `tiller` locally

1. Create a namespace called `tiller` with the following command.

```
kubectl create namespace tiller
```

### Note

By default, `tiller` stores its secrets in the `kube-system` namespace. Creating a namespace for `tiller` and specifying that namespace when you run it gives you more specific access controls to who is authorized to view the Helm chart secrets that `tiller` stores in your cluster.

2. Open a new terminal window for the `tiller` server. For the following steps, you need a terminal window for the `tiller` server and another window for the `helm` client.

### Important

You should ensure that you are the only active user for the system that you use for the `tiller` server (such as a local laptop or desktop where you are the only user that is logged in). Otherwise, any user on your system could make requests to the `tiller` server with your Kubernetes permissions. For Linux and macOS systems, you can see the current users with the following command:

```
users
```

Output:

```
ericn
```

In the above example, there is only a single user named `ericn` on the system, so it is safe to proceed. If there are more than one user logged in to your system, you should use a different system, or consider launching an Amazon EC2 instance for this procedure so that you can ensure that you are the only active user.

3. In the `tiller` server terminal, set the `TILLER_NAMESPACE` environment variable to `tiller` and then start the `tiller` server.
  - a. Set the `TILLER_NAMESPACE` environment variable to `tiller`.

- **macOS and Linux:**

```
export TILLER_NAMESPACE=tiller
```

- **Windows (PowerShell):**

```
$env:TILLER_NAMESPACE = 'tiller'
```

- b. Start the `tiller` server.

- **macOS and Linux:**

```
tiller -listen=localhost:44134 -storage=secret -logtostderr
```

- **Windows (PowerShell):**

```
tiller -listen=localhost:44134 -storage=secret
```

### Note

By default, `tiller` stores release information in ConfigMaps; however, the latest Helm documentation recommends that you use the `-storage=secret` flag to store this

information with Kubernetes secrets instead. For more information, see [Tiller's Release Information](#) in [Securing your Helm Installation](#). The `-listen=localhost:44134` flag ensures that the `tiller` server only accepts requests from your local machine (this prevents unauthorized network users from accessing your local `tiller` process).

4. In the `helm` client terminal window, set the `HELM_HOST` environment variable to `:44134`.

- **macOS and Linux:**

```
export HELM_HOST=:44134
```

- **Windows (PowerShell):**

```
$env:HELM_HOST = ':44134'
```

5. In the `helm` client terminal window, initialize the `helm` client.

```
helm init --client-only
```

6. In the `helm` client terminal window, verify that `helm` is communicating with the `tiller` server properly.

```
helm repo update
```

Output:

```
Hang tight while we grab the latest from your chart repositories...
...Skip local chart repository
...Successfully got an update from the "stable" chart repository
Update Complete. # Happy Helming!#
```

7. At this point, you can run any `helm` commands in your `helm` client terminal window (such as `helm install chart_name`) to install, modify, delete, or query Helm charts in your cluster. As you run `helm` commands, you can follow the `tiller` logs for those commands in its server terminal window. For more information, see [Helm Commands](#) and [Charts](#) in the Helm documentation.

If you're just experimenting with `helm` and you don't have a specific chart to install, you can see [Install an Example Chart](#) in the [Helm Quickstart Guide](#).

8. When you're finished, close your `helm` client and `tiller` server terminal windows. Repeat this procedure when you want to use `helm` with your cluster.

# Tutorial: Deploy the Kubernetes Web UI (Dashboard)

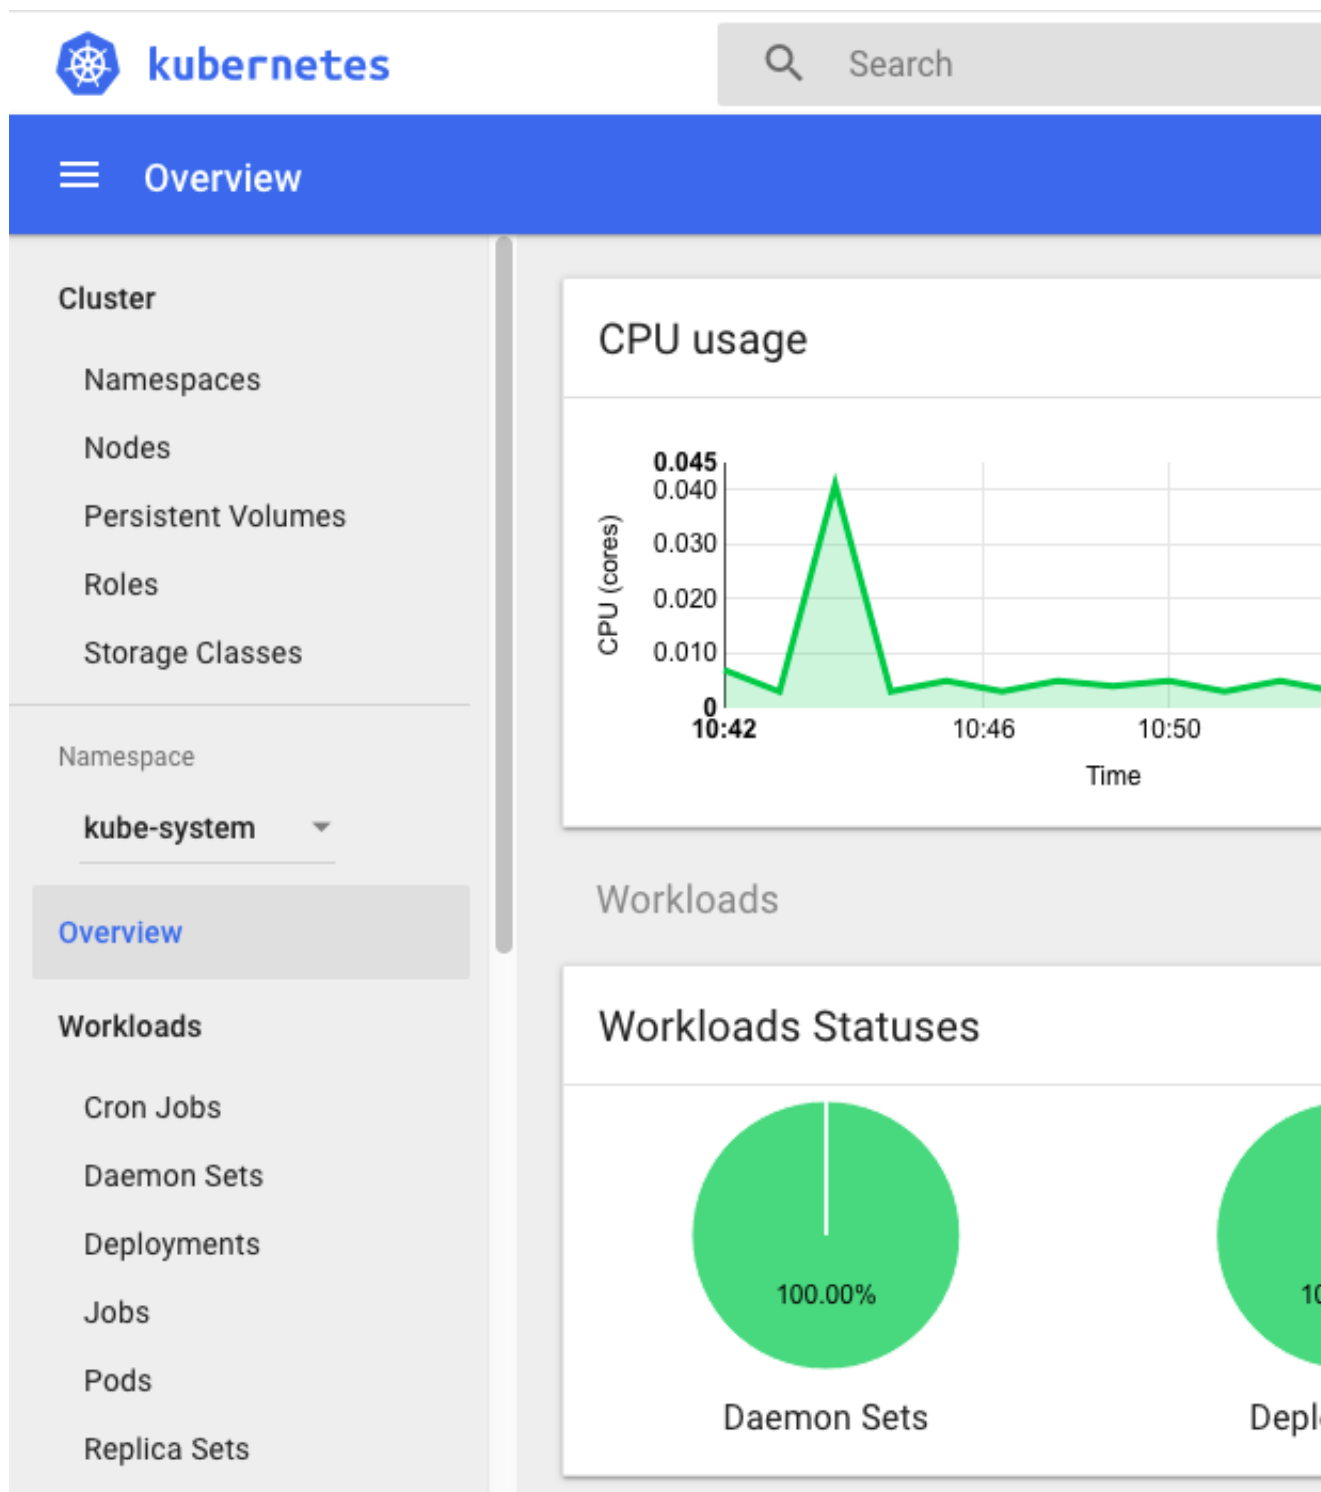
This tutorial guides you through deploying the [Kubernetes dashboard](#) to your Amazon EKS cluster, complete with CPU and memory metrics. It also helps you to create an Amazon EKS administrator service account that you can use to securely connect to the dashboard to view and control your cluster.

**Important**

Kubernetes versions 1.11 and above do not support `heapster` memory and CPU metrics in the dashboard by default. The community is working to replace `heapster` in the dashboard with the Kubernetes metrics server to fix this issue. For more information, see <https://github.com/kubernetes/dashboard/issues/3147> and <https://github.com/kubernetes/dashboard/issues/2986>.

When the dashboard project is updated to use the Kubernetes metrics server, this topic will be updated with information about how to restore the CPU and memory metrics functionality. Until that time, CPU and memory metrics are not visible in the dashboard on Amazon EKS cluster versions 1.11 and above.

There is a [potential workaround posted on GitHub](#), but the Amazon EKS team has not evaluated the risk of setting the `insecure=true` flag on the `heapster` source, so we cannot recommend the workaround at this time.



## Prerequisites

This tutorial assumes the following:

- You have created an Amazon EKS cluster by following the steps in [Getting Started with Amazon EKS \(p. 3\)](#).
- The security groups for your control plane elastic network interfaces and worker nodes follow the recommended settings in [Cluster Security Group Considerations \(p. 83\)](#).
- You are using a **kubectl** client that is [configured to communicate with your Amazon EKS cluster \(p. 13\)](#).

## Step 1: Deploy the Dashboard

Use the following steps to deploy the Kubernetes dashboard, heapster, and the influxdb backend for CPU and memory metrics to your cluster.

### To deploy the Kubernetes dashboard

1. Deploy the Kubernetes dashboard to your cluster:

```
kubectl apply -f https://raw.githubusercontent.com/kubernetes/dashboard/v1.10.1/src/
deploy/recommended/kubernetes-dashboard.yaml
```

Output:

```
secret "kubernetes-dashboard-certs" created
serviceaccount "kubernetes-dashboard" created
role "kubernetes-dashboard-minimal" created
rolebinding "kubernetes-dashboard-minimal" created
deployment "kubernetes-dashboard" created
service "kubernetes-dashboard" created
```

2. Deploy heapster to enable container cluster monitoring and performance analysis on your cluster:

```
kubectl apply -f https://raw.githubusercontent.com/kubernetes/heapster/master/deploy/
kube-config/influxdb/heapster.yaml
```

#### Note

Although heapster is deprecated, it is currently the only supported metrics provider for the Kubernetes dashboard. For more information, see <https://github.com/kubernetes/dashboard/issues/2986>.

Output:

```
serviceaccount "heapster" created
deployment "heapster" created
service "heapster" created
```

3. Deploy the influxdb backend for heapster to your cluster:

```
kubectl apply -f https://raw.githubusercontent.com/kubernetes/heapster/master/deploy/
kube-config/influxdb/influxdb.yaml
```

Output:

```
deployment "monitoring-influxdb" created
service "monitoring-influxdb" created
```

4. Create the heapster cluster role binding for the dashboard:



```
kubectl apply -f https://raw.githubusercontent.com/kubernetes/heapster/master/deploy/kube-config/rbac/heapster-rbac.yaml
```

Output:

```
clusterrolebinding "heapster" created
```

## Step 2: Create an eks-admin Service Account and Cluster Role Binding

By default, the Kubernetes dashboard user has limited permissions. In this section, you create an eks-admin service account and cluster role binding that you can use to securely connect to the dashboard with admin-level permissions. For more information, see [Managing Service Accounts](#) in the Kubernetes documentation.

### To create the eks-admin service account and cluster role binding

#### Important

The example service account created with this procedure has full cluster-admin (superuser) privileges on the cluster. For more information, see [Using RBAC Authorization](#) in the Kubernetes documentation.

1. Create a file called `eks-admin-service-account.yaml` with the text below. This manifest defines a service account and cluster role binding called `eks-admin`.

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: eks-admin
  namespace: kube-system
---
apiVersion: rbac.authorization.k8s.io/v1beta1
kind: ClusterRoleBinding
metadata:
  name: eks-admin
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: cluster-admin
subjects:
- kind: ServiceAccount
  name: eks-admin
  namespace: kube-system
```

2. Apply the service account and cluster role binding to your cluster:

```
kubectl apply -f eks-admin-service-account.yaml
```

Output:

```
serviceaccount "eks-admin" created
clusterrolebinding.rbac.authorization.k8s.io "eks-admin" created
```

## Step 3: Connect to the Dashboard

Now that the Kubernetes dashboard is deployed to your cluster, and you have an administrator service account that you can use to view and control your cluster, you can connect to the dashboard with that service account.

### To connect to the Kubernetes dashboard

1. Retrieve an authentication token for the `eks-admin` service account. Copy the `<authentication_token>` value from the output. You use this token to connect to the dashboard.

```
kubectl -n kube-system describe secret $(kubectl -n kube-system get secret | grep eks-admin | awk '{print $1}')
```

Output:

```
Name:          eks-admin-token-b5zv4
Namespace:     kube-system
Labels:        <none>
Annotations:   kubernetes.io/service-account.name=eks-admin
               kubernetes.io/service-account.uid=bcfe66ac-39be-11e8-97e8-026dce96b6e8

Type:  kubernetes.io/service-account-token

Data
====
ca.crt:      1025 bytes
namespace:   11 bytes
token:       <authentication_token>
```

2. Start the `kubectl proxy`.

```
kubectl proxy
```

3. Open the following link with a web browser to access the dashboard endpoint: <http://localhost:8001/api/v1/namespaces/kube-system/services/https:kubernetes-dashboard:/proxy/#!/login>
4. Choose **Token**, paste the `<authentication_token>` output from the previous command into the **Token** field, and choose **SIGN IN**.

## Kubernetes Dashboard

☐

### Kubeconfig

Please select the kubeconfig file that you have created to configure access to the cluster. To find out more about how to configure and use kubeconfig file, please refer to the [Configure Access to Multiple Clusters](#) section.

☒

### Token

Every Service Account has a Secret with valid Bearer Token that can be used to log in to Dashboard. To find out more about how to configure and use Bearer Tokens, please refer to the [Authentication](#) section.

Enter token

SIGN IN

SKIP

**Note**

It may take a few minutes before CPU and memory metrics appear in the dashboard.

## Step 4: Next Steps

After you have connected to your Kubernetes cluster dashboard, you can view and control your cluster using your `eks-admin` service account. For more information about using the dashboard, see the [project documentation on GitHub](#).

# Getting Started with AWS App Mesh and Kubernetes

AWS App Mesh is a service mesh based on the [Envoy](#) proxy that makes it easy to monitor and control microservices. App Mesh standardizes how your microservices communicate, giving you end-to-end visibility and helping to ensure high availability for your applications.

App Mesh gives you consistent visibility and network traffic controls for every microservice in an application. For more information, see the [App Mesh User Guide](#).

This topic helps you to use AWS App Mesh with an existing microservice application running on Amazon EKS or Kubernetes on Amazon EC2. You can either integrate Kubernetes with App Mesh resources by completing the steps in this topic, or by installing the App Mesh Kubernetes integration components. The integration components automatically complete the tasks in this topic for you, enabling you to integrate with App Mesh directly from Kubernetes. For more information, see [Configure App Mesh Integration with Kubernetes](#).

## Prerequisites

App Mesh supports microservice applications that use service discovery naming for their components. To use this getting started guide, you must have a microservice application running on Amazon EKS or Kubernetes on AWS.

Kubernetes `kube-dns` and `coredns` are supported. For more information, see [DNS for Services and Pods](#) in the Kubernetes documentation.

## Step 1: Create Your Service Mesh

A service mesh is a logical boundary for network traffic between the services that reside within it. For more information, see [Service Meshes](#) in the *AWS App Mesh User Guide*.

After you create your service mesh, you can create virtual services, virtual nodes, virtual routers, and routes to distribute traffic between the applications in your mesh.

### To create a new service mesh with the AWS Management Console

1. Open the App Mesh console at <https://console.aws.amazon.com/appmesh/>.
2. Choose **Create mesh**.
3. For **Mesh name**, specify a name for your service mesh.
4. Choose **Create mesh** to finish.

## Step 2: Create Your Virtual Nodes

A virtual node acts as a logical pointer to a particular task group, such as a Kubernetes deployment. For more information, see [Virtual Nodes](#) in the *AWS App Mesh User Guide*.

When you create a virtual node, you must specify the DNS service discovery hostname for your task group. Any inbound traffic that your virtual node expects should be specified as a *listener*. Any outbound traffic that your virtual node expects to reach should be specified as a *backend*.

You must create virtual nodes for each microservice in your application.

**To create a virtual node in the AWS Management Console.**

1. Choose the mesh that you created in the previous steps.
2. Choose **Virtual nodes** in the left navigation.
3. Choose **Create virtual node**.
4. For **Virtual node name**, choose a name for your virtual node.
5. For **Service discovery method**, choose **DNS** for services that use DNS service discovery and then specify the hostname for **DNS hostname**. Otherwise, choose **None** if your virtual node doesn't expect any ingress traffic.
6. To specify any backends (for egress traffic) for your virtual node, or to configure inbound and outbound access logging information, choose **Additional configuration**.
  - a. To specify a backend, choose **Add backend** and enter a virtual service name or full Amazon Resource Name (ARN) for the virtual service that your virtual node communicates with. Repeat this step until all of your virtual node backends are accounted for.
  - b. To configure logging, enter the HTTP access logs path that you want Envoy to use. We recommend the `/dev/stdout` path so that you can use Docker log drivers to export your Envoy logs to a service such as Amazon CloudWatch Logs.

**Note**

Logs must still be ingested by an agent in your application and sent to a destination. This file path only instructs Envoy where to send the logs.

7. If your virtual node expects ingress traffic, specify a **Port** and **Protocol** for that **Listener**.
8. If you want to configure health checks for your listener, ensure that **Health check enabled** is selected and then complete the following substeps. If not, clear this check box.
  - a. For **Health check protocol**, choose to use an HTTP or TCP health check.
  - b. For **Health check port**, specify the port that the health check should run on.
  - c. For **Healthy threshold**, specify the number of consecutive successful health checks that must occur before declaring the listener healthy.
  - d. For **Health check interval**, specify the time period in milliseconds between each health check execution.
  - e. For **Path**, specify the destination path for the health check request. This is required only if the specified protocol is HTTP. If the protocol is TCP, this parameter is ignored.
  - f. For **Timeout period**, specify the amount of time to wait when receiving a response from the health check, in milliseconds.
  - g. For **Unhealthy threshold**, specify the number of consecutive failed health checks that must occur before declaring the listener unhealthy.
9. Chose **Create virtual node** to finish.
10. Repeat this procedure as necessary to create virtual nodes for each remaining microservice in your application.

## Step 3: Create Your Virtual Routers

Virtual routers handle traffic for one or more virtual services within your mesh. After you create a virtual router, you can create and associate routes for your virtual router that direct incoming requests to different virtual nodes. For more information, see [Virtual Routers](#) in the *AWS App Mesh User Guide*.

Create virtual routers for each microservice in your application.

#### Creating a virtual router in the AWS Management Console.

1. Choose **Virtual routers** in the left navigation.
2. Choose **Create virtual router**.
3. For **Virtual router name**, specify a name for your virtual router. Up to 255 letters, numbers, hyphens, and underscores are allowed.
4. For **Listener**, specify a **Port** and **Protocol** for your virtual router.
5. Choose **Create virtual router** to finish.
6. Repeat this procedure as necessary to create virtual routers for each remaining microservice in your application.

## Step 4: Create Your Routes

A route is associated with a virtual router, and it's used to match requests for a virtual router and distribute traffic accordingly to its associated virtual nodes. For more information, see [Routes](#) in the *AWS App Mesh User Guide*.

Create routes for each microservice in your application.

#### Creating a route in the AWS Management Console.

1. Choose **Virtual routers** in the left navigation.
2. Choose the router that you want to associate a new route with.
3. In the **Routes** table, choose **Create route**.
4. For **Route name**, specify the name to use for your route.
5. For **Route type**, choose the protocol for your route.
6. For **Virtual node name**, choose the virtual node that this route will serve traffic to.
7. For **Weight**, choose a relative weight for the route. The total weight for all routes must be less than 100.
8. To use HTTP path-based routing, choose **Additional configuration** and then specify the path that the route should match. For example, if your virtual service name is `my-service.local` and you want the route to match requests to `my-service.local/metrics`, your prefix should be `/metrics`.
9. Choose **Create route** to finish.
10. Repeat this procedure as necessary to create routes for each remaining microservice in your application.

## Step 5: Create Your Virtual Services

A virtual service is an abstraction of a real service that is provided by a virtual node directly or indirectly by means of a virtual router. Dependent services call your virtual service by its `virtualServiceName`, and those requests are routed to the virtual node or virtual router that is specified as the provider for the virtual service. For more information, see [Virtual Services](#) in the *AWS App Mesh User Guide*.

Create virtual services for each microservice in your application.

#### Creating a virtual service in the AWS Management Console.

1. Choose **Virtual services** in the left navigation.

2. Choose **Create virtual service**.
3. For **Virtual service name**, choose a name for your virtual service. We recommend that you use the service discovery name of the real service that you're targeting (such as `my-service.default.svc.cluster.local`).
4. For **Provider**, choose the provider type for your virtual service:
  - If you want the virtual service to spread traffic across multiple virtual nodes, select **Virtual router** and then choose the virtual router to use from the drop-down menu.
  - If you want the virtual service to reach a virtual node directly, without a virtual router, select **Virtual node** and then choose the virtual node to use from the drop-down menu.
  - If you don't want the virtual service to route traffic at this time (for example, if your virtual nodes or virtual router doesn't exist yet), choose **None**. You can update the provider for this virtual service later.
5. Choose **Create virtual service** to finish.
6. Repeat this procedure as necessary to create virtual services for each remaining microservice in your application.

## Step 6: Updating Your Microservice Pod Specifications

App Mesh is a service mesh based on the [Envoy](#) proxy. After you create your service mesh, virtual services, virtual nodes, virtual routers, and routes, you must update your microservices to be compatible with App Mesh.

App Mesh vends the following custom container images that you must add to your Kubernetes pod specifications:

- App Mesh Envoy container image – `111345817488.dkr.ecr.us-west-2.amazonaws.com/aws-appmesh-envoy:v1.9.1.0-prod`. Envoy uses the configuration defined in the App Mesh control plane to determine where to send your application traffic.

You must use the App Mesh Envoy container image until the Envoy project team merges changes that support App Mesh. For additional details, see the [GitHub roadmap issue](#).

- App Mesh proxy route manager – `111345817488.dkr.ecr.us-west-2.amazonaws.com/aws-appmesh-proxy-route-manager:v2`. The route manager sets up a pod's network namespace with iptables rules that route ingress and egress traffic through Envoy.

The following text is an example Kubernetes pod specification that you can merge with your existing application. Substitute your mesh name and virtual node name for the `APPMESH_VIRTUAL_NODE_NAME` value, and a list of ports that your application listens on for the `APPMESH_APP_PORTS` value. Substitute the Amazon EC2 instance AWS Region for the `AWS_REGION` value.

Update each microservice pod specification in your application to include these containers, and then deploy the new specifications to update your microservices and start using App Mesh with your Kubernetes application.

### Example Kubernetes pod spec

```
spec:
  containers:
    - name: envoy
      image: 111345817488.dkr.ecr.us-west-2.amazonaws.com/aws-appmesh-envoy:v1.9.1.0-prod
      securityContext:
```

```
    runAsUser: 1337
  env:
    - name: "APPMESH_VIRTUAL_NODE_NAME"
      value: "mesh/meshName/virtualNode/virtualNodeName"
    - name: "ENVOY_LOG_LEVEL"
      value: "info"
    - name: "AWS_REGION"
      value: "aws_region_name"
  initContainers:
    - name: proxyinit
      image: 111345817488.dkr.ecr.us-west-2.amazonaws.com/aws-appmesh-proxy-route-
manager:v2
  securityContext:
    capabilities:
      add:
        - NET_ADMIN
  env:
    - name: "APPMESH_START_ENABLED"
      value: "1"
    - name: "APPMESH_IGNORE_UID"
      value: "1337"
    - name: "APPMESH_ENVOY_INGRESS_PORT"
      value: "15000"
    - name: "APPMESH_ENVOY_EGRESS_PORT"
      value: "15001"
    - name: "APPMESH_APP_PORTS"
      value: "application_port_list"
    - name: "APPMESH_EGRESS_IGNORED_IP"
      value: "169.254.169.254"
```



# Tutorial: Configure App Mesh Integration with Kubernetes

AWS App Mesh is a service mesh based on the [Envoy](#) proxy that makes it easy to monitor and control microservices. App Mesh standardizes how your microservices communicate, giving you end-to-end visibility and helping to ensure high availability for your applications.

App Mesh gives you consistent visibility and network traffic controls for every microservice in an application. For more information, see the [App Mesh User Guide](#).

When you use AWS App Mesh with Kubernetes, you manage App Mesh resources, such as virtual services and virtual nodes, that align to Kubernetes resources, such as services and deployments. You also add the App Mesh sidecar container images to Kubernetes pod specifications. This tutorial guides you through the installation of the following open source components that automatically complete these tasks for you when you work with Kubernetes resources:

- **App Mesh controller for Kubernetes** – The controller is accompanied by the deployment of three Kubernetes custom resource definitions: `mesh`, `virtual service`, and `virtual node`. The controller watches for creation, modification, and deletion of the custom resources and makes changes to the corresponding App Mesh `mesh`, `virtual service` (including `virtual router` and `route`), and `virtual node` resources through the App Mesh API. To learn more or contribute to the controller, see the [GitHub project](#).
- **App Mesh sidecar injector for Kubernetes** – The injector installs as a webhook and injects the [App Mesh sidecar container images](#) into Kubernetes pods running in specific, labeled namespaces. To learn more or contribute, see the [GitHub project](#).

*The features discussed in this topic are available as an open-source beta. This means that these features are well tested. Support for the features will not be dropped, though details may change. If the schema or schematics of a feature changes, instructions for migrating to the next version will be provided. This migration may require deleting, editing, and re-creating Kubernetes API objects.*

## Prerequisites

To use the controller and sidecar injector, you must have the following resources:

- An existing Kubernetes cluster running version 1.11 or later. If you don't have an existing cluster, you can deploy one using the [Getting Started with Amazon EKS](#) guide.
- A `kubectl` client that is configured to communicate with your Kubernetes cluster. If you're using Amazon Elastic Kubernetes Service, you can use the instructions for installing `kubectl` and configuring a `kubeconfig` file.
- `jq` and Open SSL installed.

## Step 1: Install the Controller and Custom Resources

To install the controller and Kubernetes custom resource definitions, complete the following steps.

1. The controller requires that your account and your Kubernetes worker nodes are able to work with App Mesh resources. [Attach](#) the [AWSAppMeshFullAccess](#) policy to the role that is attached to your Kubernetes worker nodes. If you are using a pod identity solution, make sure that the controller pod is bound to the policy.
2. To create the Kubernetes custom resources and launch the controller, download the following yaml file and apply it to your cluster with the following command.

```
curl https://raw.githubusercontent.com/aws/aws-app-mesh-controller-for-k8s/v0.1.1/
deploy/all.yaml | kubectl apply -f -
```

A Kubernetes namespace named `appmesh-system` is created and a container running the controller is deployed into the namespace.

3. Confirm that the controller is running with the following command.

```
kubectl rollout status deployment app-mesh-controller -n appmesh-system
```

If the controller is running, the following output is returned.

```
deployment "app-mesh-controller" successfully rolled out
```

4. Confirm that the Kubernetes custom resources for App Mesh were created with the following command.

```
kubectl get crd
```

If the custom resources were created, output similar to the following is returned.

NAME	CREATED AT
meshes.appmesh.k8s.aws	2019-05-08T14:17:26Z
virtualnodes.appmesh.k8s.aws	2019-05-08T14:17:26Z
virtualservices.appmesh.k8s.aws	2019-05-08T14:17:26Z

## Step 2: Install the Sidecar Injector

To install the sidecar injector, complete the following steps.

1. Export the name of the mesh you want to create with the following command.

```
export MESH_NAME=my-mesh
```

2. Download and execute the sidecar injector installation script with the following command.

```
curl https://raw.githubusercontent.com/aws/aws-app-mesh-inject/v0.1.4/scripts/
install.sh | bash
```

A Kubernetes namespace named `appmesh-inject` was created and a container running the injector was deployed into the namespace. If the injector successfully installed, the last several lines of the output returned are similar to the following text.

```
deployment.apps/aws-app-mesh-inject configured
mutatingwebhookconfiguration.admissionregistration.k8s.io/aws-app-mesh-inject
configured
```

```
waiting for aws-app-mesh-inject to start
deployment "aws-app-mesh-inject" successfully rolled out
Mesh name has been set up
App Mesh image has been set up
The injector is ready
```

## Step 3: Configure App Mesh

When you deploy an application in Kubernetes, you also create the Kubernetes custom resources so that the controller can create the corresponding App Mesh resources. Additionally, you must enable sidecar injection so that the [App Mesh sidecar container images](#) are deployed in each Kubernetes pod. If you'd like to see the controller and injector in action, see [the section called "Deploy a Mesh Connected Service" \(p. 152\)](#).

### Create Kubernetes Custom Resources

You can deploy mesh, virtual service, and virtual node custom resources in Kubernetes, which then triggers the controller to create the corresponding resources in App Mesh through the App Mesh API.

#### Create a Mesh

When you create a mesh custom resource, you trigger the creation of an App Mesh mesh. The mesh name that you specify must be the same as the mesh name you exported when you [installed the sidecar injector \(p. 149\)](#). If the mesh name that you specify already exists, a new mesh is not created.

```
apiVersion: appmesh.k8s.aws/v1beta1
kind: Mesh
metadata:
  name: my-mesh
```

#### Create a Virtual Service

When you create a virtual service custom resource, you trigger the creation of an App Mesh virtual service, virtual router, and one or more routes containing a route configuration. The virtual service allows requests from one application in the mesh to be routed to a number of virtual nodes that make up a service.

```
apiVersion: appmesh.k8s.aws/v1beta1
kind: VirtualService
metadata:
  name: my-svc-a
  namespace: my-namespace
spec:
  meshName: my-mesh
  routes:
    - name: route-to-svc-a
      http:
        match:
          prefix: /
        action:
          weightedTargets:
            - virtualNodeName: my-app-a
              weight: 1
```

## Create a Virtual Node

When you create a virtual node custom resource, you trigger the creation of an App Mesh virtual node. The virtual node contains listener, back-end, and service discovery configuration.

```
apiVersion: appmesh.k8s.aws/v1beta1
kind: VirtualNode
metadata:
  name: my-app-a
  namespace: my-namespace
spec:
  meshName: my-mesh
  listeners:
  - portMapping:
      port: 9000
      protocol: http
  serviceDiscovery:
    dns:
      hostname: my-app-a.my-namespace.svc.cluster.local
  backends:
  - virtualService:
      virtualServiceName: my-svc-a
```

## Sidecar Injection

You enable sidecar injection for a Kubernetes namespace. When necessary, you can override the injector's default behavior for each pod you deploy in a Kubernetes namespace that you've enabled the injector for.

### Enable Sidecar Injection for a Namespace

To enable the sidecar injector for a Kubernetes namespace, label the namespace with the following command.

```
kubectl label namespace my-namespace appmesh.k8s.aws/sidecarInjectorWebhook=enabled
```

The [App Mesh sidecar container images](#) will be automatically injected into each pod that you deploy into the namespace.

### Override Sidecar Injector Default Behavior

To override the default behavior of the injector when deploying a pod in a namespace that you've enabled the injector for, add any of the following annotations to your pod spec.

- `appmesh.k8s.aws/mesh: mesh-name` – Add when you want to use a different mesh name than the one that you specified when you installed the injector.
- `appmesh.k8s.aws/ports: "ports"` – Specify particular ports when you don't want all of the container ports defined in a pod spec passed to the sidecars as application ports.
- `appmesh.k8s.aws/egressIgnoredPorts: ports` – Specify a comma separated list of port numbers for outbound traffic that you want ignored. By default all outbound traffic ports will be routed, except port 22 (SSH).
- `appmesh.k8s.aws/virtualNode: virtual-node-name` – Specify your own name if you don't want the virtual node name passed to the sidecars to be `<deployment name>--<namespace>`.
- `appmesh.k8s.aws/sidecarInjectorWebhook: disabled` – Add when you don't want the injector enabled for a pod.

```
apiVersion: appmesh.k8s.aws/v1beta1
kind: Deployment
spec:
  metadata:
    annotations:
      appmesh.k8s.aws/mesh: my-mesh2
      appmesh.k8s.aws/ports: "8079,8080"
      appmesh.k8s.aws/egressIgnoredPorts: "3306"
      appmesh.k8s.aws/virtualNode: my-app
      appmesh.k8s.aws/sidecarInjectorWebhook: disabled
```

## Step 4: Remove Integration Components (Optional)

If you need to remove the Kubernetes integration components, run the following commands.

```
kubectl delete crd meshes.appmesh.k8s.aws
kubectl delete crd virtualnodes.appmesh.k8s.aws
kubectl delete crd virtualservices.appmesh.k8s.aws
kubectl delete namespace appmesh-system
kubectl delete namespace appmesh-inject
```

## Deploy a Mesh Connected Service

In this topic, you deploy a sample application on Kubernetes. The application deploys mesh, virtual service, and virtual node Kubernetes custom resources. Kubernetes automatically creates mesh, virtual service, and virtual node resources in App Mesh and injects the App Mesh sidecar images into Kubernetes pods.

### Prerequisites

Before you deploy the sample application, you must meet the following prerequisites:

- Meet all of the prerequisites in [Tutorial: Configure App Mesh Integration with Kubernetes \(p. 148\)](#).
- Have the App Mesh controller for Kubernetes and the App Mesh sidecar injector for Kubernetes installed and configured. When you install the sidecar injector, specify *color-mesh* as the name of your mesh. To learn more about the controller and sidecar injector and how to install and configure them, see [Tutorial: Configure App Mesh Integration with Kubernetes \(p. 148\)](#).

### Deploy a Sample Application

The sample application consists of two components:

- **ColorGateway** – A simple http service written in Go that is exposed to external clients and that responds to *http://service-name:port/color*. The gateway responds with a color retrieved from *color-teller* and a histogram of colors observed at the server that responded up to the point when you made the request.
- **ColorTeller** – A simple http service written in Go that is configured to return a color. Multiple variants of the service are deployed. Each service is configured to return a specific color.

1. To deploy the color mesh sample application, download the following file and apply it to your Kubernetes cluster with the following command.

```
curl https://raw.githubusercontent.com/aws/aws-app-mesh-controller-for-k8s/v0.1.0/examples/color.yaml | kubectl apply -f -
```

2. View the resources deployed by the sample application with the following command.

```
kubectl -n appmesh-demo get all
```

In the output, you see a collection of virtual services, virtual nodes, and mesh custom resources along with native Kubernetes deployments, pods, and services. Your output will be similar to the following output.

NAME	READY	STATUS	RESTARTS	AGE
pod/colorgateway-cc6464d75-4ktj4	2/2	Running	0	37s
pod/colorteller-86664b5956-6h26c	2/2	Running	0	36s
pod/colorteller-black-6787756c7b-dw82f	2/2	Running	0	36s
pod/colorteller-blue-55d6f99dc6-f5wgd	2/2	Running	0	36s
pod/colorteller-red-578866ffb-x9m7w	2/2	Running	0	35s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/colorgateway	ClusterIP	10.100.21.147	<none>	9080/TCP	37s
service/colorteller	ClusterIP	10.100.187.50	<none>	9080/TCP	37s
service/colorteller-black	ClusterIP	10.100.61.36	<none>	9080/TCP	36s
service/colorteller-blue	ClusterIP	10.100.254.230	<none>	9080/TCP	36s
service/colorteller-red	ClusterIP	10.100.90.38	<none>	9080/TCP	36s

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/colorgateway	1	1	1	1	37s
deployment.apps/colorteller	1	1	1	1	36s
deployment.apps/colorteller-black	1	1	1	1	36s
deployment.apps/colorteller-blue	1	1	1	1	36s
deployment.apps/colorteller-red	1	1	1	1	36s

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/colorgateway-cc6464d75	1	1	1	37s
replicaset.apps/colorteller-86664b5956	1	1	1	36s
replicaset.apps/colorteller-black-6787756c7b	1	1	1	36s
replicaset.apps/colorteller-blue-55d6f99dc6	1	1	1	36s
replicaset.apps/colorteller-red-578866ffb	1	1	1	35s

NAME	AGE
virtualservice.appmesh.k8s.aws/colorgateway.appmesh-demo	37s
virtualservice.appmesh.k8s.aws/colorteller.appmesh-demo	37s

NAME	AGE
mesh.appmesh.k8s.aws/color-mesh	38s

NAME	AGE
virtualnode.appmesh.k8s.aws/colorgateway	39s
virtualnode.appmesh.k8s.aws/colorteller	39s
virtualnode.appmesh.k8s.aws/colorteller-black	39s
virtualnode.appmesh.k8s.aws/colorteller-blue	39s
virtualnode.appmesh.k8s.aws/colorteller-red	38s

You can use the AWS Management Console or AWS CLI to see the App Mesh mesh, virtual service, virtual router, route, and virtual node resources that were automatically created by the controller. All of the resources were deployed to the `appmesh-demo` namespace, which was labelled with `appmesh.k8s.aws/sidecarInjectorWebhook: enabled`. Since the injector saw this label for the namespace, it injected the App Mesh sidecar container images into each of the

Pods. Using `kubectl describe pod <pod-name> -n appmesh-demo`, you can see that the [App Mesh sidecar container images](#) are included in each of the pods that were deployed.

## Run Application

Complete the following steps to run the application.

1. In a terminal, use the following command to create a container in the `appmesh-demo` namespace that has `curl` installed and open a shell to it. In later steps, this terminal is referred to as *Terminal A*.

```
kubectl run -n appmesh-demo -it curler --image=tutum/curl /bin/bash
```

2. From *Terminal A*, run the following command to curl the color gateway in the color mesh application 100 times. The gateway routes traffic to separate virtual nodes that return either white, black, or blue as a response.

```
for i in {1..100}; do curl colorgateway:9080/color; echo; done
```

100 responses are returned. Each response looks similar to the following text:

```
{"color":"blue", "stats": {"black":0.36,"blue":0.32,"white":0.32}}
```

In this line of output, the colorgateway routed the request to the blue virtual node. The numbers for each color denote the percentage of responses from each virtual node. The number for each color in each response is cumulative over time. The percentage is similar for each color because, by default, the weighting defined for each virtual node is the same in the `color.yaml` file you used to install the sample application.

Leave *Terminal A* open.

## Change Configuration

Change the configuration and run the application again to see the effect of the changes.

1. In a separate terminal from *Terminal A*, edit the `colorteller.appmesh-demo` virtual service with the following command.

```
kubectl edit VirtualService colorteller.appmesh-demo -n appmesh-demo
```

In the editor, you can see that the `weight` value of each `virtualNodeName` is `1`. Because the weight of each virtual node is the same, traffic routed to each virtual node is approximately even. To route all traffic to the black node only, change the values for `colorteller.appmesh-demo` and `colorteller-blue` to `0`, as shown in the following text. Save the configuration and exit the editor.

```
spec:
  meshName: color-mesh
  routes:
  - http:
      action:
        weightedTargets:
        - virtualNodeName: colorteller.appmesh-demo
          weight: 0
        - virtualNodeName: colorteller-blue
          weight: 0
```

```
- virtualNodeName: colorteller-black.appmesh-demo  
  weight: 1
```

2. In *Terminal A*, run `curl` again with the following command.

```
for i in {1..100}; do curl colorgateway:9080/color; echo; done
```

This time, all lines of output look similar to the following text.

```
{"color":"black", "stats": {"black":0.64,"blue":0.18,"white":0.19}}
```

Black is the response every time because the gateway is now routing all traffic to the black virtual node. Even though all traffic is now going to black, the white and blue virtual nodes still have response percentages, because the numbers are based on relative percentages over time. When you executed the requests in a previous step, white and blue responded, which is why they still have response percentages. You can see that the relative percentages decrease for white and blue with each response, while the percentage for black increases.

## Remove Application

When you've finished with the sample application, you can remove it by completing the following steps.

1. Use the following commands to remove the sample application and the App Mesh resources that were created.

```
kubectl delete namespace appmesh-demo  
kubectl delete mesh color-mesh
```

2. Optional: If you want to remove the controller and sidecar injector, see [Remove integration components](#) (p. ).



# Deep Learning Containers

AWS Deep Learning Containers are a set of Docker images for training and serving models in TensorFlow on Amazon EKS and Amazon Elastic Container Service (Amazon ECR). Deep Learning Containers provide optimized environments with TensorFlow, Nvidia CUDA (for GPU instances), and Intel MKL (for CPU instances) libraries and are available in Amazon ECR.

To get started using AWS Deep Learning Containers on Amazon EKS, see [AWS Deep Learning Containers on Amazon EKS](#) in the *AWS Deep Learning AMI Developer Guide*.

# Security in Amazon EKS

Cloud security at AWS is the highest priority. As an AWS customer, you benefit from a data center and network architecture that is built to meet the requirements of the most security-sensitive organizations.

Security is a shared responsibility between AWS and you. The [shared responsibility model](#) describes this as security *of* the cloud and security *in* the cloud:

- **Security of the cloud** – AWS is responsible for protecting the infrastructure that runs AWS services in the AWS Cloud. For Amazon EKS, AWS is responsible for the Kubernetes control plane, which includes the control plane nodes and etcd database. Third-party auditors regularly test and verify the effectiveness of our security as part of the [AWS compliance programs](#). To learn about the compliance programs that apply to Amazon EKS, see [AWS Services in Scope by Compliance Program](#).
- **Security in the cloud** – Your responsibility includes the following areas.
  - The security configuration of the data plane, including the configuration of the security groups that allow traffic to pass from the Amazon EKS control plane into the customer VPC
  - The configuration of the worker nodes and the containers themselves
  - The worker node guest operating system (including updates and security patches)
  - Other associated application software:
    - Setting up and managing network controls, such as firewall rules
    - Managing platform-level identity and access management, either with or in addition to IAM
  - The sensitivity of your data, your company's requirements, and applicable laws and regulations

This documentation helps you understand how to apply the shared responsibility model when using Amazon EKS. The following topics show you how to configure Amazon EKS to meet your security and compliance objectives. You also learn how to use other AWS services that help you to monitor and secure your Amazon EKS resources.

## Topics

- [Identity and Access Management for Amazon EKS \(p. 157\)](#)
- [Logging and Monitoring in Amazon EKS \(p. 170\)](#)
- [Compliance Validation for Amazon EKS \(p. 170\)](#)
- [Resilience in Amazon EKS \(p. 171\)](#)
- [Infrastructure Security in Amazon EKS \(p. 171\)](#)
- [Configuration and Vulnerability Analysis in Amazon EKS \(p. 172\)](#)

## Identity and Access Management for Amazon EKS

AWS Identity and Access Management (IAM) is an AWS service that helps an administrator securely control access to AWS resources. IAM administrators control who can be *authenticated* (signed in) and *authorized* (have permissions) to use Amazon EKS resources. IAM is an AWS service that you can use with no additional charge.

## Topics

- [Audience \(p. 158\)](#)
- [Authenticating With Identities \(p. 158\)](#)
- [Managing Access Using Policies \(p. 160\)](#)

- [How Amazon EKS Works with IAM \(p. 161\)](#)
- [Amazon EKS Identity-Based Policy Examples \(p. 163\)](#)
- [Amazon EKS Service IAM Role \(p. 166\)](#)
- [Amazon EKS Worker Node IAM Role \(p. 168\)](#)
- [Troubleshooting Amazon EKS Identity and Access \(p. 170\)](#)

## Audience

How you use AWS Identity and Access Management (IAM) differs, depending on the work you do in Amazon EKS.

**Service user** – If you use the Amazon EKS service to do your job, then your administrator provides you with the credentials and permissions that you need. As you use more Amazon EKS features to do your work, you might need additional permissions. Understanding how access is managed can help you request the right permissions from your administrator. If you cannot access a feature in Amazon EKS, see [Troubleshooting Amazon EKS Identity and Access \(p. 170\)](#).

**Service administrator** – If you're in charge of Amazon EKS resources at your company, you probably have full access to Amazon EKS. It's your job to determine which Amazon EKS features and resources your employees should access. You must then submit requests to your IAM administrator to change the permissions of your service users. Review the information on this page to understand the basic concepts of IAM. To learn more about how your company can use IAM with Amazon EKS, see [How Amazon EKS Works with IAM \(p. 161\)](#).

**IAM administrator** – If you're an IAM administrator, you might want to learn details about how you can write policies to manage access to Amazon EKS. To view example Amazon EKS identity-based policies that you can use in IAM, see [Amazon EKS Identity-Based Policy Examples \(p. 163\)](#).

## Authenticating With Identities

Authentication is how you sign in to AWS using your identity credentials. For more information about signing in using the AWS Management Console, see [The IAM Console and Sign-in Page](#) in the *IAM User Guide*.

You must be *authenticated* (signed in to AWS) as the AWS account root user, an IAM user, or by assuming an IAM role. You can also use your company's single sign-on authentication, or even sign in using Google or Facebook. In these cases, your administrator previously set up identity federation using IAM roles. When you access AWS using credentials from another company, you are assuming a role indirectly.

To sign in directly to the [AWS Management Console](#), use your password with your root user email or your IAM user name. You can access AWS programmatically using your root user or IAM user access keys. AWS provides SDK and command line tools to cryptographically sign your request using your credentials. If you don't use AWS tools, you must sign the request yourself. Do this using *Signature Version 4*, a protocol for authenticating inbound API requests. For more information about authenticating requests, see [Signature Version 4 Signing Process](#) in the *AWS General Reference*.

Regardless of the authentication method that you use, you might also be required to provide additional security information. For example, AWS recommends that you use multi-factor authentication (MFA) to increase the security of your account. To learn more, see [Using Multi-Factor Authentication \(MFA\) in AWS](#) in the *IAM User Guide*.

## AWS Account Root User

When you first create an AWS account, you begin with a single sign-in identity that has complete access to all AWS services and resources in the account. This identity is called the AWS account *root user* and

is accessed by signing in with the email address and password that you used to create the account. We strongly recommend that you do not use the root user for your everyday tasks, even the administrative ones. Instead, adhere to the [best practice of using the root user only to create your first IAM user](#). Then securely lock away the root user credentials and use them to perform only a few account and service management tasks.

## IAM Users and Groups

An *IAM user* is an identity within your AWS account that has specific permissions for a single person or application. An IAM user can have long-term credentials such as a user name and password or a set of access keys. To learn how to generate access keys, see [Managing Access Keys for IAM Users](#) in the *IAM User Guide*. When you generate access keys for an IAM user, make sure you view and securely save the key pair. You cannot recover the secret access key in the future. Instead, you must generate a new access key pair.

An *IAM group* is an identity that specifies a collection of IAM users. You can't sign in as a group. You can use groups to specify permissions for multiple users at a time. Groups make permissions easier to manage for large sets of users. For example, you could have a group named *IAMAdmins* and give that group permissions to administer IAM resources.

Users are different from roles. A user is uniquely associated with one person or application, but a role is intended to be assumable by anyone who needs it. Users have permanent long-term credentials, but roles provide temporary credentials. To learn more, see [When to Create an IAM User \(Instead of a Role\)](#) in the *IAM User Guide*.

## IAM Roles

An *IAM role* is an identity within your AWS account that has specific permissions. It is similar to an IAM user, but is not associated with a specific person. You can temporarily assume an IAM role in the AWS Management Console by [switching roles](#). You can assume a role by calling an AWS CLI or AWS API operation or by using a custom URL. For more information about methods for using roles, see [Using IAM Roles](#) in the *IAM User Guide*.

IAM roles with temporary credentials are useful in the following situations:

- **Temporary IAM user permissions** – An IAM user can assume an IAM role to temporarily take on different permissions for a specific task.
- **Federated user access** – Instead of creating an IAM user, you can use existing identities from AWS Directory Service, your enterprise user directory, or a web identity provider. These are known as *federated users*. AWS assigns a role to a federated user when access is requested through an [identity provider](#). For more information about federated users, see [Federated Users and Roles](#) in the *IAM User Guide*.
- **Cross-account access** – You can use an IAM role to allow someone (a trusted principal) in a different account to access resources in your account. Roles are the primary way to grant cross-account access. However, with some AWS services, you can attach a policy directly to a resource (instead of using a role as a proxy). To learn the difference between roles and resource-based policies for cross-account access, see [How IAM Roles Differ from Resource-based Policies](#) in the *IAM User Guide*.
- **AWS service access** – A service role is an IAM role that a service assumes to perform actions in your account on your behalf. When you set up some AWS service environments, you must define a role for the service to assume. This service role must include all the permissions that are required for the service to access the AWS resources that it needs. Service roles vary from service to service, but many allow you to choose your permissions as long as you meet the documented requirements for that service. Service roles provide access only within your account and cannot be used to grant access to services in other accounts. You can create, modify, and delete a service role from within IAM. For example, you can create a role that allows Amazon Redshift to access an Amazon S3 bucket on your behalf and then load data from that bucket into an Amazon Redshift cluster. For more information, see [Creating a Role to Delegate Permissions to an AWS Service](#) in the *IAM User Guide*.

- **Applications running on Amazon EC2** – You can use an IAM role to manage temporary credentials for applications that are running on an EC2 instance and making AWS CLI or AWS API requests. This is preferable to storing access keys within the EC2 instance. To assign an AWS role to an EC2 instance and make it available to all of its applications, you create an instance profile that is attached to the instance. An instance profile contains the role and enables programs that are running on the EC2 instance to get temporary credentials. For more information, see [Using an IAM Role to Grant Permissions to Applications Running on Amazon EC2 Instances](#) in the *IAM User Guide*.

To learn whether to use IAM roles, see [When to Create an IAM Role \(Instead of a User\)](#) in the *IAM User Guide*.

## Managing Access Using Policies

You control access in AWS by creating policies and attaching them to IAM identities or AWS resources. A policy is an object in AWS that, when associated with an identity or resource, defines their permissions. AWS evaluates these policies when an entity (root user, IAM user, or IAM role) makes a request. Permissions in the policies determine whether the request is allowed or denied. Most policies are stored in AWS as JSON documents. For more information about the structure and contents of JSON policy documents, see [Overview of JSON Policies](#) in the *IAM User Guide*.

An IAM administrator can use policies to specify who has access to AWS resources, and what actions they can perform on those resources. Every IAM entity (user or role) starts with no permissions. In other words, by default, users can do nothing, not even change their own password. To give a user permission to do something, an administrator must attach a permissions policy to a user. Or the administrator can add the user to a group that has the intended permissions. When an administrator gives permissions to a group, all users in that group are granted those permissions.

IAM policies define permissions for an action regardless of the method that you use to perform the operation. For example, suppose that you have a policy that allows the `iam:GetRole` action. A user with that policy can get role information from the AWS Management Console, the AWS CLI, or the AWS API.

## Identity-Based Policies

Identity-based policies are JSON permissions policy documents that you can attach to an identity, such as an IAM user, role, or group. These policies control what actions that identity can perform, on which resources, and under what conditions. To learn how to create an identity-based policy, see [Creating IAM Policies](#) in the *IAM User Guide*.

Identity-based policies can be further categorized as *inline policies* or *managed policies*. Inline policies are embedded directly into a single user, group, or role. Managed policies are standalone policies that you can attach to multiple users, groups, and roles in your AWS account. Managed policies include AWS managed policies and customer managed policies. To learn how to choose between a managed policy or an inline policy, see [Choosing Between Managed Policies and Inline Policies](#) in the *IAM User Guide*.

## Resource-Based Policies

Resource-based policies are JSON policy documents that you attach to a resource such as an Amazon S3 bucket. Service administrators can use these policies to define what actions a specified principal (account member, user, or role) can perform on that resource and under what conditions. Resource-based policies are inline policies. There are no managed resource-based policies.

## Access Control Lists (ACLs)

Access control policies (ACLs) control which principals (account members, users, or roles) have permissions to access a resource. ACLs are similar to resource-based policies, although they are the only

policy type that does not use the JSON policy document format. Amazon S3, AWS WAF, and Amazon VPC are examples of services that support ACLs. To learn more about ACLs, see [Access Control List \(ACL\) Overview](#) in the *Amazon Simple Storage Service Developer Guide*.

## Other Policy Types

AWS supports additional, less-common policy types. These policy types can set the maximum permissions granted to you by the more common policy types.

- **Permissions boundaries** – A permissions boundary is an advanced feature in which you set the maximum permissions that an identity-based policy can grant to an IAM entity (IAM user or role). You can set a permissions boundary for an entity. The resulting permissions are the intersection of entity's identity-based policies and its permissions boundaries. Resource-based policies that specify the user or role in the `Principal` field are not limited by the permissions boundary. An explicit deny in any of these policies overrides the allow. For more information about permissions boundaries, see [Permissions Boundaries for IAM Entities](#) in the *IAM User Guide*.
- **Service control policies (SCPs)** – SCPs are JSON policies that specify the maximum permissions for an organization or organizational unit (OU) in AWS Organizations. AWS Organizations is a service for grouping and centrally managing multiple AWS accounts that your business owns. If you enable all features in an organization, then you can apply service control policies (SCPs) to any or all of your accounts. The SCP limits permissions for entities in member accounts, including each AWS account root user. For more information about Organizations and SCPs, see [How SCPs Work](#) in the *AWS Organizations User Guide*.
- **Session policies** – Session policies are advanced policies that you pass as a parameter when you programmatically create a temporary session for a role or federated user. The resulting session's permissions are the intersection of the user or role's identity-based policies and the session policies. Permissions can also come from a resource-based policy. An explicit deny in any of these policies overrides the allow. For more information, see [Session Policies](#) in the *IAM User Guide*.

## Multiple Policy Types

When multiple types of policies apply to a request, the resulting permissions are more complicated to understand. To learn how AWS determines whether to allow a request when multiple policy types are involved, see [Policy Evaluation Logic](#) in the *IAM User Guide*.

## How Amazon EKS Works with IAM

Before you use IAM to manage access to Amazon EKS, you should understand what IAM features are available to use with Amazon EKS. To get a high-level view of how Amazon EKS and other AWS services work with IAM, see [AWS Services That Work with IAM](#) in the *IAM User Guide*.

### Topics

- [Amazon EKS Identity-Based Policies](#) (p. 161)
- [Amazon EKS Resource-Based Policies](#) (p. 162)
- [Authorization Based on Amazon EKS Tags](#) (p. 162)
- [Amazon EKS IAM Roles](#) (p. 162)

## Amazon EKS Identity-Based Policies

With IAM identity-based policies, you can specify allowed or denied actions and resources as well as the conditions under which actions are allowed or denied. Amazon EKS supports specific actions, resources, and condition keys. To learn about all of the elements that you use in a JSON policy, see [IAM JSON Policy Elements Reference](#) in the *IAM User Guide*.

## Actions

The `Action` element of an IAM identity-based policy describes the specific action or actions that will be allowed or denied by the policy. Policy actions usually have the same name as the associated AWS API operation. The action is used in a policy to grant permissions to perform the associated operation.

Policy actions in Amazon EKS use the following prefix before the action: `eks:`. For example, to grant someone permission to get descriptive information about an Amazon EKS cluster, you include the `DescribeCluster` action in their policy. Policy statements must include either an `Action` or `NotAction` element.

To specify multiple actions in a single statement, separate them with commas as follows:

```
"Action": ["eks:action1", "eks:action2"]
```

You can specify multiple actions using wildcards (\*). For example, to specify all actions that begin with the word `Describe`, include the following action:

```
"Action": "eks:Describe*"
```

To see a list of Amazon EKS actions, see [Actions Defined by Amazon Elastic Kubernetes Service](#) in the *IAM User Guide*.

## Resources

Amazon EKS does not support specifying resource ARNs in a policy.

## Condition Keys

Amazon EKS does not provide any service-specific condition keys, but it does support using some global condition keys. To see all AWS global condition keys, see [AWS Global Condition Context Keys](#) in the *IAM User Guide*.

## Examples

To view examples of Amazon EKS identity-based policies, see [Amazon EKS Identity-Based Policy Examples \(p. 163\)](#).

When you create an Amazon EKS cluster, the IAM entity user or role, such as a [federated user](#) that creates the cluster, is automatically granted `system:masters` permissions in the cluster's RBAC configuration. To grant additional AWS users or roles the ability to interact with your cluster, you must edit the `aws-auth` ConfigMap within Kubernetes.

For additional information about working with the ConfigMap, see [Managing Users or IAM Roles for your Cluster \(p. 115\)](#).

## Amazon EKS Resource-Based Policies

Amazon EKS does not support resource-based policies.

## Authorization Based on Amazon EKS Tags

Amazon EKS does not support tagging resources or controlling access based on tags.

## Amazon EKS IAM Roles

An [IAM role](#) is an entity within your AWS account that has specific permissions.

## Using Temporary Credentials with Amazon EKS

You can use temporary credentials to sign in with federation, assume an IAM role, or to assume a cross-account role. You obtain temporary security credentials by calling AWS STS API operations such as [AssumeRole](#) or [GetFederationToken](#).

Amazon EKS supports using temporary credentials.

## Service-Linked Roles

Amazon EKS does not support service-linked roles.

## Service Roles

This feature allows a service to assume a [service role](#) on your behalf. This role allows the service to access resources in other services to complete an action on your behalf. Service roles appear in your IAM account and are owned by the account. This means that an IAM administrator can change the permissions for this role. However, doing so might break the functionality of the service.

Amazon EKS supports service roles. For more information, see [the section called "Service IAM Role" \(p. 166\)](#) and [the section called "Worker Node IAM Role" \(p. 168\)](#).

## Choosing an IAM Role in Amazon EKS

When you create a cluster resource in Amazon EKS, you must choose a role to allow Amazon EKS to access several other AWS resources on your behalf. If you have previously created a service role, then Amazon EKS provides you with a list of roles to choose from. It's important to choose a role that has the Amazon EKS managed policies attached to it. For more information, see [the section called "Check for an Existing Service Role" \(p. 166\)](#) and [the section called "Check for an Existing Worker Node Role" \(p. 168\)](#).

# Amazon EKS Identity-Based Policy Examples

By default, IAM users and roles don't have permission to create or modify Amazon EKS resources. They also can't perform tasks using the AWS Management Console, AWS CLI, or AWS API. An IAM administrator must create IAM policies that grant users and roles permission to perform specific API operations on the specified resources they need. The administrator must then attach those policies to the IAM users or groups that require those permissions.

To learn how to create an IAM identity-based policy using these example JSON policy documents, see [Creating Policies on the JSON Tab](#) in the *IAM User Guide*.

When you create an Amazon EKS cluster, the IAM entity user or role, such as a [federated user](#) that creates the cluster, is automatically granted `system:masters` permissions in the cluster's RBAC configuration. To grant additional AWS users or roles the ability to interact with your cluster, you must edit the `aws-auth` ConfigMap within Kubernetes.

For additional information about working with the ConfigMap, see [Managing Users or IAM Roles for your Cluster \(p. 115\)](#).

### Topics

- [Policy Best Practices \(p. 164\)](#)
- [Using the Amazon EKS Console \(p. 164\)](#)
- [Allow Users to View Their Own Permissions \(p. 164\)](#)
- [Update a Kubernetes cluster \(p. 165\)](#)
- [List or describe all clusters \(p. 165\)](#)



## Policy Best Practices

Identity-based policies are very powerful. They determine whether someone can create, access, or delete Amazon EKS resources in your account. These actions can incur costs for your AWS account. When you create or edit identity-based policies, follow these guidelines and recommendations:

- **Get Started Using AWS Managed Policies** – To start using Amazon EKS quickly, use AWS managed policies to give your employees the permissions they need. These policies are already available in your account and are maintained and updated by AWS. For more information, see [Get Started Using Permissions With AWS Managed Policies](#) in the *IAM User Guide*.
- **Grant Least Privilege** – When you create custom policies, grant only the permissions required to perform a task. Start with a minimum set of permissions and grant additional permissions as necessary. Doing so is more secure than starting with permissions that are too lenient and then trying to tighten them later. For more information, see [Grant Least Privilege](#) in the *IAM User Guide*.
- **Enable MFA for Sensitive Operations** – For extra security, require IAM users to use multi-factor authentication (MFA) to access sensitive resources or API operations. For more information, see [Using Multi-Factor Authentication \(MFA\) in AWS](#) in the *IAM User Guide*.
- **Use Policy Conditions for Extra Security** – To the extent that it's practical, define the conditions under which your identity-based policies allow access to a resource. For example, you can write conditions to specify a range of allowable IP addresses that a request must come from. You can also write conditions to allow requests only within a specified date or time range, or to require the use of SSL or MFA. For more information, see [IAM JSON Policy Elements: Condition](#) in the *IAM User Guide*.

## Using the Amazon EKS Console

To access the Amazon EKS console, you must have a minimum set of permissions. These permissions must allow you to list and view details about the Amazon EKS resources in your AWS account. If you create an identity-based policy that is more restrictive than the minimum required permissions, the console won't function as intended for entities (IAM users or roles) with that policy.

To ensure that those entities can still use the Amazon EKS console, create a policy with your own unique name, such as `AmazonEKSAAdminPolicy`. Attach the policy to the entities. For more information, see [Adding Permissions to a User](#) in the *IAM User Guide*:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "eks:*"
      ],
      "Resource": "*"
    }
  ]
}
```

You don't need to allow minimum console permissions for users that are making calls only to the AWS CLI or the AWS API. Instead, allow access to only the actions that match the API operation that you're trying to perform.

## Allow Users to View Their Own Permissions

This example shows how you might create a policy that allows IAM users to view the inline and managed policies that are attached to their user identity. This policy includes permissions to complete this action on the console or programmatically using the AWS CLI or AWS API.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "ViewOwnUserInfo",
      "Effect": "Allow",
      "Action": [
        "iam:GetUserPolicy",
        "iam:ListGroupsForUser",
        "iam:ListAttachedUserPolicies",
        "iam:ListUserPolicies",
        "iam:GetUser"
      ],
      "Resource": [
        "arn:aws:iam::*:user/${aws:username}"
      ]
    },
    {
      "Sid": "NavigateInConsole",
      "Effect": "Allow",
      "Action": [
        "iam:GetGroupPolicy",
        "iam:GetPolicyVersion",
        "iam:GetPolicy",
        "iam:ListAttachedGroupPolicies",
        "iam:ListGroupPolicies",
        "iam:ListPolicyVersions",
        "iam:ListPolicies",
        "iam:ListUsers"
      ],
      "Resource": "*"
    }
  ]
}
```

## Update a Kubernetes cluster

This example shows how you can create a policy that allows a user to update the Kubernetes version of any *dev* cluster for an account, in any region.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": "eks:UpdateClusterVersion",
      "Resource": "arn:aws:eks:*:11112223333:cluster/dev"
    }
  ]
}
```

## List or describe all clusters

This example shows how you can create a policy that allows a user read-only access to list or describe all clusters. An account must be able to list and describe clusters to use the `update-kubeconfig` AWS CLI command.

```
{
  "Version": "2012-10-17",
```

```
"Statement": [  
  {  
    "Effect": "Allow",  
    "Action": [  
      "eks:DescribeCluster",  
      "eks:ListClusters"  
    ],  
    "Resource": "*"   
  }  
]
```

## Amazon EKS Service IAM Role

Amazon EKS makes calls to other AWS services on your behalf to manage the resources that you use with the service. Before you can create Amazon EKS clusters, you must create an IAM role with the following IAM policies:

- [AmazonEKSServicePolicy](#)
- [AmazonEKSClusterPolicy](#)

## Check for an Existing Service Role

You can use the following procedure to check and see if your account already has the Amazon EKS service role.

### To check for the `eksServiceRole` in the IAM console

1. Open the IAM console at <https://console.aws.amazon.com/iam/>.
2. In the navigation pane, choose **Roles**.
3. Search the list of roles for `eksServiceRole` or `AWSServiceRoleForAmazonEKS`. If the role does not exist, see [Creating the Amazon EKS Service Role \(p. 167\)](#) to create the role. If the role does exist, select the role to view the attached policies.
4. Choose **Permissions**.
5. Ensure that the **AmazonEKSServicePolicy** and **AmazonEKSClusterPolicy** managed policies are attached to the role. If the policies are attached, your Amazon EKS service role is properly configured.
6. Choose **Trust Relationships, Edit Trust Relationship**.
7. Verify that the trust relationship contains the following policy. If the trust relationship matches the policy below, choose **Cancel**. If the trust relationship does not match, copy the policy into the **Policy Document** window and choose **Update Trust Policy**.

```
{  
  "Version": "2012-10-17",  
  "Statement": [  
    {  
      "Effect": "Allow",  
      "Principal": {  
        "Service": "eks.amazonaws.com"  
      },  
      "Action": "sts:AssumeRole"  
    }  
  ]  
}
```

## Creating the Amazon EKS Service Role

You can use the following procedure to create the Amazon EKS service role if you do not already have one for your account.

### To create your Amazon EKS service role in the IAM console

1. Open the IAM console at <https://console.aws.amazon.com/iam/>.
2. Choose **Roles**, then **Create role**.
3. Choose **EKS** from the list of services, then **Allows Amazon EKS to manage your clusters on your behalf** for your use case, then **Next: Permissions**.
4. Choose **Next: Tags**.
5. (Optional) Add metadata to the role by attaching tags as key–value pairs. For more information about using tags in IAM, see [Tagging IAM Entities](#) in the *IAM User Guide*.
6. Choose **Next: Review**.
7. For **Role name**, enter a unique name for your role, such as `eksServiceRole`, then choose **Create role**.

### To create your Amazon EKS service role with AWS CloudFormation

1. Save the following AWS CloudFormation template to a text file on your local system.

```
---
AWSTemplateFormatVersion: '2010-09-09'
Description: 'Amazon EKS Service Role'

Resources:

  eksServiceRole:
    Type: AWS::IAM::Role
    Properties:
      AssumeRolePolicyDocument:
        Version: '2012-10-17'
        Statement:
          - Effect: Allow
            Principal:
              Service:
                - eks.amazonaws.com
            Action:
              - sts:AssumeRole
      ManagedPolicyArns:
        - arn:aws:iam::aws:policy/AmazonEKSServicePolicy
        - arn:aws:iam::aws:policy/AmazonEKSClusterPolicy

Outputs:

  RoleArn:
    Description: The role that Amazon EKS will use to create AWS resources for
    Kubernetes clusters
    Value: !GetAtt eksServiceRole.Arn
    Export:
      Name: !Sub "${AWS::StackName}-RoleArn"
```

2. Open the AWS CloudFormation console at <https://console.aws.amazon.com/cloudformation>.
3. Choose **Create stack**.
4. For **Specify template**, select **Upload a template file**, and then choose **Choose file**.
5. Choose the file you created earlier, and then choose **Next**.

6. For **Stack name**, enter a name for your role, such as `eksServiceRole`, and then choose **Next**.
7. On the **Configure stack options** page, choose **Next**.
8. On the **Review** page, review your information, acknowledge that the stack might create IAM resources, and then choose **Create stack**.

## Amazon EKS Worker Node IAM Role

The Amazon EKS worker node `kubelet` daemon makes calls to AWS APIs on your behalf. Worker nodes receive permissions for these API calls through an IAM instance profile and associated policies. Before you can launch worker nodes and register them into a cluster, you must create an IAM role for those worker nodes to use when they are launched. This requirement applies to worker nodes launched with the Amazon EKS-optimized AMI provided by Amazon, or with any other worker node AMIs that you intend to use. Before you create worker nodes, you must create an IAM role with the following IAM policies:

- [AmazonEKSWorkerNodePolicy](#)
- [AmazonEKS\\_CNI\\_Policy](#)
- [AmazonEC2ContainerRegistryReadOnly](#)

## Check for an Existing Worker Node Role

You can use the following procedure to check and see if your account already has the Amazon EKS worker node role.

### To check for the `NodeInstanceRole` in the IAM console

1. Open the IAM console at <https://console.aws.amazon.com/iam/>.
2. In the navigation pane, choose **Roles**.
3. Search the list of roles for `NodeInstanceRole`. If the role does not exist, see [Creating the Amazon EKS Worker Node Role](#) (p. 169) to create the role. If the role does exist, select the role to view the attached policies.
4. Choose **Permissions**.
5. Ensure that the **AmazonEKSWorkerNodePolicy**, **AmazonEKS\_CNI\_Policy**, and **AmazonEC2ContainerRegistryReadOnly** managed policies are attached to the role. If the policies are attached, your Amazon EKS worker node role is properly configured.
6. Choose **Trust Relationships**, **Edit Trust Relationship**.
7. Verify that the trust relationship contains the following policy. If the trust relationship matches the policy below, choose **Cancel**. If the trust relationship does not match, copy the policy into the **Policy Document** window and choose **Update Trust Policy**.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Principal": {
        "Service": "ec2.amazonaws.com"
      },
      "Action": "sts:AssumeRole"
    }
  ]
}
```

## Creating the Amazon EKS Worker Node Role

If you created your worker nodes by following the steps in the [Getting Started with the AWS Management Console](#) (p. 8) or [Getting Started with eksctl](#) (p. 3) topics, then the worker node role account already exists and you don't need to manually create it. You can use the following procedure to create the Amazon EKS worker node role if you do not already have one for your account.

### To create your Amazon EKS worker node role in the IAM console

1. Open the IAM console at <https://console.aws.amazon.com/iam/>.
2. Choose **Roles**, then **Create role**.
3. Choose **EC2** from the list of services, then **Next: Permissions**.
4. Select the following policies:
  - **AmazonEKSWorkerNodePolicy**
  - **AmazonEKS\_CNI\_Policy**
  - **AmazonEC2ContainerRegistryReadOnly**
5. Choose **Next: Tags**.
6. (Optional) Add metadata to the role by attaching tags as key–value pairs. For more information about using tags in IAM, see [Tagging IAM Entities](#) in the *IAM User Guide*.
7. Choose **Next: Review**.
8. For **Role name**, enter a unique name for your role, such as `NodeInstanceRole`, then choose **Create role**.

### To create your Amazon EKS instance role with AWS CloudFormation

1. Save the following AWS CloudFormation template to a text file on your local system.

```
---
AWSTemplateFormatVersion: '2010-09-09'
Description: 'Amazon EKS Worker Node Role'

Resources:

  NodeInstanceRole:
    Type: AWS::IAM::Role
    Properties:
      AssumeRolePolicyDocument:
        Version: 2012-10-17
        Statement:
          - Effect: Allow
            Principal:
              Service: ec2.amazonaws.com
            Action: sts:AssumeRole
      Path: "/"
      ManagedPolicyArns:
        - arn:aws:iam::aws:policy/AmazonEKSWorkerNodePolicy
        - arn:aws:iam::aws:policy/AmazonEKS_CNI_Policy
        - arn:aws:iam::aws:policy/AmazonEC2ContainerRegistryReadOnly

Outputs:

  RoleArn:
    Description: The role that the worker node kubelet uses to make calls to the Amazon
    EKS API on your behalf
    Value: !GetAtt NodeInstanceRole.Arn
    Export:
```

```
Name: !Sub "${AWS::StackName}-RoleArn"
```

2. Open the AWS CloudFormation console at <https://console.aws.amazon.com/cloudformation>.
3. Choose **Create stack**.
4. For **Specify template**, select **Upload a template file**, and then choose **Choose file**.
5. Choose the file you created earlier, and then choose **Next**.
6. For **Stack name**, enter a name for your role, such as `NodeInstanceRole`, and then choose **Next**.
7. On the **Configure stack options** page, choose **Next**.
8. On the **Review** page, review your information, acknowledge that the stack might create IAM resources, and then choose **Create stack**.

## Troubleshooting Amazon EKS Identity and Access

To diagnose and fix common issues that you might encounter when working with Amazon EKS and IAM see [Troubleshooting IAM \(p. 181\)](#).

## Logging and Monitoring in Amazon EKS

Amazon EKS control plane logging provides audit and diagnostic logs directly from the Amazon EKS control plane to CloudWatch Logs in your account. These logs make it easy for you to secure and run your clusters. You can select the exact log types you need, and logs are sent as log streams to a group for each Amazon EKS cluster in CloudWatch. For more information, see [Amazon EKS Control Plane Logging \(p. 38\)](#).

Amazon EKS is integrated with AWS CloudTrail, a service that provides a record of actions taken by a user, role, or an AWS service in Amazon EKS. CloudTrail captures all API calls for Amazon EKS as events. The calls captured include calls from the Amazon EKS console and code calls to the Amazon EKS API operations. For more information, see [Logging Amazon EKS API Calls with AWS CloudTrail \(p. 173\)](#).

## Compliance Validation for Amazon EKS

Third-party auditors assess the security and compliance of Amazon EKS as part of multiple AWS compliance programs. These include SOC, PCI, ISO, HIPAA, and others.

For a list of AWS services in scope of specific compliance programs, see [AWS Services in Scope by Compliance Program](#). For general information, see [AWS Compliance Programs](#).

You can download third-party audit reports using AWS Artifact. For more information, see [Downloading Reports in AWS Artifact](#).

Your compliance responsibility when using Amazon EKS is determined by the sensitivity of your data, your company's compliance objectives, and applicable laws and regulations. AWS provides the following resources to help with compliance:

- [Security and Compliance Quick Start Guides](#) – These deployment guides discuss architectural considerations and provide steps for deploying security- and compliance-focused baseline environments on AWS.
- [Architecting for HIPAA Security and Compliance Whitepaper](#) – This whitepaper describes how companies can use AWS to create HIPAA-compliant applications.
- [AWS Compliance Resources](#) – This collection of workbooks and guides might apply to your industry and location.

- [AWS Config](#) – This AWS service assesses how well your resource configurations comply with internal practices, industry guidelines, and regulations.
- [AWS Security Hub](#) – This AWS service provides a comprehensive view of your security state within AWS that helps you check your compliance with security industry standards and best practices.

## Resilience in Amazon EKS

The AWS global infrastructure is built around AWS Regions and Availability Zones. AWS Regions provide multiple physically separated and isolated Availability Zones, which are connected with low-latency, high-throughput, and highly redundant networking. With Availability Zones, you can design and operate applications and databases that automatically fail over between Availability Zones without interruption. Availability Zones are more highly available, fault tolerant, and scalable than traditional single or multiple data center infrastructures.

Amazon EKS runs Kubernetes control plane instances across multiple Availability Zones to ensure high availability. Amazon EKS automatically detects and replaces unhealthy control plane instances, and it provides automated version upgrades and patching for them.

This control plane consists of at least two API server nodes and three `etcd` nodes that run across three Availability Zones within a Region. Amazon EKS automatically detects and replaces unhealthy control plane instances, restarting them across the Region as needed. Amazon EKS leverages the architecture of AWS Regions in order to maintain high availability. Because of this, Amazon EKS is able to offer an [SLA for API server endpoint availability](#).

For more information about AWS Regions and Availability Zones, see [AWS Global Infrastructure](#).

## Infrastructure Security in Amazon EKS

As a managed service, Amazon EKS is protected by the AWS global network security procedures that are described in the [Amazon Web Services: Overview of Security Processes](#) whitepaper.

You use AWS published API calls to access Amazon EKS through the network. Clients must support Transport Layer Security (TLS) 1.0 or later. We recommend TLS 1.2 or later. Clients must also support cipher suites with perfect forward secrecy (PFS) such as Ephemeral Diffie-Hellman (DHE) or Elliptic Curve Ephemeral Diffie-Hellman (ECDHE). Most modern systems such as Java 7 and later support these modes.

Additionally, requests must be signed by using an access key ID and a secret access key that is associated with an IAM principal. Or you can use the [AWS Security Token Service](#) (AWS STS) to generate temporary security credentials to sign requests.

When you create an Amazon EKS cluster, you specify the Amazon VPC subnets for your cluster to use. Amazon EKS requires subnets in at least two Availability Zones. We recommend a network architecture that uses private subnets for your worker nodes and public subnets for Kubernetes to create internet-facing load balancers within.

For more information about VPC considerations, see [Cluster VPC Considerations \(p. 81\)](#).

If you create your VPC and worker node groups with the AWS CloudFormation templates provided in the [Getting Started with Amazon EKS \(p. 3\)](#) walkthrough, then your control plane and worker node security groups are configured with our recommended settings.

For more information about security group considerations, see [Cluster Security Group Considerations \(p. 83\)](#).

When you create a new cluster, Amazon EKS creates an endpoint for the managed Kubernetes API server that you use to communicate with your cluster (using Kubernetes management tools such as `kubectl`).



By default, this API server endpoint is public to the internet, and access to the API server is secured using a combination of AWS Identity and Access Management (IAM) and native Kubernetes [Role Based Access Control](#) (RBAC).

You can enable private access to the Kubernetes API server so that all communication between your worker nodes and the API server stays within your VPC. You can also completely disable public access to your API server so that it's not accessible from the internet.

For more information about modifying cluster endpoint access, see [Modifying Cluster Endpoint Access](#) (p. 35).

You can implement network policies with tools such as [Project Calico](#) (p. 97). Project Calico is a third party open source project. For more information, see the [Project Calico documentation](#).

## Configuration and Vulnerability Analysis in Amazon EKS

Amazon EKS platform versions represent the capabilities of the cluster control plane, including which Kubernetes API server flags are enabled and the current Kubernetes patch version. New clusters are deployed with the latest platform version. For details, see [Platform Versions](#) (p. 44).

You can [update an Amazon EKS cluster](#) (p. 26) to newer Kubernetes versions. As new Kubernetes versions become available in Amazon EKS, we recommend that you proactively update your clusters to use the latest available version. For more information about Kubernetes versions in EKS, see [Amazon EKS Kubernetes Versions](#) (p. 43).

Track security or privacy events for Amazon Linux 2 at the [Amazon Linux Security Center](#) or subscribe to the associated [RSS feed](#). Security and privacy events include an overview of the issue affected, packages, and instructions for updating your instances to correct the issue.

You can use [Amazon Inspector](#) to check for unintended network accessibility of your worker nodes and for vulnerabilities on those Amazon EC2 instances.

# Logging Amazon EKS API Calls with AWS CloudTrail

Amazon EKS is integrated with AWS CloudTrail, a service that provides a record of actions taken by a user, role, or an AWS service in Amazon EKS. CloudTrail captures all API calls for Amazon EKS as events. The calls captured include calls from the Amazon EKS console and code calls to the Amazon EKS API operations.

If you create a trail, you can enable continuous delivery of CloudTrail events to an Amazon S3 bucket, including events for Amazon EKS. If you don't configure a trail, you can still view the most recent events in the CloudTrail console in **Event history**. Using the information collected by CloudTrail, you can determine the request that was made to Amazon EKS, the IP address from which the request was made, who made the request, when it was made, and additional details.

To learn more about CloudTrail, see the [AWS CloudTrail User Guide](#).

## Amazon EKS Information in CloudTrail

CloudTrail is enabled on your AWS account when you create the account. When activity occurs in Amazon EKS, that activity is recorded in a CloudTrail event along with other AWS service events in **Event history**. You can view, search, and download recent events in your AWS account. For more information, see [Viewing Events with CloudTrail Event History](#).

For an ongoing record of events in your AWS account, including events for Amazon EKS, create a trail. A *trail* enables CloudTrail to deliver log files to an Amazon S3 bucket. By default, when you create a trail in the console, the trail applies to all AWS Regions. The trail logs events from all Regions in the AWS partition and delivers the log files to the Amazon S3 bucket that you specify. Additionally, you can configure other AWS services to further analyze and act upon the event data collected in CloudTrail logs. For more information, see the following:

- [Overview for Creating a Trail](#)
- [CloudTrail Supported Services and Integrations](#)
- [Configuring Amazon SNS Notifications for CloudTrail](#)
- [Receiving CloudTrail Log Files from Multiple Regions](#) and [Receiving CloudTrail Log Files from Multiple Accounts](#)

All Amazon EKS actions are logged by CloudTrail and are documented in the [Amazon EKS API Reference](#). For example, calls to the [CreateCluster](#), [ListClusters](#) and [DeleteCluster](#) sections generate entries in the CloudTrail log files.

Every event or log entry contains information about who generated the request. The identity information helps you determine the following:

- Whether the request was made with root or AWS Identity and Access Management (IAM) user credentials.
- Whether the request was made with temporary security credentials for a role or federated user.
- Whether the request was made by another AWS service.

For more information, see the [CloudTrail userIdentity Element](#).

# Understanding Amazon EKS Log File Entries

A trail is a configuration that enables delivery of events as log files to an Amazon S3 bucket that you specify. CloudTrail log files contain one or more log entries. An event represents a single request from any source and includes information about the requested action, the date and time of the action, request parameters, and so on. CloudTrail log files aren't an ordered stack trace of the public API calls, so they don't appear in any specific order.

The following example shows a CloudTrail log entry that demonstrates the `CreateCluster` action.

```
{
  "eventVersion": "1.05",
  "userIdentity": {
    "type": "IAMUser",
    "principalId": "AKIAIOSFODNN7EXAMPLE",
    "arn": "arn:aws:iam::111122223333:user/ericn",
    "accountId": "111122223333",
    "accessKeyId": "AKIAIOSFODNN7EXAMPLE",
    "userName": "ericn"
  },
  "eventTime": "2018-05-28T19:16:43Z",
  "eventSource": "eks.amazonaws.com",
  "eventName": "CreateCluster",
  "awsRegion": "us-west-2",
  "sourceIPAddress": "205.251.233.178",
  "userAgent": "PostmanRuntime/6.4.0",
  "requestParameters": {
    "resourcesVpcConfig": {
      "subnetIds": [
        "subnet-a670c2df",
        "subnet-4f8c5004"
      ]
    }
  },
  "roleArn": "arn:aws:iam::111122223333:role/AWSServiceRoleForAmazonEKS-CAC1G1VH3ZKZ",
  "clusterName": "test"
},
"responseElements": {
  "cluster": {
    "clusterName": "test",
    "status": "CREATING",
    "createdAt": 1527535003.208,
    "certificateAuthority": {},
    "arn": "arn:aws:eks:us-west-2:111122223333:cluster/test",
    "roleArn": "arn:aws:iam::111122223333:role/AWSServiceRoleForAmazonEKS-CAC1G1VH3ZKZ",
    "version": "1.10",
    "resourcesVpcConfig": {
      "securityGroupIds": [],
      "vpcId": "vpc-21277358",
      "subnetIds": [
        "subnet-a670c2df",
        "subnet-4f8c5004"
      ]
    }
  }
}
},
"requestID": "a7a0735d-62ab-11e8-9f79-81ce5b2b7d37",
"eventID": "eab22523-174a-499c-9dd6-91e7be3ff8e3",
"readOnly": false,
"eventType": "AwsApiCall",
"recipientAccountId": "111122223333"
}
```

# Related Projects

These open source projects extend the functionality of Kubernetes clusters running on AWS, including clusters managed by Amazon EKS.

## Management Tools

Related management tools for Amazon EKS and Kubernetes clusters.

### eksctl

`eksctl` is a simple CLI tool for creating clusters on Amazon EKS.

- Project URL: <https://eksctl.io/>
- Project documentation: <https://eksctl.io/>
- AWS open source blog: [eksctl: Amazon EKS Cluster with One Command](#)

### AWS Service Operator

AWS Service Operator allows you to create AWS resources using `kubectl`.

- Project URL: <https://github.com/awslabs/aws-service-operator>
- Project documentation: <https://github.com/awslabs/aws-service-operator/blob/master/readme.adoc>
- AWS open source blog: [AWS Service Operator for Kubernetes Now Available](#)

## Networking

Related networking projects for Amazon EKS and Kubernetes clusters.

### Amazon VPC CNI plugin for Kubernetes

Amazon EKS supports native VPC networking via the Amazon VPC CNI plugin for Kubernetes. Using this CNI plugin allows Kubernetes pods to have the same IP address inside the pod as they do on the VPC network. For more information, see [Pod Networking \(p. 85\)](#) and [CNI Configuration Variables \(p. 87\)](#).

- Project URL: <https://github.com/aws/amazon-vpc-cni-k8s>
- Project documentation: <https://github.com/aws/amazon-vpc-cni-k8s/blob/master/README.md>

### AWS Application Load Balancer (ALB) Ingress Controller for Kubernetes

The AWS ALB Ingress Controller satisfies Kubernetes ingress resources by provisioning Application Load Balancers.

- Project URL: <https://github.com/kubernetes-sigs/aws-alb-ingress-controller>
- Project documentation: <https://github.com/kubernetes-sigs/aws-alb-ingress-controller/tree/master/docs>
- AWS open source blog: [Kubernetes Ingress with AWS ALB Ingress Controller](#)

## ExternalDNS

ExternalDNS synchronizes exposed Kubernetes services and ingresses with DNS providers including Amazon Route 53 and AWS Service Discovery.

- Project URL: <https://github.com/kubernetes-incubator/external-dns>
- Project documentation: <https://github.com/kubernetes-incubator/external-dns/blob/master/docs/tutorials/aws.md>

## Security

Related security projects for Amazon EKS and Kubernetes clusters.

### AWS IAM Authenticator

A tool to use AWS IAM credentials to authenticate to a Kubernetes cluster. For more information, see [Installing aws-iam-authenticator](#) (p. 108).

- Project URL: <https://github.com/kubernetes-sigs/aws-iam-authenticator>
- Project documentation: <https://github.com/kubernetes-sigs/aws-iam-authenticator/blob/master/README.md>
- AWS open source blog: [Deploying the AWS IAM Authenticator to kops](#)

## Storage

Related storage projects for Amazon EKS and Kubernetes clusters.

### Amazon EFS CSI Driver

The Amazon Elastic File System Container Storage Interface (CSI) Driver implements the CSI specification for container orchestrators to manage the lifecycle of Amazon EFS resources.

- Project URL: <https://github.com/aws/csi-driver-aws-efs>
- Project documentation: <https://github.com/aws/aws-efs-csi-driver/blob/master/docs/README.md>

## Machine Learning

Related machine learning projects for Amazon EKS and Kubernetes clusters.

### Kubeflow

A machine learning toolkit for Kubernetes.

- Project URL: <https://www.kubeflow.org/>
- Project documentation: <https://www.kubeflow.org/docs/>
- AWS open source blog: [Kubeflow on Amazon EKS](#)

## Auto Scaling

Related auto scaling projects for Amazon EKS and Kubernetes clusters.

### Cluster Autoscaler

Cluster Autoscaler is a tool that automatically adjusts the size of the Kubernetes cluster based on CPU and memory pressure.

- Project URL: <https://github.com/kubernetes/autoscaler/tree/master/cluster-autoscaler>
- Project documentation: <https://github.com/kubernetes/autoscaler/blob/master/cluster-autoscaler/cloudprovider/aws/README.md>
- Amazon EKS workshop: [https://eksworkshop.com/scaling/deploy\\_ca/](https://eksworkshop.com/scaling/deploy_ca/)

### Escalator

Escalator is a batch or job optimized horizontal autoscaler for Kubernetes.

- Project URL: <https://github.com/atlassian/escalator>
- Project documentation: <https://github.com/atlassian/escalator/blob/master/docs/README.md>

## Monitoring

Related monitoring projects for Amazon EKS and Kubernetes clusters.

### Prometheus

Prometheus is an open-source systems monitoring and alerting toolkit.

- Project URL: <https://prometheus.io/>
- Project documentation: <https://prometheus.io/docs/introduction/overview/>
- Amazon EKS workshop: <https://eksworkshop.com/monitoring/>

## Continuous Integration / Continuous Deployment

Related CI/CD projects for Amazon EKS and Kubernetes clusters.

### Jenkins X

CI/CD solution for modern cloud applications on Amazon EKS and Kubernetes clusters.

- Project URL: <https://jenkins-x.io/>
- Project documentation: <https://jenkins-x.io/documentation/>

- AWS open source blog: [Continuous Delivery with Amazon EKS and Jenkins X](#)

# Amazon EKS Troubleshooting

This chapter covers some common errors that you may see while using Amazon EKS and how to work around them.

## Insufficient Capacity

If you receive the following error while attempting to create an Amazon EKS cluster, then one of the Availability Zones you specified does not have sufficient capacity to support a cluster.

Cannot create cluster '*example-cluster*' because *us-east-1d*, the targeted availability zone, does not currently have sufficient capacity to support the cluster. Retry and choose from these availability zones: *us-east-1a*, *us-east-1b*, *us-east-1c*

Retry creating your cluster with subnets in your cluster VPC that are hosted in the Availability Zones returned by this error message.

## aws-iam-authenticator Not Found

If you receive the error "aws-iam-authenticator": executable file not found in \$PATH, then your **kubectl** is not configured for Amazon EKS. For more information, see [Installing aws-iam-authenticator](#) (p. 108).

## Worker Nodes Fail to Join Cluster

There are two common reasons that prevent worker nodes from joining the cluster:

- The `aws-auth-cm.yaml` file does not have the correct IAM role ARN for your worker nodes. Ensure that the worker node IAM role ARN (not the instance profile ARN) is specified in your `aws-auth-cm.yaml` file. For more information, see [Launching Amazon EKS Worker Nodes](#) (p. 56).
- The **ClusterName** in your worker node AWS CloudFormation template does not exactly match the name of the cluster you want your worker nodes to join. Passing an incorrect value to this field results in an incorrect configuration of the worker node's `/var/lib/kubelet/kubeconfig` file, and the nodes will not join the cluster.

## Unauthorized or Access Denied (kubectl)

If you receive one of the following errors while running **kubectl** commands, then your **kubectl** is not configured properly for Amazon EKS or the IAM user or role credentials that you are using do not map to a Kubernetes RBAC user with sufficient permissions in your Amazon EKS cluster.

- `could not get token: AccessDenied: Access denied`
- `error: You must be logged in to the server (Unauthorized)`
- `error: the server doesn't have a resource type "svc"`



This could be because the cluster was created with one set of AWS credentials (from an IAM user or role), and **kubectl** is using a different set of credentials.

When an Amazon EKS cluster is created, the IAM entity (user or role) that creates the cluster is added to the Kubernetes RBAC authorization table as the administrator (with `system:master` permissions). Initially, only that IAM user can make calls to the Kubernetes API server using **kubectl**. For more information, see [Managing Users or IAM Roles for your Cluster \(p. 115\)](#). Also, the [AWS IAM Authenticator for Kubernetes](#) uses the AWS SDK for Go to authenticate against your Amazon EKS cluster. If you use the console to create the cluster, you must ensure that the same IAM user credentials are in the AWS SDK credential chain when you are running **kubectl** commands on your cluster.

If you install and configure the AWS CLI, you can configure the IAM credentials for your user. If the AWS CLI is configured properly for your user, then the [AWS IAM Authenticator for Kubernetes](#) can find those credentials as well. For more information, see [Configuring the AWS CLI](#) in the *AWS Command Line Interface User Guide*.

If you assumed a role to create the Amazon EKS cluster, you must ensure that **kubectl** is configured to assume the same role. Use the following command to update your kubeconfig file to use an IAM role. For more information, see [Create a kubeconfig for Amazon EKS \(p. 111\)](#).

```
aws --region region eks update-kubeconfig --name cluster_name --role-arn  
arn:aws:iam::aws_account_id:role/role_name
```

To map an IAM user to a Kubernetes RBAC user, see [Managing Users or IAM Roles for your Cluster \(p. 115\)](#).

## hostname doesn't match

Your system's Python version must be 2.7.9 or greater. Otherwise, you receive `hostname doesn't match` errors with AWS CLI calls to Amazon EKS. For more information, see [What are "hostname doesn't match" errors?](#) in the Python Requests FAQ.

## getsockopt: no route to host

Docker runs in the 172.17.0.0/16 CIDR range in Amazon EKS clusters. We recommend that your cluster's VPC subnets do not overlap this range. Otherwise, you will receive the following error:

```
Error: : error upgrading connection: error dialing backend: dial tcp 172.17.nn.nn:10250:  
getsockopt: no route to host
```

## CNI Log Collection Tool

The Amazon VPC CNI plugin for Kubernetes has its own troubleshooting script (which is available on worker nodes at `/opt/cni/bin/aws-cni-support.sh`) that you can use to collect diagnostic logs for support cases and general troubleshooting.

The script collects the following diagnostic information:

- L-IPAMD introspection data
- Metrics
- Kubelet introspection data

- `ifconfig` output
- `ip rule show` output
- `iptables-save` output
- `iptables -nvL` output
- `iptables -nvL -t nat` output
- A dump of the CNI configuration
- Kubelet logs
- Stored `/var/log/messages`
- Worker node's route table information (via `ip route`)
- The `sysctl`s output of `/proc/sys/net/ipv4/conf/{all,default,eth0}/rp_filter`

Use the following command to run the script on your worker node:

```
sudo bash /opt/cni/bin/aws-cni-support.sh
```

**Note**

If the script is not present at that location, then the CNI container failed to run. You can manually download and run the script with the following command:

```
curl https://raw.githubusercontent.com/aws/amazon-vpc-cni-k8s/master/scripts/aws-cni-support.sh | sudo bash
```

The diagnostic information is collected and stored at `/var/log/aws-routed-eni/aws-cni-support.tar.gz`.

## Troubleshooting IAM

This topic covers some common errors that you may see while using Amazon EKS with IAM and how to work around them.

### AccessDeniedException

If you receive an `AccessDeniedException` when calling an AWS API operation, then the AWS Identity and Access Management (IAM) user or role credentials that you are using do not have the required permissions to make that call.

```
An error occurred (AccessDeniedException) when calling the DescribeCluster operation:
User: arn:aws:iam::111122223333:user/user_name is not authorized to perform:
eks:DescribeCluster on resource: arn:aws:eks:us-west-2:111122223333:cluster/cluster_name
```

In the above example message, the user does not have permissions to call the Amazon EKS `DescribeCluster` API operation. To provide Amazon EKS admin permissions to a user, see [Amazon EKS Identity-Based Policy Examples \(p. 163\)](#).

For more general information about IAM, see [Controlling Access Using Policies](#) in the *IAM User Guide*.

### I Am Not Authorized to Perform iam:PassRole

If you receive an error that you're not authorized to perform the `iam:PassRole` action, then you must contact your administrator for assistance. Your administrator is the person that provided you with your

user name and password. Ask that person to update your policies to allow you to pass a role to Amazon EKS.

Some AWS services allow you to pass an existing role to that service, instead of creating a new service role or service-linked role. To do this, you must have permissions to pass the role to the service.

The following example error occurs when an IAM user named `marymajor` tries to use the console to perform an action in Amazon EKS. However, the action requires the service to have permissions granted by a service role. Mary does not have permissions to pass the role to the service.

```
User: arn:aws:iam::123456789012:user/marymajor is not authorized to perform: iam:PassRole
```

In this case, Mary asks her administrator to update her policies to allow her to perform the `iam:PassRole` action.

## I Want to View My Access Keys

After you create your IAM user access keys, you can view your access key ID at any time. However, you can't view your secret access key again. If you lose your secret key, you must create a new access key pair.

Access keys consist of two parts: an access key ID (for example, `AKIAIOSFODNN7EXAMPLE`) and a secret access key (for example, `wJalrXUtnFEMI/K7MDENG/bPxrFcYEXAMPLEKEY`). Like a user name and password, you must use both the access key ID and secret access key together to authenticate your requests. Manage your access keys as securely as you do your user name and password.

### Important

Do not provide your access keys to a third party, even to help [find your canonical user ID](#). By doing this, you might give someone permanent access to your account.

When you create an access key pair, you are prompted to save the access key ID and secret access key in a secure location. The secret access key is available only at the time you create it. If you lose your secret access key, you must add new access keys to your IAM user. You can have a maximum of two access keys. If you already have two, you must delete one key pair before creating a new one. To view instructions, see [Managing Access Keys](#) in the *IAM User Guide*.

## I'm an Administrator and Want to Allow Others to Access Amazon EKS

To allow others to access Amazon EKS, you must create an IAM entity (user or role) for the person or application that needs access. They will use the credentials for that entity to access AWS. You must then attach a policy to the entity that grants them the correct permissions in Amazon EKS.

To get started right away, see [Creating Your First IAM Delegated User and Group](#) in the *IAM User Guide*.

## I Want to Allow People Outside of My AWS Account to Access My Amazon EKS Resources

You can create a role that users in other accounts or people outside of your organization can use to access your resources. You can specify who is trusted to assume the role. For services that support resource-based policies or access control lists (ACLs), you can use those policies to grant people access to your resources.

To learn more, consult the following:

- To learn whether Amazon EKS supports these features, see [How Amazon EKS Works with IAM](#) (p. 161).

- To learn how to provide access to your resources across AWS accounts that you own, see [Providing Access to an IAM User in Another AWS Account That You Own](#) in the *IAM User Guide*.
- To learn how to provide access to your resources to third-party AWS accounts, see [Providing Access to AWS Accounts Owned by Third Parties](#) in the *IAM User Guide*.
- To learn how to provide access through identity federation, see [Providing Access to Externally Authenticated Users \(Identity Federation\)](#) in the *IAM User Guide*.
- To learn the difference between using roles and resource-based policies for cross-account access, see [How IAM Roles Differ from Resource-based Policies](#) in the *IAM User Guide*.

# Amazon EKS Service Limits

The following table provides the default limits for Amazon EKS for an AWS account that can be changed. For more information, see [AWS Service Limits](#) in the *Amazon Web Services General Reference*.

Resource	Default Limit
Maximum number of Amazon EKS clusters per region, per account	50

The following table provides limitations for Amazon EKS that cannot be changed.

Resource	Default Limit
Maximum number of control plane security groups per cluster (these are specified when you create the cluster)	5

# Document History for Amazon EKS

The following table describes the major updates and new features for the Amazon EKS User Guide. We also update the documentation frequently to address the feedback that you send us.

update-history-change	update-history-description	update-history-date
<a href="#">Amazon EKS platform version update</a>	New platform versions to address <a href="#">CVE-2019-11247</a> .	August 5, 2019
<a href="#">Amazon EKS region expansion (p. 185)</a>	Amazon EKS is now available in the Asia Pacific (Hong Kong) ( <code>ap-east-1</code> ) region.	July 31, 2019
<a href="#">Kubernetes 1.10 deprecated on Amazon EKS</a>	Kubernetes version 1.10 is no longer supported on Amazon EKS. Please update any 1.10 clusters to version 1.11 or higher in order to avoid service interruption.	July 30, 2019
<a href="#">Added topic on ALB Ingress Controller</a>	The AWS ALB Ingress Controller for Kubernetes is a controller that triggers the creation of an Application Load Balancer when Ingress resources are created.	July 11, 2019
<a href="#">New Amazon EKS-optimized AMI</a>	Removing unnecessary <code>kubect1</code> binary from AMIs.	July 3, 2019
<a href="#">Kubernetes Version 1.13</a>	Added Kubernetes version 1.13 support for new clusters and version upgrades.	June 18, 2019
<a href="#">New Amazon EKS-optimized AMI patched for AWS-2019-005</a>	Amazon EKS has updated the Amazon EKS-optimized AMI to address the vulnerabilities described in <a href="#">AWS-2019-005</a> .	June 17, 2019
<a href="#">Announcing deprecation of Kubernetes 1.10 in Amazon EKS</a>	Amazon EKS will deprecate Kubernetes version 1.10 on July 22, 2019. On this day, you will no longer be able to create new 1.10 clusters and all Amazon EKS clusters running Kubernetes version 1.10 will be updated to the latest available platform version of Kubernetes version 1.11.	May 21, 2019
<a href="#">Amazon EKS platform version update</a>	New platform version for Kubernetes 1.11 and 1.10 clusters to support custom DNS names in the Kubelet certificate and improve <code>etcd</code> performance.	May 21, 2019

<a href="#">Getting Started with eksctl</a>	This getting started guide helps you to install all of the required resources to get started with Amazon EKS using <code>eksctl</code> , a simple command line utility for creating and managing Kubernetes clusters on Amazon EKS.	May 10, 2019
<a href="#">AWS CLI get-token command (p. 185)</a>	The <b>aws eks get-token</b> command was added to the AWS CLI so that you no longer need to install the AWS IAM Authenticator for Kubernetes to create client security tokens for cluster API server communication. Upgrade your AWS CLI installation to the latest version to take advantage of this new functionality. For more information, see <a href="#">Installing the AWS Command Line Interface</a> in the <i>AWS Command Line Interface User Guide</i> .	May 10, 2019
<a href="#">Amazon EKS platform version update</a>	New platform version for Kubernetes 1.12 clusters to support custom DNS names in the Kubelet certificate and improve <code>etcd</code> performance. This fixes a bug that caused worker node Kubelet daemons to request a new certificate every few seconds.	May 8, 2019
<a href="#">Prometheus tutorial</a>	Added topic for deploying Prometheus to your Amazon EKS cluster.	April 5, 2019
<a href="#">Amazon EKS Control Plane Logging</a>	Amazon EKS control plane logging makes it easy for you to secure and run your clusters by providing audit and diagnostic logs directly from the Amazon EKS control plane to CloudWatch Logs in your account.	April 4, 2019
<a href="#">Kubernetes Version 1.12 (p. 185)</a>	Added Kubernetes version 1.12 support for new clusters and version upgrades.	March 28, 2019
<a href="#">Added App Mesh Getting Started Guide</a>	Added documentation for getting started with App Mesh and Kubernetes.	March 27, 2019

<a href="#">Amazon EKS API server endpoint private access</a>	Added documentation for disabling public access for your Amazon EKS cluster's Kubernetes API server endpoint.	March 19, 2019
<a href="#">Added topic for installing the Kubernetes metrics server</a>	The Kubernetes metrics server is an aggregator of resource usage data in your cluster.	March 18, 2019
<a href="#">Added list of related open source projects</a>	These open source projects extend the functionality of Kubernetes clusters running on AWS, including clusters managed by Amazon EKS.	March 15, 2019
<a href="#">Added topic for installing Helm locally</a>	The <code>helm</code> package manager for Kubernetes helps you install and manage applications on your Kubernetes cluster. This topic helps you install and run the <code>helm</code> and <code>tiller</code> binaries locally so that you can install and manage charts using the <code>helm</code> CLI on your local system.	March 11, 2019
<a href="#">Amazon EKS platform version update</a>	New platform version updating Amazon EKS Kubernetes 1.11 clusters to patch level 1.11.8 to address <a href="#">CVE-2019-1002100</a> .	March 8, 2019
<a href="#">Increased cluster limit</a>	Amazon EKS has increased the number of clusters that you can create in a region from 3 to 50.	February 13, 2019
<a href="#">Amazon EKS region expansion (p. 185)</a>	Amazon EKS is now available in the EU (London) ( <code>eu-west-2</code> ), EU (Paris) ( <code>eu-west-3</code> ), and Asia Pacific (Mumbai) ( <code>ap-south-1</code> ) regions.	February 13, 2019
<a href="#">New Amazon EKS-optimized AMI patched for ALAS-2019-1156</a>	Amazon EKS has updated the Amazon EKS-optimized AMI to address the vulnerability described in <a href="#">ALAS-2019-1156</a> .	February 11, 2019
<a href="#">New Amazon EKS-optimized AMI patched for ALAS2-2019-1141</a>	Amazon EKS has updated the Amazon EKS-optimized AMI to address the CVEs referenced in <a href="#">ALAS2-2019-1141</a> .	January 9, 2019
<a href="#">Amazon EKS region expansion (p. 185)</a>	Amazon EKS is now available in the Asia Pacific (Seoul) ( <code>ap-northeast-2</code> ) region.	January 9, 2019



<a href="#">Amazon EKS region expansion (p. 185)</a>	Amazon EKS is now available in the following additional regions: EU (Frankfurt) (eu-central-1), Asia Pacific (Tokyo) (ap-northeast-1), Asia Pacific (Singapore) (ap-southeast-1), and Asia Pacific (Sydney) (ap-southeast-2).	December 19, 2018
<a href="#">Amazon EKS cluster updates</a>	Added documentation for Amazon EKS <a href="#">cluster Kubernetes version updates</a> and <a href="#">worker node replacement</a> .	December 12, 2018
<a href="#">Amazon EKS region expansion (p. 185)</a>	Amazon EKS is now available in the EU (Stockholm) (eu-north-1) region.	December 11, 2018
<a href="#">Amazon EKS platform version update</a>	New platform version updating Kubernetes to patch level 1.10.11 to address <a href="#">CVE-2018-1002105</a> .	December 4, 2018
<a href="#">Added version 1.0.0 support for the Application Load Balancer ingress controller</a>	The Application Load Balancer ingress controller releases version 1.0.0 with formal support from AWS.	November 20, 2018
<a href="#">Added support for CNI network configuration</a>	The Amazon VPC CNI plugin for Kubernetes version 1.2.1 now supports custom network configuration for secondary pod network interfaces.	October 16, 2018
<a href="#">Added support for MutatingAdmissionWebhook and ValidatingAdmissionWebhook</a>	Amazon EKS platform version 1.10-eks.2 now supports MutatingAdmissionWebhook and ValidatingAdmissionWebhook admission controllers.	October 10, 2018
<a href="#">Added Partner AMI information</a>	Canonical has partnered with Amazon EKS to create worker node AMIs that you can use in your clusters.	October 3, 2018
<a href="#">Added instructions for AWS CLI update-kubeconfig command</a>	Amazon EKS has added the update-kubeconfig to the AWS CLI to simplify the process of creating a kubeconfig file for accessing your cluster.	September 21, 2018
<a href="#">New Amazon EKS-optimized AMIs</a>	Amazon EKS has updated the Amazon EKS-optimized AMIs (with and without GPU support) to provide various security fixes and AMI optimizations.	September 13, 2018

<a href="#">Amazon EKS region expansion (p. 185)</a>	Amazon EKS is now available in the EU (Ireland) ( <code>eu-west-1</code> ) region.	September 5, 2018
<a href="#">Amazon EKS platform version update</a>	New platform version with support for Kubernetes <a href="#">aggregation layer</a> and the <a href="#">Horizontal Pod Autoscaler</a> (HPA).	August 31, 2018
<a href="#">New Amazon EKS-optimized AMIs and GPU support</a>	Amazon EKS has updated the Amazon EKS-optimized AMI to use a new AWS CloudFormation worker node template and <a href="#">bootstrap script</a> . In addition, a new <a href="#">Amazon EKS-optimized AMI with GPU support</a> is available.	August 22, 2018
<a href="#">New Amazon EKS-optimized AMI patched for ALAS2-2018-1058</a>	Amazon EKS has updated the Amazon EKS-optimized AMI to address the CVEs referenced in <a href="#">ALAS2-2018-1058</a> .	August 14, 2018
<a href="#">Amazon EKS-optimized AMI build scripts</a>	Amazon EKS has open-sourced the build scripts that are used to build the Amazon EKS-optimized AMI. These build scripts are now available on GitHub.	July 10, 2018
<a href="#">Amazon EKS initial release (p. 185)</a>	Initial documentation for service launch	June 5, 2018

# AWS Glossary

For the latest AWS terminology, see the [AWS Glossary](#) in the *AWS General Reference*.